Conditional Latent-Prior Generative Compressed Sensing for Petrophysical Target-Log Reconstruction

Romulo Brito da Silva[1,3,*], Irineu de Azevedo Lima Neto[3]
Diego Passos[4], Herson Oliveira da Rocha[1,2,3]
Roseane Marchezi Missagia[1],
Marco Antonio Rodrigues de Ceia[1]

[1]LENEP, North Fluminense State University (UENF), Rodovia Amaral Peixoto, km 163, Imboassica, Macaé, RJ 27925-310, Brazil

[2]Polytechnic Institute (IPLOI), Campus Macaé, Federal University of Rio de Janeiro (UFRJ), Macaé, Brazil

[3]Scientific Computing Supporting Foundation (FACC), Petrópolis, Brazil

[4]Instituto Superior de Engenharia de Lisboa, Instituto Politécnico de Lisboa, Lisbon, Portugal

[*]Corresponding author at: LENEP/UENF, Rodovia Amaral Peixoto, km 163, Imboassica, Macaé, RJ 27925-310, Brazil. E-mail: romulobrito@lenep.uenf.br

Abstract

Reconstructing a petrophysical target curve is difficult when dense auxiliary logs are available along the well but the target is observed only at selected depths. Direct models can use those sparse observations as input features, yet they do not require the reconstruction to honor them as measurements. We introduce Conditional Latent-Prior Compressed Sensing with Generative Models (CLP-CSGM). A decoder trained on windows of the target curve supplies the windows the reconstruction may use. Dense auxiliary logs predict a conditional latent code, z_0. Latent refinement balances disagreement with the sparse measurements against proximity to z_0. That code is the reference for this balance, not only a starting value. A residual at the observed depths can remain. The cross-well evaluation uses public Dutch North Sea logs, with three wells for model development and one well held out. The main targets are derived: clay volume from gamma ray by the Larionov relation, and a processed, density-derived porosity. Sparse observations are coordinate subsamples of these same curves rather than independently acquired

core or laboratory measurements. CLP-CSGM has lower mean root-mean-square error than the direct autoencoder in seven of nine cross-well configurations, with the gains concentrated in the smallest training set. This comparison is a conditional advantage over that autoencoder, not superiority over every comparator. At the reference noise standard deviation 0.01, depth interpolation of the sparse measurements has lower error than CLP-CSGM in all nine configurations. Under larger perturbations that interpolation comparison changes. These noise levels were used for sensitivity analysis and were not calibrated to field measurement uncertainties. In a supplementary laboratory-porosity check, the direct autoencoder has the lower mean error. The benefit depends on the regime, and broader validation is still required.

Keywords: compressed sensing; generative model; inverse problem; well-log reconstruction; petrophysics; sparse assimilation

Computer Code Availability

CLP-CSGM (Conditional Latent-Prior Compressed Sensing with Generative Models) was developed by Romulo Brito da Silva, LENEP, North Fluminense State University, Rodovia Amaral Peixoto, km 163, Imboassica, Macae, RJ 27925-310, Brazil (contact: `romulobrito@lenep.uenf.br` ; telephone: +55 21 99101-3308). First publicly released in 2026, it is implemented in Python 3.10 using NumPy 1.26, pandas 2.3, Matplotlib 3.7, scikit-learn 1.6, PyTorch 2.6, and PyWavelets 1.4. The source repository occupies approximately 10 MB, excluding datasets and experiment outputs, and runs on a multi-core CPU with at least 8 GB of RAM; CUDA is optional.

Source code, installation instructions, dependencies, benchmark scripts, and documentation are available at `https://github.com/romulobrito/CLP--CSGM` under the MIT License. The repository covers the cross-well clay-volume and F03-4 porosity benchmarks, structural ablations, latent-prior sensitivity, and supplementary diagnostics. Processed well-log tables are not redistributed; the README specifies the required inputs and reproduction steps.

Data availability. The F3-derived benchmarks follow Zhang et al. [2025]; users must obtain the tabulated inputs from the cited releases or portals under the applicable licenses. The Lapa dataset follows Alvarez et al. [2025], whose underlying field data cannot be shared publicly. Supplementary analyses therefore require locally authorized inputs.

1 Introduction

Petrophysical target-log reconstruction builds a continuous property profile in a target well. Dense wireline logs run along the well, but they measure the target property only indirectly [Ellis and Singer, 2007, Asquith and Krygowski, 2004, Wang et al., 2024, Mondal and Singh, 2022, Sun et al., 2026]. A few depths already carry a value of that property. The task is to reconstruct the remaining depths while penalizing disagreement with those known values.

Supervised regressors learn a nonlinear map from logs to the target property. When the sparse observations are available, they are commonly joined to the dense logs and used as extra inputs [Bishop, 2006, Ahmadi and Chen, 2019, Wood, 2020, Pan et al., 2023, Mondal and Singh, 2022, Koray et al., 2024, Matinkia et al., 2023, Abdel Azim and Aljehani, 2022, Shi et al., 2024, He et al., 2025, Markus et al., 2026]. After training, that map is not required to reproduce the observations it was given. In this work those observations stay measurements: the reconstructed window is penalized for disagreeing with them, and that penalty does not force every observation to be met exactly.

Recent studies complete a missing log interval, translate a geomodel property map, or regress porosity and shale volume from the log suite [Koeshidayatullah et al., 2026, Al-Fakih et al., 2026, 2025, 2023]. Each of those predictors returns the missing property from a trained map. Section 2 compares those trained maps with the measurement penalty used here.

Using that penalty still requires a representation of the target window. A fixed sparsity basis is a weak description of this window when the logs change from one well to another [Donoho, 2006, Candes et al., 2006a,b, Bora et al., 2017]. Conditional Latent-Prior CSGM (CLP-CSGM) learns which target windows the decoder can produce. The conditional map predicts a latent code $\boldsymbol{z}_0$ from the dense auxiliary logs. This code provides an informed initialization and serves as the reference for the quadratic penalty during latent refinement. Latent refinement balances agreement with the sparse observations against deviation from $\boldsymbol{z}_0$. A residual at the observed depths can remain. The decoder range and that penalty are separate: the range limits which windows can appear, and the penalty keeps the latent code near $\boldsymbol{z}_0$. Section 3.3 states this step, including the first trajectory and the perturbed restarts. Ridge is the primary reference for the map from dense logs to $\boldsymbol{z}_0$, because it is a linear low-capacity map, not because it is best in every regime. An MLP is tested with the

same pipeline.

The contributions are threefold. First, the dense logs predict $\boldsymbol{z}_0$, which initializes the first latent trajectory and remains the center of the quadratic penalty, while the sparse values enter as a penalized measurement residual. Second, the objective has a MAP/Tikhonov reading and a local error bound that separates decoder error, a code $\boldsymbol{z}_0$ far from a good code, noise, and whether nearby windows can be told apart from the measurements. Third, the tests use three development wells and hold out F03-4. They compare direct baselines, the decoded prediction of the conditional prior, a measurement-penalty-only search, and a ridge map against a multilayer perceptron for $\boldsymbol{z}_0$. That local bound separates those error sources under stated assumptions. It is not a verified S-REC guarantee for the trained decoder.

2 Related work

2.1 Petrophysical log prediction

Well logs provide dense but indirect measurements of rock and fluid properties, and data-driven methods are widely used to predict petrophysical targets [Ellis and Singer, 2007, Asquith and Krygowski, 2004, Ahmadi and Chen, 2019, Wood, 2020]. Recent studies address three neighboring tasks.

Log completion reconstructs missing intervals of a wireline curve from that curve. Sequence-based GANs generate synthetic logs and fill gaps in gamma ray, sonic, neutron porosity, deep resistivity, and bulk density [Al-Fakih et al., 2025]. A fine-tuned time-series foundation model forecasts the next interval of the same suite and flags anomalies, including a zero-shot check in a second basin [Koeshidayatullah et al., 2026]. The unknown is a missing piece of the log, and the trained predictor fills it in one pass.

Geomodel translation reconstructs a reservoir-property image from another image. A Pix2Pix conditional GAN maps two-dimensional slices of facies, porosity, permeability, and water saturation, taking either a masked image of the same property or a different property as generator input [Al-Fakih et al., 2026]. Agreement with the supplied image is learned in training.

Direct property regression maps wireline logs to a petrophysical target. ANN and random-forest models predict total porosity, effective porosity, and shale volume from logs that include neutron

porosity [Al-Fakih et al., 2023]. Missing log values are treated in preprocessing. At inference the model returns the property from the logs, and target values act as training labels.

CLP-CSGM recovers a window of a target property when dense logs are available along the window and only sparse observations of that property are in hand. The conditional map predicts $\boldsymbol{z}_0 = h(\boldsymbol{u})$ from the dense logs. This code initializes the first latent trajectory and remains the center of the quadratic penalty. The sparse values are measurements $\boldsymbol{b}$. After training, latent refinement balances agreement with $\boldsymbol{b}$ against deviation from $\boldsymbol{z}_0$. A residual at the observed depths can remain. Log completion and geomodel translation fill a curve or a map from context. Direct regression predicts the property from the logs. CLP-CSGM keeps an explicit measurement penalty at inference. The difference is which quantity is unknown. In log completion, geomodel translation, and direct regression, the missing property values are what the trained model must produce. In CLP-CSGM those sparse values are already known measurements, and the unknown is the remainder of the target window. The inference step penalizes disagreement with those measurements.

The experiments use the processed F3 block data of Zhang et al. [2025], who combine supervised prediction with Gardner and Larionov relations and SHAP-based interpretation. We use these logs to study reconstruction from dense context while penalizing disagreement with sparse target-well observations.

2.2 Classical and generative compressed sensing

Classical compressed sensing combines an explicit measurement operator with a sparsity or compressibility prior [Donoho, 2006, Candes et al., 2006a,b]. Here, M_Ω is a coordinate restriction selecting observed target depths. A fixed transform-domain prior may be restrictive for cross-well petrophysical windows, whose representation can vary across wells. Restriction operators are also common in seismic, ultrasonic, and sparse-sensor reconstruction [Herrmann and Hennenfent, 2008, Hennenfent and Herrmann, 2008, de Assis et al., 2022, Muir and Zhan, 2021, da Silva et al., 2025].

CSGM replaces fixed-basis sparsity with the range $\mathrm{Im}(G)$ of a learned generator or decoder [Bora et al., 2017]. Reconstruction searches for a latent code $\boldsymbol{z}$ whose decoded signal reduces the measurement residual. Recovery properties and limitations of this learned-range prior have been studied under set-restricted conditions and related generative formulations [Bora et al., 2017, Kamath et al., 2020, Asim et al., 2020]. CLP-CSGM retains this decoder-range view and adds

dense-log conditioning of the latent search.

2.3 Conditional generative inverse problems and positioning

Conditional approaches differ in where side information and measurements enter. Direct conditional predictors map side information, alone or with observations, to an output or latent code [Mirza and Osindero, 2014, Sohn et al., 2015], while measurement-conditional CSGM conditions the generator directly on compressed measurements and their sampling pattern [Kim et al., 2021]. Other formulations retain an explicit inverse problem: side information may regularize a measurement-consistent reconstruction [Ongie et al., 2020, Lucas et al., 2018], approximate or sample inverse posteriors [Adler and Öktem, 2018, Altekrüger et al., 2024, Bendel et al., 2024], or use score/diffusion models for inverse reconstruction [Dasgupta et al., 2024, Jacobsen et al., 2025].

CLP-CSGM belongs to this explicit conditional inverse-reconstruction setting. Dense logs predict $\boldsymbol{z}_0 = h(\boldsymbol{u})$ before the search. This code initializes the first trajectory and remains the center of the quadratic penalty. The other restarts begin at small perturbations of $\boldsymbol{z}_0$. The sparse target values enter only through the measurement residual. The experiments therefore compare a direct map from $[\boldsymbol{u}, \boldsymbol{b}]$ to the window, the decoded prediction of the conditional prior $G(\boldsymbol{z}_0)$, and a reconstruction that uses only the measurements. Measurement-conditional generators and posterior-sampling architectures remain outside the empirical scope.

3 Method

CLP-CSGM uses three pieces. A decoder, trained on windows of the target curve, supplies the windows the reconstruction may use. The conditional map predicts a latent code $\boldsymbol{z}_0$ from the dense auxiliary logs. This code provides an informed initialization and serves as the reference for the quadratic penalty during latent refinement. Latent refinement balances agreement with the sparse observations against deviation from $\boldsymbol{z}_0$. A residual at the observed depths can remain. The decoder range and the quadratic penalty are different components. This section defines the measurement model, the objective, the comparisons that drop one piece, and the error bound.

3.1 Problem setup

Let $\boldsymbol{y} \in \mathbb{R}^N$ denote a target-property window of N depth samples and let $\boldsymbol{u} \in \mathbb{R}^p$ denote the vectorized dense-log channels over the same interval. Target values are observed at the ordered coordinate set $\Omega = \{\omega_1, \ldots, \omega_m\} \subset \{1, \ldots, N\}$, where $1 \leq m < N$ and the measurement ratio is $\rho = m/N$.

The coordinate-restriction operator $M_\Omega : \mathbb{R}^N \to \mathbb{R}^m$ is represented by

$$M_\Omega = \begin{bmatrix} \boldsymbol{e}_{\omega_1}^\top \\ \vdots \\ \boldsymbol{e}_{\omega_m}^\top \end{bmatrix} \in \mathbb{R}^{m \times N}, \qquad M_\Omega \boldsymbol{y} = \begin{bmatrix} y_{\omega_1} \\ \vdots \\ y_{\omega_m} \end{bmatrix}, \tag{1}$$

where $\boldsymbol{e}_i$ is the i-th canonical basis vector of $\mathbb{R}^N$. Coordinate restriction is widely used to represent incomplete seismic, ultrasonic, and sparse-sensor observations [Herrmann and Hennenfent, 2008, de Assis et al., 2022, da Silva et al., 2025].

The sparse target measurements satisfy

$$\boldsymbol{b} = M_\Omega \boldsymbol{y} + \boldsymbol{\eta}, \qquad \boldsymbol{b}, \boldsymbol{\eta} \in \mathbb{R}^m, \tag{2}$$

where $\boldsymbol{\eta}$ is additive noise. Under the Gaussian model, $\boldsymbol{\eta} \sim \mathcal{N}(\boldsymbol{0}, \sigma_\eta^2 I_m)$; the noiseless case has $\boldsymbol{\eta} = \boldsymbol{0}$.

The dense logs and the sparse measurements do different jobs. The dense-log vector $\boldsymbol{u}$ predicts the code $\boldsymbol{z}_0 = h(\boldsymbol{u})$. The measurements $\boldsymbol{b}$ enter only as the penalized mismatch in Eq. (10). The first latent trajectory starts at $\boldsymbol{z}_0$, and the penalty center remains $\boldsymbol{z}_0$.

Standardization and modeling assumptions. All preprocessing statistics are computed from training windows only. Let $\boldsymbol{\mu}_y, \boldsymbol{s}_y$ and $\boldsymbol{\mu}_u, \boldsymbol{s}_u$ be the componentwise training means and standard deviations, with $D_y = \mathrm{diag}(\boldsymbol{s}_y)$ and $D_u = \mathrm{diag}(\boldsymbol{s}_u)$. These two matrices rescale each coordinate by its training standard deviation. They are not estimated covariance matrices of the profiles, and this rescaling leaves any correlation already present between coordinates unchanged. We use $\widetilde{\boldsymbol{y}} = D_y^{-1}(\boldsymbol{y} - \boldsymbol{\mu}_y)$ and $\widetilde{\boldsymbol{u}} = D_u^{-1}(\boldsymbol{u} - \boldsymbol{\mu}_u)$. Sparse measurements are standardized with $D_\Omega = M_\Omega D_y M_\Omega^\top$ and $\widetilde{\boldsymbol{b}} = D_\Omega^{-1}(\boldsymbol{b} - M_\Omega \boldsymbol{\mu}_y)$; near-zero standard deviations use the implementation safeguard. The

decoder G returns a standardized window. The same window in the original target units is

$$G_{\text{orig}}(\boldsymbol{z}) = \boldsymbol{\mu}_y + D_y G(\boldsymbol{z}), \tag{3}$$

and a standardized prediction returns to that scale by

$$\widehat{\boldsymbol{y}} = \boldsymbol{\mu}_y + D_y \widehat{\widetilde{\boldsymbol{y}}}. \tag{4}$$

For a fixed set of observed coordinates, Eq. (2) adds noise with $\boldsymbol{\eta} \sim \mathcal{N}(\mathbf{0}, \sigma_\eta^2 I_m)$. Independence across coordinates and a common variance are assumptions of this synthetic benchmark noise. They are not a demonstrated description of instrumental noise, and this manuscript does not identify them with the 30dB filename label. The standardized noise is $\widetilde{\boldsymbol{\eta}} = D_\Omega^{-1}\boldsymbol{\eta}$. Because D_Ω is diagonal,

$$\operatorname{Cov}(\widetilde{\boldsymbol{\eta}}) = \sigma_\eta^2 D_\Omega^{-1} D_\Omega^{-\top} = \sigma_\eta^2 D_\Omega^{-2}. \tag{5}$$

This diagonal transformation preserves the independence of the added Gaussian noise. When the training standard deviations differ across coordinates, the standardized variances differ as well. For $\boldsymbol{r} = M_\Omega G(\boldsymbol{z}) - \widetilde{\boldsymbol{b}}$,

$$\boldsymbol{r}^\top \operatorname{Cov}(\widetilde{\boldsymbol{\eta}})^{-1} \boldsymbol{r} = \frac{1}{\sigma_\eta^2} \|D_\Omega \boldsymbol{r}\|_2^2. \tag{6}$$

The right-hand side is the data term of Eq. (10) divided by σ_η^2. Proposition 1 places the MAP weight σ_η^2/σ_z^2 on that original-scale term. The factor $\sigma_z^2 I_k$ in that proposition is the isotropic hypothesis behind the Euclidean latent penalty. It is not an estimated covariance of the residuals of h. Ridge standardizes the latent codes while fitting h and returns $\boldsymbol{z}_0$ to the encoder scale.

Where a later equation omits tildes, its coordinates are those written there. The refinement objective keeps the standardized residual and the diagonal weight D_Ω. The decoder $G : \mathbb{R}^k \to \mathbb{R}^N$, with $k < N$, supplies standardized windows, and G_{orig} supplies the corresponding windows in the original target units. We assume that those original-scale windows are adequately approximated by $\operatorname{Im}(G_{\text{orig}})$, and that $h : \mathbb{R}^p \to \mathbb{R}^k$ provides a useful anchor $\boldsymbol{z}_0 = h(\boldsymbol{u})$. These assumptions are assessed empirically through reconstruction, ablations, and prior sensitivity.

3.2 Learned model components

Offline training fits an autoencoder on standardized target windows and a latent-prior predictor on paired dense-log windows and encoded targets. All learned components remain frozen during target-well reconstruction.

Generative decoder. Let $E_{\theta_E} : \mathbb{R}^N \to \mathbb{R}^k$ and $G_{\theta_G} : \mathbb{R}^k \to \mathbb{R}^N$, with $k < N$, denote the encoder and decoder. Given standardized training targets $\{\widetilde{\boldsymbol{y}}_i\}_{i=1}^{n_{\mathrm{tr}}}$, their parameters are obtained from

$$(\widehat{\theta}_E, \widehat{\theta}_G) \in \underset{\theta_E,\theta_G}{\arg\min} \frac{1}{n_{\mathrm{tr}}} \sum_{i=1}^{n_{\mathrm{tr}}} \|\widetilde{\boldsymbol{y}}_i - G_{\theta_G}\left(E_{\theta_E}\left(\widetilde{\boldsymbol{y}}_i\right)\right)\|_2^2 + \gamma_{\mathrm{AE}}\left(\|\theta_E\|_2^2 + \|\theta_G\|_2^2\right), \tag{7}$$

where $\gamma_{\mathrm{AE}} \geq 0$ is the regularization weight. After training, we write $E = E_{\widehat{\theta}_E}$ and $G = G_{\widehat{\theta}_G}$. The frozen decoder defines the learned target-window class

$$\mathcal{M} = \mathrm{Im}(G) = \left\{G(\boldsymbol{z}) : \boldsymbol{z} \in \mathbb{R}^k\right\} \subseteq \mathbb{R}^N. \tag{8}$$

This set is the range of the standardized decoder. The windows returned to the original target units form $\mathrm{Im}(G_{\mathrm{orig}})$. Target-well reconstruction therefore optimizes only $\boldsymbol{z} \in \mathbb{R}^k$ and produces $\widehat{\boldsymbol{y}} = G_{\mathrm{orig}}(\widehat{\boldsymbol{z}})$; it does not update the network parameters or search over arbitrary vectors in $\mathbb{R}^N$. The encoder is used to train the decoder and to build the codes on which the dense-log map is fit. Inference does not encode the unknown window. In the reported model, $k = 16$ is the dimension of that code. A reconstruction $G(E(\boldsymbol{y}))$ that is given the full target window checks whether the decoder can represent that window. It uses target values that inference does not have, so it is not one of the compared predictors. Appendix B.12 reports that check.

Log-conditioned latent prior. Each training target is encoded as $\boldsymbol{z}_i = E(\widetilde{\boldsymbol{y}}_i)$, and $h : \mathbb{R}^p \to \mathbb{R}^k$ is fitted from $\{(\widetilde{\boldsymbol{u}}_i, \boldsymbol{z}_i)\}_{i=1}^{n_{\mathrm{tr}}}$. For the Ridge variant, $h(\widetilde{\boldsymbol{u}}) = W\widetilde{\boldsymbol{u}} + \boldsymbol{c}$, with parameters

$$(\widehat{W}, \widehat{\boldsymbol{c}}) \in \underset{W,\boldsymbol{c}}{\arg\min} \left\{ \frac{1}{n_{\mathrm{tr}}} \sum_{i=1}^{n_{\mathrm{tr}}} \|\boldsymbol{z}_i - W\widetilde{\boldsymbol{u}}_i - \boldsymbol{c}\|_2^2 + \alpha\|W\|_F^2 \right\}, \tag{9}$$

where $\alpha \geq 0$ controls the Ridge penalty [Hoerl and Kennard, 1970]. After fitting, $h(\widetilde{\boldsymbol{u}}) = \widehat{W}\widetilde{\boldsymbol{u}} + \widehat{\boldsymbol{c}}$,

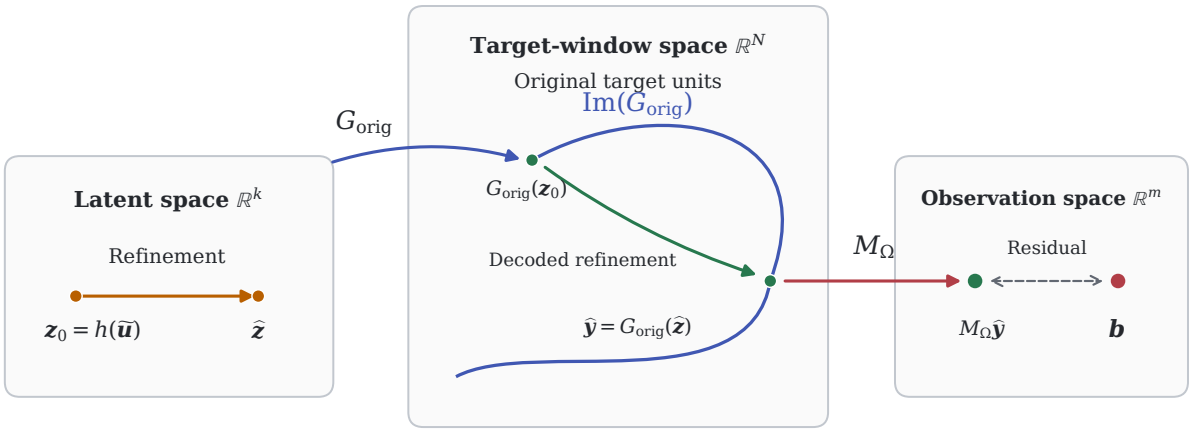


Figure 1: Geometry of the conditional latent code, in the original target units. Dense logs determine $z_0 = h(u)$. That code initializes the first trajectory and remains the center of the quadratic penalty. Test-time refinement updates the latent code to $\hat{z}$. The decoded window $\hat{y} = G_{\mathrm{orig}}(\hat{z})$ belongs to $\mathrm{Im}(G_{\mathrm{orig}})$. M_Ω selects the observed coordinates. The measurement residual is the gap between $M_\Omega \hat{y}$ and b. Refinement balances that residual against deviation from z_0, and the observations need not be met exactly.

and, suppressing tildes, the contextual anchor is $z_0 = h(u) \in \mathbb{R}^k$. The decoder G determines which target windows can be represented, and h uses dense-log context to center the latent search. CLP-CSGM is this scheme; the class of h is a configuration choice. Ridge is the primary reference because it is a linear map from the dense logs to z_0, not because it is the best map in every test. An MLP variant of h is reported with the same pipeline in Section 6.4.

3.3 Test-time reconstruction

At inference time, G and h remain fixed, and only the latent code of the current target window is optimized. Given standardized dense logs $\tilde{u} \in \mathbb{R}^p$ and sparse target observations $\tilde{b} \in \mathbb{R}^m$, define $z_0 = h(\tilde{u})$. CLP-CSGM computes

$$\begin{aligned} \hat{z}_{\mathrm{num}} &\in \arg\min_{z \in \mathbb{R}^k} \left\{ \left\| D_\Omega \left(M_\Omega G(z) - \tilde{b} \right) \right\|_2^2 + \lambda \left\| z - z_0 \right\|_2^2 \right\}, \\ \hat{y} &= G_{\mathrm{orig}} \left(\hat{z}_{\mathrm{num}} \right). \end{aligned} \tag{10}$$

The first term penalizes disagreement with the sparse observations and the second penalizes

deviation from $\boldsymbol{z}_0$. Their sum does not force the observations to be met exactly, so a residual at the observed depths can remain. Because each row of M_Ω selects one window coordinate,

$$\left\| D_\Omega \left(M_\Omega G(\boldsymbol{z}) - \widetilde{\boldsymbol{b}} \right) \right\|_2^2 = \| M_\Omega G_{\mathrm{orig}}(\boldsymbol{z}) - \boldsymbol{b} \|_2^2 . \tag{11}$$

The left-hand side is the data term in Eq. (10). The right-hand side is the squared residual in the original target units, which is the residual evaluated by the implementation. The diagonal entries of D_y are the training standard deviations of the window coordinates. On the cross-well V_c training windows those entries are not equal: the ratio of the largest to the smallest is about 1.13 at step 8, 1.23 at step 16, and 1.40 at step 32. A single factor therefore does not convert a penalty weight between the unweighted standardized residual and this objective. The implementation also divides both terms by the number of windows. Multiplying the entire objective by the same positive factor leaves its minimizers unchanged. $\lambda \geq 0$ is selected on validation windows for each measurement ratio and seed. Proposition 1 identifies λ with σ_η^2/σ_z^2 when the likelihood is an isotropic Gaussian in those original target units. That is the scale of Eq. (11). The same ratio does not apply to the unweighted standardized residual $\|M_\Omega G(\boldsymbol{z}) - \widetilde{\boldsymbol{b}}\|_2^2$. The reported experiments do not compute λ from those variances. For each seed and each measurement ratio, validation full-window RMSE selects one candidate from $\{10^{-4}, 3 \cdot 10^{-4}, 10^{-3}, 3 \cdot 10^{-3}, 10^{-2}, 3 \cdot 10^{-2}, 10^{-1}\}$, and that choice is fixed before the test score is computed. The selected λ is the validation choice for the original-scale objective. Because G is nonlinear, the objective is generally nonconvex and is solved with Adam [Kingma and Ba, 2015]. The first run starts from $\boldsymbol{z}_0$. Each additional restart starts at $\boldsymbol{z}_0$ plus an independent draw from $\mathcal{N}(\mathbf{0}, 0.05^2 I)$, with the same shape as $\boldsymbol{z}_0$. The penalty center remains $\boldsymbol{z}_0$ on every restart and every iteration, and the lowest-objective candidate is retained. Thus $\widehat{\boldsymbol{z}}_{\mathrm{num}}$ is a numerical solution, not a certified global minimizer.

Ablation variants. The *latent-prior-only* variant suppresses refinement and uses $\widehat{\boldsymbol{z}}_{\mathrm{prior}} = \boldsymbol{z}_0$, followed by $\widehat{\boldsymbol{y}}_{\mathrm{prior}} = G_{\mathrm{orig}}(\widehat{\boldsymbol{z}}_{\mathrm{prior}})$. The *measurement-only CSGM* variant sets $\lambda = 0$ and computes

$$\widehat{\boldsymbol{z}}_{\mathrm{meas}} \in \operatorname*{arg\,min}_{\boldsymbol{z} \in \mathbb{R}^k} \left\| D_\Omega \left(M_\Omega G(\boldsymbol{z}) - \widetilde{\boldsymbol{b}} \right) \right\|_2^2 . \tag{12}$$

The latent-prior-only variant returns the decoded prediction of the conditional prior, $G(\boldsymbol{z}_0)$,

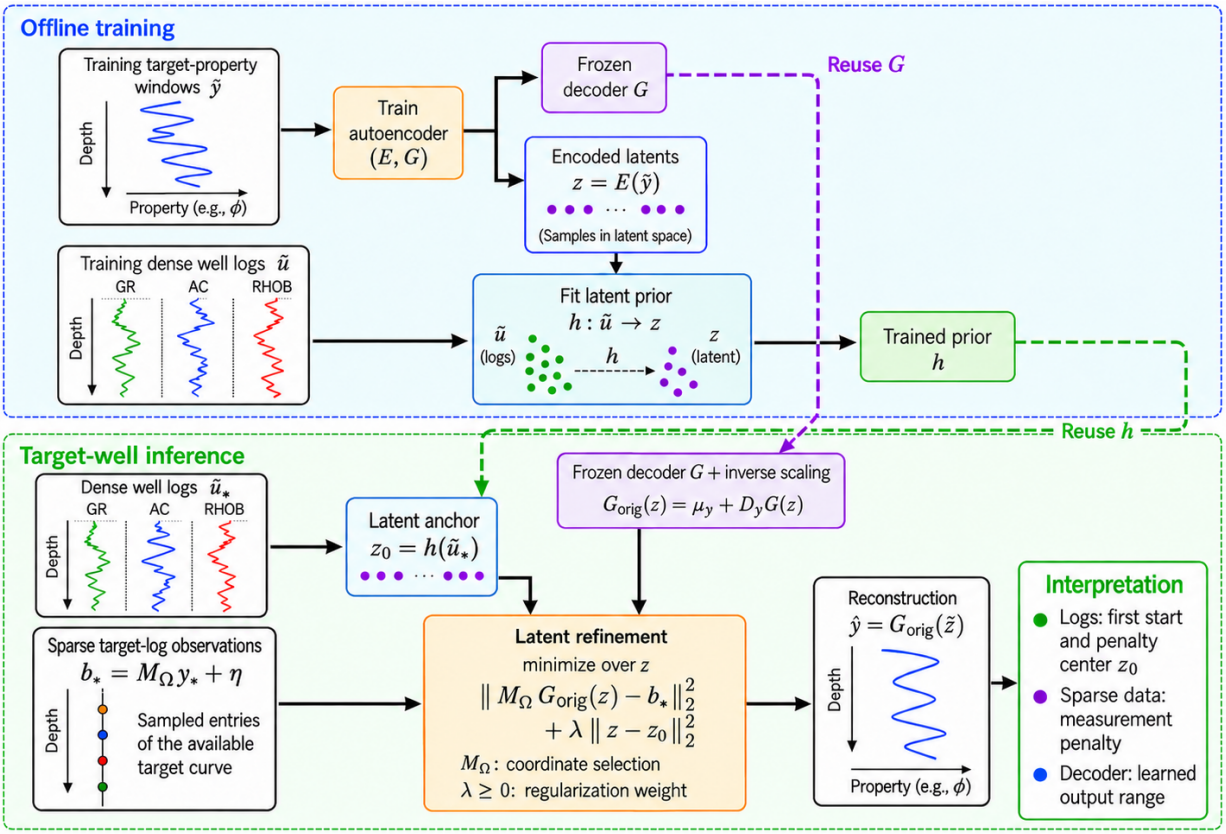


Figure 2: CLP-CSGM workflow. Offline, standardized target windows train the autoencoder (E, G), and paired dense-log windows fit the latent-prior predictor h. At inference, $\boldsymbol{u}_*$ determines $\boldsymbol{z}_0 = h(\boldsymbol{u}_*)$, sparse observations $\boldsymbol{b}_*$ enter latent refinement as a measurement penalty through M_Ω, and the decoded result is returned to the original target scale. OpenAI's ChatGPT, using the GPT-5.5 model, was used to assist in generating the initial explanatory schematic and visual layout. The final figure was reviewed, edited, and validated by the authors.

with no latent search. In that variant $\boldsymbol{z}_0$ is not an initialization, because no trajectory is optimized. The measurement-only variant uses only the sparse measurements and recovers measurement-only generative compressed sensing [Bora et al., 2017]. Comparison with the full objective in Eq. (10) shows what each piece contributes. Results are reported in Section 6.3.

3.4 Interpretation and theoretical context

The theoretical results required to interpret CLP-CSGM are stated below. Complete proofs appear in Appendix A.

Proposition 1 (MAP and Tikhonov interpretation). *Assume* $p(\boldsymbol{b} \mid \boldsymbol{z}, \boldsymbol{u}) = p(\boldsymbol{b} \mid \boldsymbol{z})$ *and*

$\boldsymbol{b} \mid \boldsymbol{z} \sim \mathcal{N}\left(M_\Omega G_{\text{orig}}(\boldsymbol{z}), \sigma_\eta^2 I_m\right)$, *while* $\boldsymbol{z} \mid \boldsymbol{u} \sim \mathcal{N}\left(\boldsymbol{z}_0, \sigma_z^2 I_k\right)$, *where* $\boldsymbol{z}_0 = h(\boldsymbol{u})$, $\sigma_\eta > 0$, *and*

$\sigma_z > 0$. *Then every MAP estimator satisfies*

$$\widehat{\boldsymbol{z}}_{\mathrm{MAP}} \in \operatorname*{arg\,min}_{\boldsymbol{z}\in\mathbb{R}^k} \left\{ \left\| D_\Omega \left(M_\Omega G(\boldsymbol{z}) - \widetilde{\boldsymbol{b}} \right) \right\|_2^2 + \lambda \left\| \boldsymbol{z} - \boldsymbol{z}_0 \right\|_2^2 \right\}, \qquad \lambda = \frac{\sigma_\eta^2}{\sigma_z^2}. \tag{13}$$

By Eq. (11)*, this data term is the original-scale residual. The ratio* σ_η^2/σ_z^2 *is the MAP weight on that scale. It is not the weight of the unweighted standardized residual, because the diagonal entries of* D_y *vary by coordinate. Thus, CLP-CSGM is a Tikhonov-regularized latent inverse problem. Its center is the log-conditioned anchor, not the origin [Tikhonov, 1963, Burger and Lucka, 2014]. The likelihood induces the data-fidelity term, while the conditional prior induces the latent penalty.*

Remark 1 (Relation to classical CSGM). Setting $\lambda = 0$ recovers the measurement-only CSGM problem in Eq. (12) [Bora et al., 2017]. Increasing λ shifts the reconstruction toward the decoded prediction of the conditional prior, $G(\boldsymbol{z}_0)$, by penalizing departures from $\boldsymbol{z}_0 = h(\boldsymbol{u})$ more strongly.

Proposition 2 (Existence of a minimizer)**.** *Assume that* $G : \mathbb{R}^k \to \mathbb{R}^N$ *is continuous and* $\lambda > 0$*. For fixed* M_Ω*,* D_Ω*,* $\widetilde{\boldsymbol{b}} \in \mathbb{R}^m$*, and* $\boldsymbol{z}_0 \in \mathbb{R}^k$*, the objective* $\mathcal{J}(\boldsymbol{z}) = \left\| D_\Omega \left(M_\Omega G(\boldsymbol{z}) - \widetilde{\boldsymbol{b}} \right) \right\|_2^2 + \lambda \|\boldsymbol{z} - \boldsymbol{z}_0\|_2^2$ *is continuous and coercive. Consequently,*

$$\operatorname*{arg\,min}_{\boldsymbol{z}\in\mathbb{R}^k} \mathcal{J}(\boldsymbol{z}) \neq \varnothing. \tag{14}$$

For any global minimizer $\widehat{\boldsymbol{z}}$, $\mathcal{J}(\widehat{\boldsymbol{z}}) \leq \mathcal{J}(\boldsymbol{z}_0) = \left\| D_\Omega \left(M_\Omega G(\boldsymbol{z}_0) - \widetilde{\boldsymbol{b}} \right) \right\|_2^2$. Exact minimization therefore cannot worsen the optimized objective relative to the latent-prior-only candidate, although it does not guarantee lower error with respect to the unknown target. For $\widehat{\boldsymbol{z}}_{\mathrm{num}}$, this comparison holds only when $\mathcal{J}(\widehat{\boldsymbol{z}}_{\mathrm{num}}) \leq \mathcal{J}(\boldsymbol{z}_0)$.

Remark 2 (Representation error and decoder mismatch). Every CLP-CSGM reconstruction in the original target units satisfies $\widehat{\boldsymbol{y}} = G_{\mathrm{orig}}(\widehat{\boldsymbol{z}}) \in \mathrm{Im}(G_{\mathrm{orig}})$. Consequently, for an unknown target window $\boldsymbol{y}^\star \in \mathbb{R}^N$, the error of any such reconstruction is bounded below by the oracle representation error

$$\mathrm{dist}\left(\boldsymbol{y}^\star, \mathrm{Im}(G_{\mathrm{orig}})\right) := \inf_{\boldsymbol{y}'\in\mathrm{Im}(G_{\mathrm{orig}})} \left\| \boldsymbol{y}^\star - \boldsymbol{y}' \right\|_2 = \inf_{\boldsymbol{z}\in\mathbb{R}^k} \left\| \boldsymbol{y}^\star - G_{\mathrm{orig}}(\boldsymbol{z}) \right\|_2. \tag{15}$$

Thus, even with noiseless observations and exact global optimization, latent refinement cannot

remove mismatch between $\boldsymbol{y}^\star$ and $\mathrm{Im}(G_{\mathrm{orig}})$. This structural error is represented by $\boldsymbol{r}^\dagger$ in Proposition 3 and is distinct from the latent-prior mismatch $\|\boldsymbol{z}^\dagger - \boldsymbol{z}_0\|_2$.

An inaccurate code $\boldsymbol{z}_0 = h(\boldsymbol{u})$ may bias the search toward a suboptimal latent region, because it initializes the first trajectory and remains the center of the penalty. The data-fidelity term can partially correct this bias when the sparse observations are informative and λ does not dominate the objective. The decoder range determines what can be represented; the conditional code $\boldsymbol{z}_0$ determines where the latent penalty is centered.

Proposition 3 (Conditional local recovery interpretation)**.** *Let $\boldsymbol{y}^\star \in \mathbb{R}^N$ satisfy $\boldsymbol{b} = M_\Omega \boldsymbol{y}^\star + \boldsymbol{\eta}$, let $\boldsymbol{z}_0 = h(\boldsymbol{u})$, $\lambda > 0$, and let $\widehat{\boldsymbol{z}}$ be any global minimizer of Eq.* (10)*. The residual $\|M_\Omega G_{\mathrm{orig}}(\boldsymbol{z}) - \boldsymbol{b}\|_2^2$ is the right-hand side of Eq.* (11)*. For any reference code $\boldsymbol{z}^\dagger \in \mathbb{R}^k$, define*

$$\begin{aligned} \boldsymbol{r}^\dagger &= \boldsymbol{y}^\star - G_{\mathrm{orig}}(\boldsymbol{z}^\dagger), & \boldsymbol{e}^\dagger &= M_\Omega \boldsymbol{r}^\dagger + \boldsymbol{\eta}, \\ \Delta^\dagger &= \left(\|\boldsymbol{e}^\dagger\|_2^2 + \lambda \|\boldsymbol{z}^\dagger - \boldsymbol{z}_0\|_2^2\right)^{1/2}, & R^\dagger &= \frac{\Delta^\dagger}{\sqrt{\lambda}}. \end{aligned} \tag{16}$$

Let $\mathcal{M}_{R^\dagger} = G_{\mathrm{orig}}(B(\boldsymbol{z}_0, R^\dagger))$, where $B(\boldsymbol{z}_0, R^\dagger) = \left\{\boldsymbol{z} \in \mathbb{R}^k : \|\boldsymbol{z} - \boldsymbol{z}_0\|_2 \le R^\dagger\right\}$. The constants $\gamma_R > 0$ and $\delta_R \ge 0$ below refer to pairs of windows in this original-scale image. They are not constants of the standardized decoder G. If

$$\|M_\Omega(\boldsymbol{y}_1 - \boldsymbol{y}_2)\|_2 \ge \gamma_R \|\boldsymbol{y}_1 - \boldsymbol{y}_2\|_2 - \delta_R, \qquad \forall\, \boldsymbol{y}_1, \boldsymbol{y}_2 \in \mathcal{M}_{R^\dagger}, \tag{17}$$

then

$$\|G_{\mathrm{orig}}(\widehat{\boldsymbol{z}}) - \boldsymbol{y}^\star\|_2 \le \|\boldsymbol{r}^\dagger\|_2 + \frac{\Delta^\dagger + \|\boldsymbol{e}^\dagger\|_2 + \delta_R}{\gamma_R}. \tag{18}$$

Proposition 3 is a sufficient conditional interpretation, not a verified guarantee for the trained decoder. The latent penalty localizes the S-REC condition around $\boldsymbol{z}_0$, while the bound separates representation error, measurement residual/noise, prior mismatch, and local separability, following the decoder-range view of generative compressed sensing [Candes et al., 2006b, Bora et al., 2017]. Coordinate restriction is not proved to satisfy S-REC on the trained range; this depends on local decoder secants. The proof, inexact-optimizer extension, and empirical separability diagnostics appear in Appendix A.3 and Appendix B.1.

4 Data and exploratory analysis

We evaluate two F3 settings. In the first, clay volume is reconstructed in a held-out well. In the second, porosity in one well is reconstructed from gamma ray alone.

4.1 Field settings and benchmark provenance

The experiments use the processed F3 well-log curves of Zhang et al. [2025]; no new field data were acquired. We use the FS2–FS6 intervals of F02-1, F03-2, F03-4, and F06-1. Sonic and gamma ray are available in all wells; density is measured in F02-1/F03-2 and predicted in F03-4/F06-1. Porosity is density-derived and V_c is Larionov-derived from gamma ray; further geological and processing details are given in the source study.

The cross-well benchmark reconstructs V_c from 30 dB noisy logs, using F02-1, F03-2, and F06-1 for development and F03-4 for testing. The single-well benchmark reconstructs F03-4 porosity from gamma ray alone. In both, $\boldsymbol{b}$ is obtained by coordinate subsampling of the processed target. Appendix D provides a limited-sample Lapa permeability case based on Alvarez et al. [2025].

Table 1: Summary of the two F3 benchmarks. The cross-well target V_c is derived from gamma ray through the Larionov relation, whereas the single-well target is processed porosity.

dataset	well	target	n_rows	depth_min	depth_max	target_mean	target_std	target_min	target_max
cross_train	F02-1	V_c	165	1149.7056	1174.6992	0.1215	0.0407	0.0513	0.1934
cross_train	F03-2	V_c	811	951.9000	1073.4000	0.1184	0.0490	0.0393	0.2688
cross_train	F06-1	V_c	122	1133.1000	1151.2500	0.1223	0.0437	0.0427	0.2175
cross_test	F03-4	V_c	832	946.5000	1071.1500	0.0824	0.0341	0.0112	0.1699
real_f03_gr	F03-4	Porosity	1435	856.0500	1071.1500	0.2789	0.0142	0.2471	0.3264

Table 1 summarizes the two settings. Held-out F03-4 has mean $V_c = 0.0824$, versus about 0.12 in the training wells, and Fig. 3 shows a lower-valued but overlapping distribution. Larger sliding-window steps further reduce training support, whereas CLP-CSGM can still refine the held-out-well estimate using $\boldsymbol{b}$ at inference.

Channel-target correlations and input design. Table 2 reports descriptive Pearson correlations used only to guide input design. Clay volume in this benchmark is computed from gamma ray by the Larionov relation [Zhang et al., 2025]. The main cross-well input therefore leaves gamma ray out and uses sonic, density, acoustic impedance, and V_p. With gamma ray included, the model would read the log from which the target was computed. This is not the usual field interpretation, in

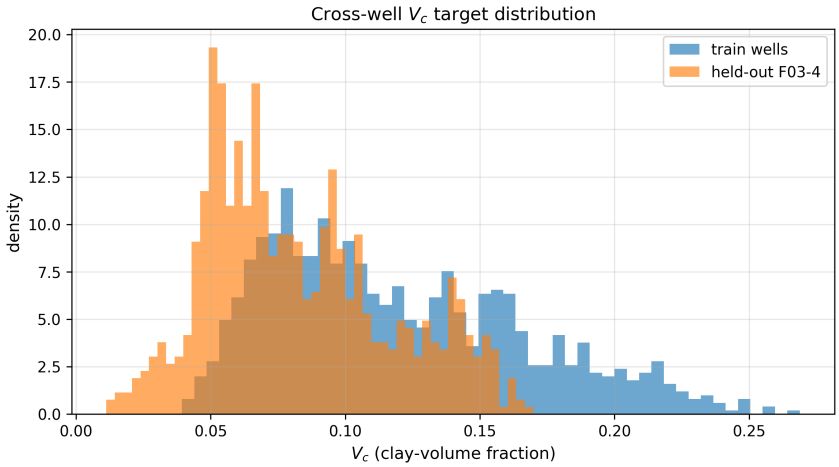


Figure 3: Empirical cross-well V_c distributions for the pooled training wells and the held-out F03-4 well. The held-out distribution is shifted toward lower clay-volume values while retaining partial overlap with the training distribution.

which gamma ray is logged and clay volume is obtained from that relation directly. The gamma-ray index between the clean and shale endpoints is linear in gamma ray. The Larionov map from that index to clay volume is not a straight line. When gamma ray has been logged, that relation is the calculation, and the model is not a substitute for it. Gamma ray is put back only in Appendix B.7, as a check on the effect of that channel.

Table 2: Pearson correlations between the available input channels and the target properties in the cross-well V_c and single-well F03-4 porosity benchmarks.

dataset	channel	target	pearson_r
cross_train	Sonic	V_c	-0.1642
cross_train	RHOB	V_c	0.0807
cross_train	GR	V_c	0.9733
cross_train	AI	V_c	0.1672
cross_train	V_p	V_c	0.1819
cross_test	Sonic	V_c	-0.1606
cross_test	RHOB	V_c	-0.0061
cross_test	GR	V_c	0.9772
cross_test	AI	V_c	0.1330
cross_test	V_p	V_c	0.1077
real_f03	AC	Porosity	0.9996
real_f03	GR	Porosity	-0.1759

In F03-4, sonic travel time follows the processed porosity almost one to one, whereas gamma ray is only weakly related to it. Sonic is left out for that reason. The density curve behind this porosity was itself predicted from sonic and gamma ray. The test that remains uses gamma ray as the only dense log and keeps a few values of the porosity curve as measurements. It is not a porosity

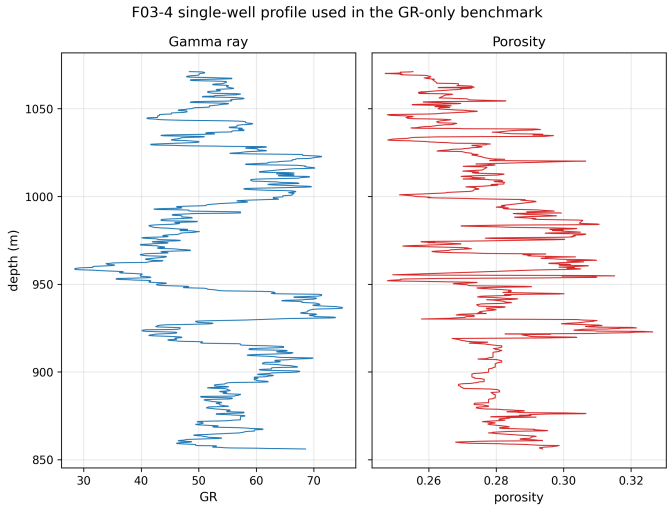


Figure 4: F03-4 depth profiles used in the GR-only porosity benchmark. Gamma ray is the only dense input, while porosity is reconstructed from the log-conditioned latent prior and sparse target observations.

interpretation from density, neutron, and sonic together (Fig. 4).

5 Experimental setup

The evaluation uses paired protocols for the cross-well low-data V_c benchmark and the single-well F03-4 GR-only porosity benchmark. Within each cell, all methods share the same partitions, measurement ratios, random seeds, sampled coordinates, and noise realizations.

5.1 Cross-well clay-volume (V_c) low-data benchmark

The cross-well benchmark reconstructs length-$N = 64$ processed V_c windows in the held-out F03-4 well. F02-1, F03-2, and F06-1 provide the model-development data, and the dense input $\boldsymbol{u}$ concatenates sonic, bulk density, acoustic impedance, and V_p. Reported cross-well performance refers to this single held-out well. As justified in Section 4, gamma ray is left out because clay volume was computed from that log [Zhang et al., 2025]. This is not the log suite for a well in which gamma ray is available and clay volume is obtained from it directly.

Sparse observations follow $\boldsymbol{b} = M_\Omega \boldsymbol{y} + \boldsymbol{\eta}$ at $\rho \in \{0.05, 0.10, 0.20\}$. Sliding-window steps $\{8, 16, 32\}$ control training support, and each step–ρ cell uses three paired seed realizations. This is a controlled sparse-assimilation benchmark on a processed target curve, not recovery against

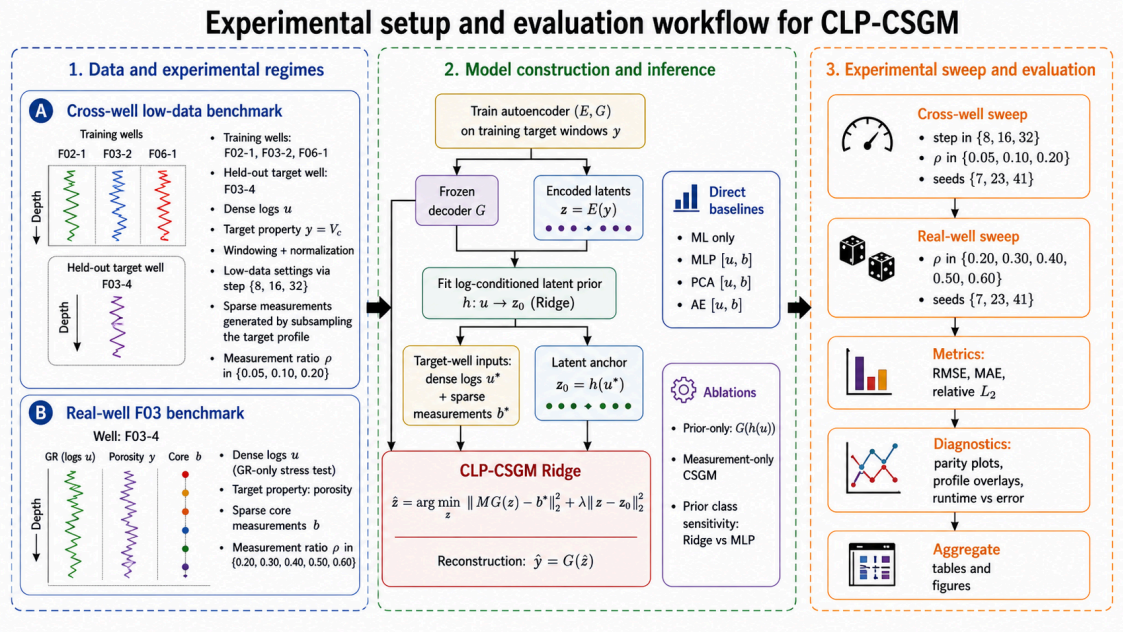


Figure 5: Experimental protocol. Offline training fits the decoder and the map from dense logs to z_0. At test time the first trajectory starts at that code and the penalty remains centered there. Partitions, measurement ratios, seeds, and sparse-observation draws are shared across methods, which are evaluated using RMSE, MAE, and relative ℓ_2 error. OpenAI's ChatGPT, using the GPT-5.5 model, was used to assist in generating the initial explanatory schematic and visual layout. The final figure was reviewed, edited, and validated by the authors.

an independently acquired measurement. Split and sampling details are given in Section 5.3; GR inclusion and the embargo-style boundary check are reported in Appendix B.7.

5.2 Single-well F03-4 GR-only porosity benchmark

The second benchmark reconstructs length-$N = 64$ windows of the processed density-derived F03-4 porosity curve using gamma ray as the only dense input. Acoustic travel time is excluded because it is nearly collinear with the target (Pearson $r \approx 0.9996$) and would provide a near-direct proxy. Density porosity is linear in bulk density, and a standard sonic porosity is linear in slowness. If those logs have been recorded, the processed porosity is already determined, and this test is not a substitute for those calculations. Gamma ray is then the only dense log, so the dense context is intentionally incomplete. The target follows the processing of Zhang et al. [2025], in which the density curve underlying F03-4 porosity is itself predicted from sonic and gamma ray.

Sparse observations are sampled at $\rho \in \{0.20, 0.30, 0.40, 0.50, 0.60\}$. Depth is partitioned before window construction, preventing overlap among training, validation, and test windows. The bench-

mark is a controlled test with gamma ray as the only dense log. It is not a recommendation to interpret porosity from gamma ray alone, and its errors do not transfer to a suite that also includes density, neutron, and sonic.

The smaller Lapa pre-salt permeability case [Alvarez et al., 2025] is retained as supplementary evidence in Appendix D. Appendix E reports a separate single-well check on laboratory porosity. The direct AE has the lower mean held-out RMSE. Laboratory measurements correct the decoded prediction of the conditional prior in every split, and that correction is enough where the test logs remain close to the support used to fit $h(\boldsymbol{u})$. The check is not part of the main grid.

5.3 Training, validation, and test protocol

All preprocessing statistics and model choices are determined without test-target leakage. Training statistics are reused unchanged on validation and test windows. The training split fits the direct baselines, autoencoder (E, G), and latent-prior predictor h. Validation targets select λ and baseline-specific hyperparameters, and test targets are used only for final evaluation.

In the cross-well benchmark, F02-1, F03-2, and F06-1 form the development pool, F03-4 is held out entirely for testing, and 10% of the development windows are reserved for validation. Steps $\{8, 16, 32\}$ produce approximately $\{104, 52, 27\}$ windows before this split. At step 32 the development pool is small; Appendix B.8 redraws that pool ten times with the measurement pattern held fixed. CLP-CSGM Ridge has lower RMSE than AE $[\boldsymbol{u}, \boldsymbol{b}]$ at every ρ and in all ten splits. Because the validation tail is selected after window construction, boundary windows may overlap when step $< N$; Appendix B.7 reports an embargo-style check that removes these overlap-prone windows.

For F03-4 porosity, depth samples are divided into contiguous training, validation, and test blocks using fractions 60%, 20%, and 20% before windows are constructed independently. This depth-first partition prevents overlap and preserves contiguous-depth generalization.

For each seed and measurement ratio ρ, the number of observed coordinates is

$$m = \max\{4, \mathrm{round}(\rho L)\}.$$

We use nearest-integer rounding, with half-integer ties resolved to the nearest even integer. One

coordinate set $\Omega = \{\omega_1, \ldots, \omega_m\}$ is sampled for each (seed, ρ) cell and held fixed across all training, validation, and test windows in that cell. The set defines

$$\boldsymbol{b} = \begin{bmatrix} y_{\omega_1} & \cdots & y_{\omega_m} \end{bmatrix}^\top + \boldsymbol{\eta}. \tag{19}$$

Within each (seed, ρ) cell, all methods share the partitions, the fixed sampled coordinates, their ordering, and the noise realization. Direct baselines receive $[\boldsymbol{u}, \boldsymbol{b}] \in \mathbb{R}^{p+m}$, with only the measured target values in the ordering of M_Ω. Because Ω is fixed within the cell, the j-th entry of $\boldsymbol{b}$ consistently corresponds to coordinate ω_j across windows, making coordinate identity implicit in the stable input ordering. No zero-padded target, binary mask, or measurement matrix is added. CLP-CSGM uses M_Ω explicitly in the measurement term of Eq. (10). The protocol therefore compares direct fusion with operator-based assimilation under a fixed sampling pattern, giving the direct models a stable coordinate-to-input mapping within each cell. Appendix B.10 gives one direct AE the same measurements on the window coordinates plus a binary mask. That mask does not remove the CLP-CSGM advantage at steps 16 and 32; at step 8 the masked AE remains better at $\rho = 0.05$ and $\rho = 0.10$, and the mask does not uniformly lower the direct AE error. This comparator is not part of the main grid.

Appendix B.11 redraws the coordinate set five times with the model seeds held fixed. Where training support is smallest (steps 16 and 32), CLP-CSGM Ridge keeps the lower RMSE in every draw. At step 8, where the development pool is largest, the ranking is less stable: AE $[\boldsymbol{u}, \boldsymbol{b}]$ stays ahead at $\rho = 0.05$ and on the majority of draws at $\rho = 0.10$, while CLP-CSGM stays ahead on the majority of draws at $\rho = 0.20$. The cell-level winner still matches the published ranking in all nine cells. That check is not part of the main grid.

Appendix B.12 adds direct linear and ridge maps from $[\boldsymbol{u}, \boldsymbol{b}]$ to the window, and depth interpolation of the sparse measurements, on this same cross-well grid. It also raises the measurement-noise standard deviation above the published value 0.01. Section 6 summarizes that noise dependence for interpolation and CLP-CSGM Ridge. The linear maps and the per-seed tables are not merged into the submitted nine-cell grid.

Noise is applied after coordinate subsampling. Each observed coordinate receives an independent standard-normal draw, scaled to a fixed standard deviation and added to the subsampled target

Table 3: Implementation settings for offline training, target-well latent refinement, and direct supervised baselines.

Phase	Module	Implementation
Offline training	Standardization	Training-set mean and standard deviation for u and y; reused unchanged on validation and test windows.
	Autoencoder (E, G)	MLP with one hidden layer (width 128) in encoder and decoder, latent dimension $k = 16$, ReLU activations, linear output; 200 epochs, Adam [Kingma and Ba, 2015] (10^{-3}), weight decay 10^{-5}, batch size 64.
	Latent prior $h(u)$	Ridge regressor with $\alpha = 1.0$ on standardized (u, z) pairs (main experiments). Sensitivity variant: MLP with two hidden layers (width 128), L_2 penalty 10^{-4}, learning rate 10^{-3}, max. 500 iterations, early stopping.
Target-well inference	Latent refinement	Adam [Kingma and Ba, 2015] for 400 iterations, learning rate 0.05, three restarts (first start at $z_0 = h(u)$; two further starts at z_0 plus $\mathcal{N}(0, 0.05^2 I)$; penalty center remains z_0); λ selected by validation RMSE from $\{10^{-4}, 3{\cdot}10^{-4}, 10^{-3}, 3{\cdot}10^{-3}, 10^{-2}, 3{\cdot}10^{-2}, 10^{-1}\}$ for the original-scale residual, not computed as σ_η^2/σ_z^2.
	Reconstruction	$\hat{y} = G(\hat{z})$ after inverse standardization to the original target scale.
Direct baselines	$[u, b]$ regressors	ML only and MLP $[u, b]$: sklearn MLP with two hidden layers (width 128), max. 500 iterations, learning rate 10^{-3}, L_2 penalty 10^{-4}, early stopping. PCA $[u, b]$: PCA score grid $\{8, 16, 24, 32, 48, 64\}$ with the same MLP head. AE $[u, b]$: separate MLP autoencoder ($k = 32$, hidden width 128, 120 epochs) plus MLP from $[u, b]$ to latent codes.

value. That standard deviation is $\sigma_\eta = 0.01$ for cross-well V_c and $\sigma_\eta = 0.02$ for F03-4 porosity, in the units of the target curve. It is not a percentage of the curve. The held-out V_c curve has standard deviation about 0.035, so 0.01 is about 0.29 times that deviation. These two levels were adopted for the controlled benchmarks. No calibration to instrumental measurement uncertainty is documented. The input files are labeled 30dB. That label is not this added standard deviation, and this manuscript does not define the procedure behind the label. Within one run, every method uses the same noisy observations. Direct $[\boldsymbol{u}, \boldsymbol{b}]$ models standardize the concatenated features after the noise has been added.

5.4 Implementation details

Table 3 reports the architectures, optimization settings, and validation grids. The compact backbone uses latent dimension $k = 16$ and a decoder that remains frozen during target-well reconstruction.

For CLP-CSGM, the first latent-refinement run starts from $\boldsymbol{z}_0 = h(\boldsymbol{u})$. Each additional run

starts at $\boldsymbol{z}_0$ plus an independent draw from $\mathcal{N}(\mathbf{0}, 0.05^2 I)$. The penalty center remains $\boldsymbol{z}_0$ on every run and every iteration, and the candidate with the lowest objective is retained. The measurement-only ablation sets $\lambda = 0$ and starts from zero, whereas the prior-only variant returns the decoded prediction of the conditional prior, $G(h(\boldsymbol{u}))$, with no latent search. Predictions are returned to the original scale using Eq. (4).

5.5 Baselines and evaluation metrics

The direct baselines are ML only, which uses only $\boldsymbol{u}$; MLP $[\boldsymbol{u}, \boldsymbol{b}]$, which predicts the target directly; PCA $[\boldsymbol{u}, \boldsymbol{b}]$, which predicts target principal-component scores; and AE $[\boldsymbol{u}, \boldsymbol{b}]$, which predicts codes of a separately trained target autoencoder. CLP-CSGM Ridge is the proposed reference configuration. Prior-only Ridge and measurement-only CSGM are structural ablations, while the MLP latent prior is used only for sensitivity analysis.

Metrics are computed after inverse standardization on the original target scale. RMSE is the primary metric and selects λ on validation data; MAE and relative ℓ_2 error provide complementary summaries. Results are computed under the paired protocol for each seed and then summarized across seeds, with lower values indicating better reconstruction.

Uncertainty summaries. All seed-level intervals use the three fixed seeds $\{7, 23, 41\}$. Reported 95% half-widths are Student-t intervals with $\mathrm{df} = n_{\mathrm{seeds}} - 1 = 2$ (critical value $t_{2,0.975} \approx 4.303$), not normal $z = 1.96$ intervals. Where CLP-CSGM Ridge is compared with AE $[\boldsymbol{u}, \boldsymbol{b}]$, we also report the paired seed-wise difference

$$\Delta\mathrm{RMSE} = \mathrm{RMSE}_{\mathrm{AE}} - \mathrm{RMSE}_{\mathrm{CLP}},$$

together with its seed-level standard error and Student-t half-width. Positive ΔRMSE favors CLP-CSGM Ridge. With only three seeds, these intervals describe seed-level reproducibility within the paired protocol; their interpretation remains limited by the small seed panel. Two supplements widen that evidence and leave the three-seed tables as the main result. Appendix B.9 repeats the comparison on ten seeds. Each extra seed refits the models and draws a new coordinate set, so training randomness and observation placement change together. None of those nine intervals

include zero. The CLP-CSGM advantage remains where training support is smallest (steps 16 and 32); AE remains better at step 8 for $\rho = 0.05$ and $\rho = 0.10$. Appendix B.11 holds the same three model seeds and redraws only the coordinate set. The advantage at steps 16 and 32 holds in every draw, and the step-8 ranking is the less stable one. Overlapping windows inside a well are not counted as further independent samples.

6 Results

All comparisons follow the paired protocol of Section 5: within each (seed, ρ) cell, the methods share partitions, sampled coordinates, and noisy sparse observations. Method orderings in this section refer to the reported mean metrics across the paired seeds.

6.1 Cross-well low-data results

Table 4 compares CLP-CSGM Ridge with AE $[\boldsymbol{u}, \boldsymbol{b}]$, the direct baseline with the lowest unweighted mean RMSE across the grid. CLP-CSGM Ridge has lower mean RMSE in seven of the nine cells; the exceptions are step 8 with $\rho \in \{0.05, 0.10\}$. It has lower mean RMSE at every tested ratio for steps 16 and 32. Across the three tested step sizes, the descriptive RMSE gap shifts progressively in favor of CLP-CSGM as training support decreases. At step 32, the RMSE reductions relative to AE are approximately 32% at $\rho = 0.10$ and 57% at $\rho = 0.20$.

Across the complete grid, the unweighted mean of the nine cell-level mean RMSE values decreases from 0.03496 for AE $[\boldsymbol{u}, \boldsymbol{b}]$ to 0.02862 for CLP-CSGM Ridge, a reduction of approximately 18.1%. The ordering still depends on the amount of training data: the direct AE has lower mean RMSE in the two cells with the largest training pool, step 8 at $\rho = 0.05$ and $\rho = 0.10$.

Table 4: Cross-well V_c comparison between CLP-CSGM Ridge and AE $[\boldsymbol{u}, \boldsymbol{b}]$. Values are RMSE mean and standard deviation across seeds; a negative gap favors CLP-CSGM.

step	n_train_approx	rho	clp_csgm_rmse_mean	clp_csgm_rmse_std	ae_ub_rmse_mean	ae_ub_rmse_std	gap_vs_ae_pct
8	104	0.05000	0.04483	0.00524	0.03567	0.00100	25.69014
8	104	0.10000	0.03646	0.00194	0.03195	0.00208	14.10258
8	104	0.20000	0.02267	0.00381	0.02600	0.00164	-12.83374
16	52	0.05000	0.03344	0.00151	0.03712	0.00098	-9.91538
16	52	0.10000	0.02689	0.00235	0.03420	0.00291	-21.35910
16	52	0.20000	0.01771	0.00138	0.03214	0.00742	-44.89074
32	27	0.05000	0.03217	0.00268	0.03901	0.00259	-17.53826
32	27	0.10000	0.02640	0.00342	0.03894	0.00228	-32.21019
32	27	0.20000	0.01705	0.00119	0.03959	0.00234	-56.94200

Table 5 reports the paired seed-wise difference $\Delta\text{RMSE} = \text{RMSE}_{\text{AE}} - \text{RMSE}_{\text{CLP}}$ with Student-t

half-widths (df = 2). At steps 16 and 32, where training support is smallest, the mean paired gap is positive at every ρ and CLP-CSGM Ridge wins all three seeds; at step 8, where the development pool is largest, AE $[\boldsymbol{u}, \boldsymbol{b}]$ wins the two smallest ρ cells, matching the mean-RMSE ordering in Table 4. Several of these three-seed intervals include zero. Appendix B.9 repeats the same pairwise comparison on ten seeds: none of those intervals include zero. The CLP-CSGM advantage remains in the few-window cells (steps 16 and 32); AE remains better at step 8 for $\rho = 0.05$ and $\rho = 0.10$. Appendix B.10 repeats the three-seed grid with one extra direct AE that sees an explicit window mask. The mask does not remove the CLP-CSGM advantage at steps 16 and 32. At step 8 the masked AE remains better at $\rho = 0.05$ and $\rho = 0.10$, and the mask does not uniformly lower the direct AE error.

Appendix B.11 redraws the sparse coordinates five times while the model seeds stay at $\{7, 23, 41\}$. The published cell winner is unchanged in all nine cells. At steps 16 and 32, CLP-CSGM Ridge has lower RMSE than AE $[\boldsymbol{u}, \boldsymbol{b}]$ in every draw. At step 8, AE remains better at $\rho = 0.05$ in every draw and at $\rho = 0.10$ on the majority of draws; CLP-CSGM remains better at $\rho = 0.20$ on the majority of draws. The advantage where training support is smallest therefore holds for every tested placement. Step 8, where that support is largest, remains the less stable regime.

Appendix B.12 compares those same nine cells with ordinary least squares and ridge regression from $[\boldsymbol{u}, \boldsymbol{b}]$, and with linear interpolation of the sparse measurements along depth. The linear maps are scored at the published noise standard deviation 0.01. At that reference level, interpolation has the lowest full-window RMSE in all nine cells and in all 27 seed comparisons. CLP-CSGM Ridge is lower than both direct linear maps at step 16, $\rho = 0.20$, and at step 32 for every ρ.

Figure 6 uses the same windows and changes only the noise. Each panel is one step and each curve is one measurement fraction, so the nine cells are not averaged together. At noise 0.01, interpolation has the lower error in all nine configurations. At noise 0.05, CLP-CSGM Ridge has the lower mean RMSE at steps 16 and 32 and at step 8 for $\rho = 0.20$. At step 8 for $\rho = 0.05$ and $\rho = 0.10$, interpolation keeps the lower mean. At 0.10 and 0.20, CLP-CSGM Ridge has the lower error in all nine configurations. The held-out V_c curve has standard deviation 0.035, on the same scale as this noise. Relative to that deviation, 0.05 is about 1.43 times larger, and 0.10 and 0.20 are about 2.86 and 5.71 times larger. These noise levels were used for sensitivity analysis and were not calibrated to field measurement uncertainties. The decoder and the dense-log map are fit

once per step and seed. The penalty weight is then selected on validation for each measurement fraction and each noise level, so the curves are not a fixed-penalty reconstruction. The sensitivity sweep uses a different, wider grid. The weight selected at 0.01 and at 0.05 also belongs to the published grid. That membership does not mean the two searches were the same. At noise 0.20, some cells select 100, the largest candidate in the sweep grid. That choice is the best value among the candidates that were tried, not an optimum outside the grid. The points at noise 0.01 belong to this same refit: their ordering matches the published grid, and the cell means can differ slightly from that grid. The interpolation error grows with the noise because the interpolated curve follows the measurements. The CLP-CSGM error moves less and stays near the window predicted from the dense logs alone. At step 8 for $\rho = 0.05$ and $\rho = 0.10$, that starting window already has higher error than the interpolation at noise 0.05, and the CLP-CSGM mean stays between the two. The decoder can still represent the curve. That pattern does not isolate the dense-log window as the only cause. These runs are not merged into the submitted tables.

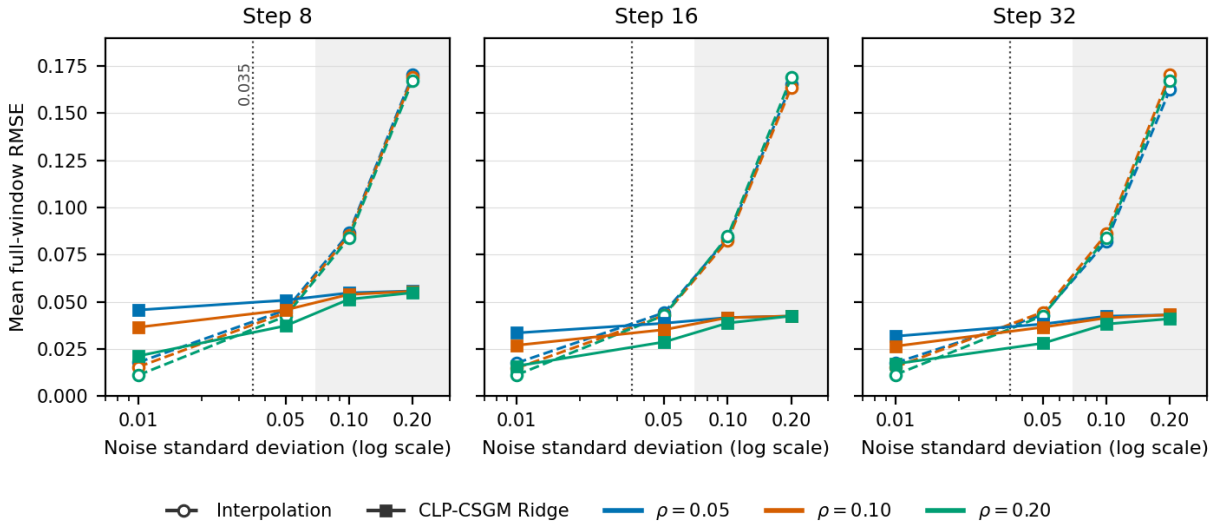


Figure 6: Mean full-window RMSE of depth interpolation and CLP-CSGM Ridge. The horizontal axis is logarithmic. Each panel is one step and each measurement fraction is a separate curve. Open circles are interpolation and squares are CLP-CSGM Ridge. Values are means over seeds $\{7, 23, 41\}$. The dotted line is the held-out V_c standard deviation 0.035, on the same scale as the noise. Noise 0.05 is about 1.43 times that deviation. The shaded region marks 0.10 and 0.20, about 2.86 and 5.71 times that deviation. These noise levels were used for sensitivity analysis and were not calibrated to field measurement uncertainties. The decoder and the dense-log map are fit once per step and seed. The penalty weight is selected on validation for each measurement fraction and each noise level. The points at noise 0.01 belong to this refit. Cell means and seed-level win counts are reported in Table 20.

Taken with the ten-seed panel, the two checks separate the randomness behind the three-seed intervals: more seeds, each with its own coordinate draw, and more coordinate draws at fixed model seeds. Both leave the submitted three-seed tables unchanged. Overlapping windows are not counted as extra independent samples.

Table 5: Paired seed-level $\Delta\text{RMSE} = \text{RMSE}_{\text{AE}} - \text{RMSE}_{\text{CLP}}$ on the cross-well V_c grid ($n = 3$ seeds). SE is the seed-level standard error; the 95% half-width uses Student-t with df = 2. Positive mean ΔRMSE favors CLP-CSGM Ridge. Wins count seeds with strictly lower RMSE for each method.

step	ρ	mean ΔRMSE	SE	t_2 95% half-width	CLP wins	AE wins
8	0.05	-0.00916	0.00350	0.01507	0	3
8	0.10	-0.00451	0.00231	0.00992	1	2
8	0.20	0.00334	0.00200	0.00860	2	1
16	0.05	0.00368	0.00123	0.00531	3	0
16	0.10	0.00730	0.00187	0.00803	3	0
16	0.20	0.01443	0.00362	0.01556	3	0
32	0.05	0.00684	0.00120	0.00515	3	0
32	0.10	0.01254	0.00216	0.00930	3	0
32	0.20	0.02254	0.00202	0.00869	3	0

Figure 7 places the pairwise result within the full benchmark. The mean RMSE of direct MLP and PCA increases sharply as training support decreases, whereas AE has the lowest unweighted mean RMSE among the direct baselines. Table 4 therefore focuses on the most competitive direct comparator in this grid.

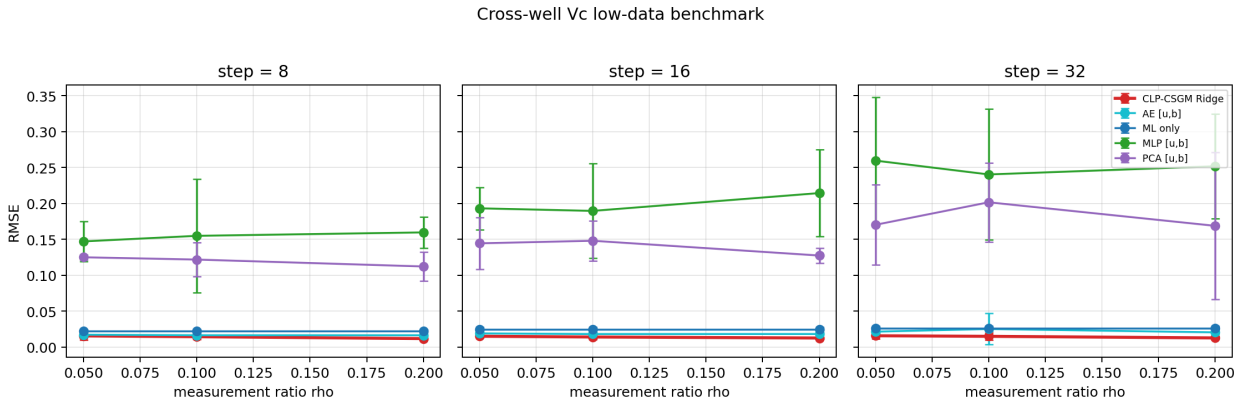


Figure 7: Cross-well V_c RMSE for the main methods. CLP-CSGM Ridge has the lowest mean RMSE in the medium-to-severe low-data cells, while the mean RMSE of direct MLP and PCA increases as training windows become scarce. Error bars are Student-t 95% half-widths across seeds (df = 2).

Figure 8 shows one held-out F03-4 profile for step 16, seed 7, and $\rho = 0.20$. In this displayed example, CLP-CSGM Ridge follows the broad variations of the processed reference more closely

than AE $[\boldsymbol{u}, \boldsymbol{b}]$. Direct MLP and PCA predictions exhibit larger oscillations. Their panels use wider horizontal scales than the reference, so that comparison is qualitative. The individual window in Fig. 15 is a different configuration: step 16, seed 7, $\rho = 0.10$, and window index 0. There, CLP-CSGM Ridge is closer to the reference, the autoencoder prediction is positively offset, and local deviations remain. These examples illustrate differences in reconstruction behavior within the displayed configurations. They do not establish a uniform advantage across the benchmark, and they are not a statistical-significance claim. The profile is consistent with the aggregate RMSE for this cell. Uniform window-level comparisons remain based on the aggregate metrics.

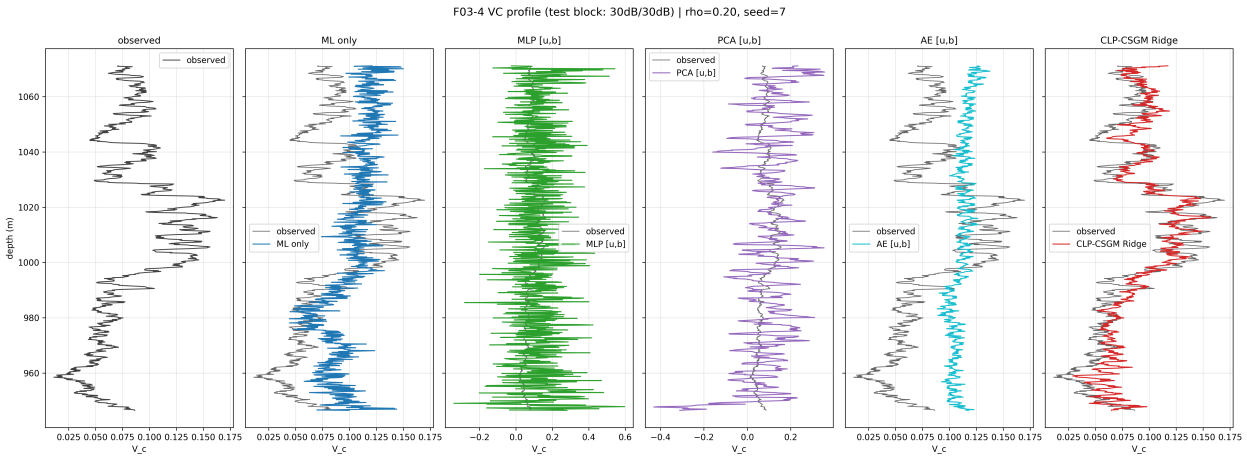


Figure 8: Held-out F03-4 V_c reconstruction at step 16, seed 7, and $\rho = 0.20$. The target is the processed Larionov-derived curve, and the sparse observations are coordinate subsamples of that curve. The MLP and PCA panels use wider horizontal scales than the reference panel.

The paired step-16 checks in Appendix B.7 (Table 14) show that including gamma ray reduces absolute error and that the embargo changes training support, but neither modification reverses the ordering of CLP-CSGM Ridge against AE $[\boldsymbol{u}, \boldsymbol{b}]$ and measurement-only CSGM at the tested ratios.

6.2 Single-well F03 GR-only results

The F03-4 benchmark uses gamma ray as the only dense log. Table 6 shows that CLP-CSGM Ridge achieves mean RMSE 0.01390 across the tested grid, compared with 0.01434 for AE $[\boldsymbol{u}, \boldsymbol{b}]$. CLP-CSGM Ridge has lower mean RMSE at $\rho \in \{0.20, 0.30\}$, the methods are nearly tied at $\rho = 0.50$, and AE has lower mean RMSE at $\rho \in \{0.40, 0.60\}$. In this benchmark, the largest advantage of explicit latent refinement occurs at the sparsest ratios. Direct AE fusion becomes competitive as more entries of $\boldsymbol{b}$ are available.

Table 6: Single-well F03-4 GR-only porosity benchmark: mean RMSE by measurement ratio ρ. The prior-only and measurement-only rows are structural reference variants.

method	$\rho = 0.20$	$\rho = 0.30$	$\rho = 0.40$	$\rho = 0.50$	$\rho = 0.60$
CLP-CSGM Ridge	0.01596	0.01475	0.01377	0.01281	0.01220
AE [u,b]	0.01850	0.01554	0.01304	0.01286	0.01178
Prior-only Ridge	0.01958	0.01958	0.01958	0.01958	0.01958
Measurement-only CSGM	0.02519	0.02285	0.02033	0.01622	0.01494
ML only	0.02359	0.02359	0.02359	0.02359	0.02359
PCA [u,b]	0.02149	0.02187	0.02261	0.02409	0.02426
MLP [u,b]	0.03022	0.02997	0.03712	0.04199	0.04627

Figure 9 displays the complete comparison with seed-level Student-t 95% half-widths (df = 2). Table 7 reports the paired $\Delta\text{RMSE} = \text{RMSE}_{\text{AE}} - \text{RMSE}_{\text{CLP}}$ on the same three seeds: the mean gap is positive at $\rho = 0.20$ and $\rho = 0.30$ (CLP wins all seeds), near zero at $\rho = 0.50$, and negative at $\rho \in \{0.40, 0.60\}$ (AE wins all seeds). With $n = 3$, several paired intervals include zero. The F03 ordering therefore changes with how many depths are measured, and the three-seed intervals are not a large-sample significance claim. RMSE variability, MAE, and relative ℓ_2 error are reported in Appendix B.2 (Table 13) and Appendix B.3 (Figs. 12–13); these metrics retain the same regime-dependent interpretation.

Table 7: Paired seed-level $\Delta\text{RMSE} = \text{RMSE}_{\text{AE}} - \text{RMSE}_{\text{CLP}}$ on the F03-4 GR-only porosity benchmark ($n = 3$ seeds). SE is the seed-level standard error; the 95% half-width uses Student-t with df = 2. Positive mean ΔRMSE favors CLP-CSGM Ridge.

ρ	mean ΔRMSE	SE	t_2 95% half-width	CLP wins	AE wins
0.20	0.00254	0.00042	0.00180	3	0
0.30	0.00079	0.00032	0.00139	3	0
0.40	-0.00073	0.00035	0.00151	0	3
0.50	0.00006	0.00033	0.00144	2	1
0.60	-0.00041	0.00024	0.00103	0	3

The stress tests in Appendix C (Table 22) extend the sparse grid to $\rho \in \{0.05, 0.10, 0.20\}$. CLP-CSGM Ridge retains lower mean RMSE than AE under increased noise and a smaller training pool, while the ordering becomes less stable with a coarse sliding step and fewer test windows.

Figure 10 shows the depth-profile reconstruction at $\rho = 0.20$, the main-grid ratio at which CLP-CSGM Ridge has its largest mean RMSE advantage over the direct AE baseline.

The porosity example is more mixed than the clay-volume profile. In this displayed profile, CLP-CSGM Ridge follows parts of the broad trend and smooths local variations. Because the profile combines predictions from overlapping windows, the view is qualitative. Pointwise recovery

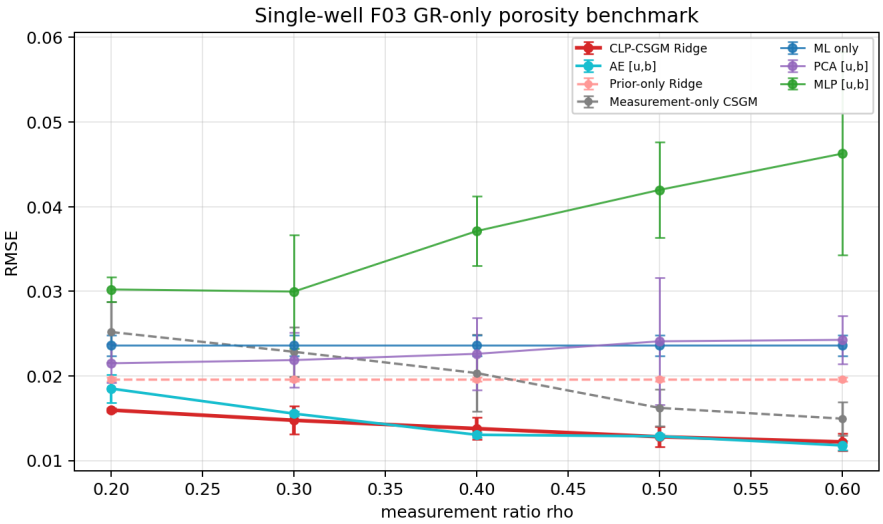


Figure 9: F03-4 GR-only porosity RMSE by measurement ratio. CLP-CSGM Ridge has the lowest mean RMSE at the sparsest ratios, while AE $[\boldsymbol{u}, \boldsymbol{b}]$ becomes competitive as more target observations are available. Error bars are Student-t 95% half-widths across seeds (df = 2).

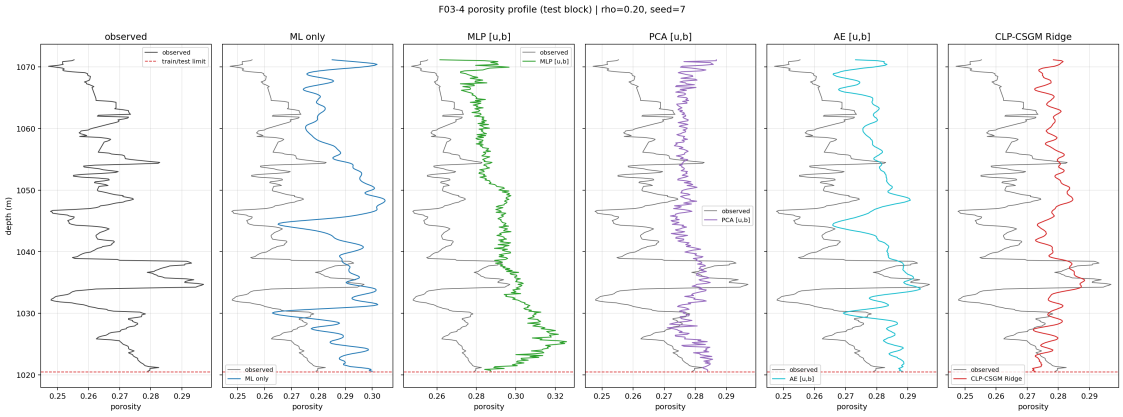


Figure 10: F03-4 depth-profile reconstruction at $\rho = 0.20$. The target is the processed density-derived porosity curve, and the sparse observations are coordinate subsamples of that curve. In this displayed profile, CLP-CSGM Ridge follows parts of the broad trend and smooths local variations. The panels do not share one horizontal scale.

and validation against independent core-porosity measurements fall outside this visualization. The displayed window in Fig. 14 uses the corresponding configuration seed 7, $\rho = 0.20$, and window index 0. There, both CLP-CSGM Ridge and AE $[\boldsymbol{u}, \boldsymbol{b}]$ miss the central depression of the reference. The pooled parity plot in Fig. 17 prints a correlation of 0.4632 for CLP-CSGM Ridge and 0.3720 for the autoencoder. Those values describe the points drawn on that figure. They do not establish statistical significance, and overlapping-window predictions are not independent observations. The method uses $m = 13$ observations per length-64 window and achieves a mean RMSE of approximately 0.016 at $\rho = 0.20$. This profile is consistent with the lower mean RMSE in Table 6 at that ratio. The profile is illustrative. The aggregate results still place the benefit of explicit latent refinement at the sparse end of this weak-context grid.

The limited-sample Lapa permeability case in Appendix D provides supplementary evidence. The prior-guided CLP-CSGM variants have lower mean RMSE than the direct baselines that use $[\boldsymbol{u}, \boldsymbol{b}]$ across the tested grid. Their comparison with measurement-only CSGM is ratio-dependent.

Appendix E scores laboratory porosity on held-out coordinates of one cored interval. CLP-CSGM is not superior on the mean: the direct AE has the lower RMSE at $\rho = 0.25$ and $\rho = 0.50$ (0.0537 and 0.0523, against 0.0602 and 0.0580 for CLP-CSGM Ridge). Laboratory measurements still correct the log-only anchor, which CLP-CSGM Ridge beats in every split. That correction is enough on the two shallower blocks, where the test logs remain close to the support that fits the anchor and CLP-CSGM Ridge has the lower block-mean RMSE. It does not close the gap on the deepest block, where the logs shift and the decoded prediction $G(\boldsymbol{z}_0)$ differs substantially from the target windows. The block RMSE, the log summaries, and the split counts are in Appendix E.

6.3 Ablation analysis

The structural ablations drop one term at a time. Prior-only returns the decoded prediction of the conditional prior, $G(h(\boldsymbol{u}))$, with no latent search. Measurement-only sets $\lambda = 0$ and searches the decoder range using only the sparse measurements [Bora et al., 2017]. The comparison shows what each term contributes (Remark 1).

In the cross-well benchmark, CLP-CSGM Ridge has lower mean RMSE than prior-only in all nine cells, with reductions from 20.5% to 60.6% (Table 8). It also has lower mean RMSE than measurement-only CSGM in seven cells. The exceptions are step 8 with $\rho \in \{0.05, 0.10\}$, where

Table 8: Cross-well V_c ablation results. Gains are relative RMSE reductions of CLP-CSGM Ridge with respect to the corresponding structural ablation; negative gain indicates that measurement-only CSGM has lower mean RMSE.

Step	ρ	CLP-CSGM	Prior-only	Meas.-only	AE $[u, b]$	Gain vs prior (%)	Gain vs meas. (%)
8	0.0500	0.0448	0.0564	0.0368	0.0357	20.4913	-21.7274
8	0.1000	0.0365	0.0564	0.0297	0.0320	35.3433	-22.7163
8	0.2000	0.0227	0.0564	0.0251	0.0260	59.8016	9.5581
16	0.0500	0.0334	0.0426	0.0362	0.0371	21.5310	7.6727
16	0.1000	0.0269	0.0426	0.0295	0.0342	36.8933	8.9585
16	0.2000	0.0177	0.0426	0.0213	0.0321	58.4391	16.7449
32	0.0500	0.0322	0.0432	0.0341	0.0390	25.5774	5.6452
32	0.1000	0.0264	0.0432	0.0271	0.0389	38.9248	2.4363
32	0.2000	0.0170	0.0432	0.0219	0.0396	60.5584	22.0073

Table 9: Single-well F03-4 GR-only ablation results. Gains are relative RMSE reductions of CLP-CSGM Ridge with respect to the corresponding structural ablation.

ρ	CLP-CSGM	Prior-only	Meas.-only	AE $[u, b]$	Gain vs prior (%)	Gain vs meas. (%)
0.2000	0.0160	0.0196	0.0252	0.0185	18.4646	36.6254
0.3000	0.0148	0.0196	0.0228	0.0155	24.6433	35.4375
0.4000	0.0138	0.0196	0.0203	0.0130	29.6514	32.2641
0.5000	0.0128	0.0196	0.0162	0.0129	34.5823	21.0300
0.6000	0.0122	0.0196	0.0149	0.0118	37.6999	18.3817

measurement-only CSGM has lower mean RMSE and the Ridge code $\boldsymbol{z}_0$ is not beneficial. At steps 16 and 32, the full model has lower mean RMSE than both ablations at every tested ratio.

In the F03-4 GR-only benchmark, CLP-CSGM Ridge has lower mean RMSE than both ablations at every tested ratio (Table 9). Its reduction relative to prior-only increases from approximately 18.5% at $\rho = 0.20$ to 37.7% at $\rho = 0.60$, while its largest gains over measurement-only occur at the sparsest ratios. Both terms contribute. The cross-well exceptions show that the conditional code $\boldsymbol{z}_0$ is not helpful in every cell.

6.4 Latent-prior model sensitivity

Tables 10 and 11 compare Ridge and MLP implementations of h while keeping the remaining pipeline fixed. This sensitivity analysis is limited to these two prior classes. Signed Ridge gain is defined as $100(\mathrm{RMSE}_{\mathrm{MLP}} - \mathrm{RMSE}_{\mathrm{Ridge}})/\mathrm{RMSE}_{\mathrm{MLP}}$; positive values favor Ridge.

In the cross-well benchmark, MLP has lower mean RMSE in all nine cells. The relative difference is largest at sparse measurement ratios and decreases as ρ increases; at step 16, for example, signed Ridge gain changes from approximately -82% at $\rho = 0.05$ to -26% at $\rho = 0.20$. Ridge remains

Table 10: Cross-well V_c latent-prior sensitivity. Positive signed Ridge gain favors Ridge, whereas negative values favor MLP.

Step	ρ	CLP-CSGM MLP	CLP-CSGM Ridge	Ridge gain (%)
8	0.0500	0.0183	0.0448	-144.8574
8	0.1000	0.0171	0.0365	-112.7948
8	0.2000	0.0138	0.0227	-64.3792
16	0.0500	0.0183	0.0334	-82.4082
16	0.1000	0.0175	0.0269	-53.3581
16	0.2000	0.0141	0.0177	-25.8526
32	0.0500	0.0221	0.0322	-45.4853
32	0.1000	0.0212	0.0264	-24.5143
32	0.2000	0.0158	0.0170	-7.7698

Table 11: Single-well F03-4 GR-only latent-prior sensitivity. Signed Ridge gain follows the definition used in Table 10.

ρ	CLP-CSGM MLP	CLP-CSGM Ridge	Ridge gain (%)
0.2000	0.0186	0.0160	14.2682
0.3000	0.0167	0.0148	11.4868
0.4000	0.0152	0.0138	9.2929
0.5000	0.0132	0.0128	2.6767
0.6000	0.0124	0.0122	1.6615

the headline comparator for the reason in Section 7: a fixed linear reference, not a claim that it dominates MLP on this grid.

In the F03-4 GR-only benchmark, Ridge has lower mean RMSE at every tested ratio, but the relative difference decreases from approximately 14.3% at $\rho = 0.20$ to 1.7% at $\rho = 0.60$.

Which map is better depends on the dataset and on how many depths are measured, and the difference shrinks as more target values enter $\boldsymbol{b}$. Ridge stays the reported reference because it is a fixed linear map, not because it wins every cell. The rankings support choosing that map on validation data for the intended test, rather than after seeing the held-out result.

Supplementary diagnostics. Appendix B.6 audits validation, selected λ, and per-window errors; Appendix B.4 shows realized sparse-measurement windows; Appendix B.1 assesses empirical decoder-range separation; and Appendix B.13 reports inference cost. These analyses serve interpretive and reproducibility purposes and are separate from the benchmark comparisons.

7 Interpretation, scope, and limitations

CLP-CSGM uses three pieces that act together. The decoder range $\mathrm{Im}(G_{\mathrm{orig}})$ is the set of target windows, in the original units, that the decoder can produce. The conditional map predicts $\boldsymbol{z}_0 = h(\boldsymbol{u})$ from the dense logs. This code initializes the first latent trajectory and remains the center of the quadratic penalty. The measurement term in Eq. (10) penalizes disagreement with the sparse observations in the original target units and does not force those observations to be met exactly. The decoder range limits which windows can appear. The code $\boldsymbol{z}_0$ carries the dense logs into that search, and the measurement term uses the sparse target values. Unlike direct $[\boldsymbol{u}, \boldsymbol{b}]$ predictors, CLP-CSGM retains $\boldsymbol{b}$ as an explicit measurement penalty during inference. It is therefore most relevant when dense logs are informative but incomplete and sparse target entries should enter that penalty. Log-completion models fill a missing interval of a wireline curve, geomodel translation fills a property image from another image, and direct regression maps logs to porosity or shale volume [Koeshidayatullah et al., 2026, Al-Fakih et al., 2026, 2025, 2023]. CLP-CSGM is the relevant formulation when those dense logs are already available and a few observations of the target property enter the reconstructed window through that penalty.

The depth profiles in Figs. 8 and 10 illustrate how the balance of these information sources affects reconstruction detail in the displayed examples. The multi-channel cross-well context supports closer tracking of the broad held-out V_c variations, whereas the deliberately weak GR-only context produces a smoother F03-4 porosity reconstruction that omits some local variations. The displayed porosity window shows that both CLP-CSGM and the direct autoencoder can miss a local depression. These examples do not establish a uniform advantage across the benchmark. The attainable detail is governed jointly by the information in $\boldsymbol{u}$, the adequacy of $\mathrm{Im}(G_{\mathrm{orig}})$, and the number and placement of observations in $\boldsymbol{b}$.

The ablations in Section 6.3 drop one piece at a time. Using only the decoded prediction of the conditional prior, or only the measurements, raises the error in most cells, and which piece matters more changes with the number of training windows. The full Ridge model has lower mean RMSE than prior-only throughout the cross-well grid and lower mean RMSE than measurement-only CSGM in seven of nine cells. In the two step-8 exceptions, measurement-only inversion has lower mean RMSE, so the conditional code $\boldsymbol{z}_0$ is not helpful in those cells. On the F03-4 grid, where

gamma ray is the only dense log, the full model has lower mean RMSE than both ablations at every tested ratio. The penalty around $\boldsymbol{z}_0$ keeps the search from drifting, and the measurement term can still correct it.

The sensitivity results in Section 6.4 show that no single architecture for h is uniformly preferable. MLP has lower mean RMSE in all nine multi-channel cross-well cells, whereas Ridge has lower mean RMSE at every tested ratio under the single-channel F03-4 setting. In both cases, the difference decreases as ρ increases, which is consistent with less dependence on $\boldsymbol{z}_0$ as more target observations become available. These comparisons distinguish the CLP-CSGM framework from its reference implementation. Ridge remains the primary reported configuration because it is a fixed linear map: the main tables then use one h in every cell, and h is not chosen after seeing the test ranking. The MLP results have the same experimental status. They show that a more flexible map improves the cross-well grid and that the preferred class changes with the test. The class of h, the latent dimension, and the regularization should be selected on validation data for the intended measurements, not by adopting whichever variant wins on the held-out test set. The experiments establish this sensitivity but do not identify which architectural property causes the opposite rankings.

The empirical scope is limited by the construction of the benchmark targets. In the cross-well experiment, V_c is a processed Larionov-derived curve. The gamma-ray index between the clean and shale endpoints is linear in gamma ray. The Larionov map from that index to clay volume is not a straight line. If gamma ray has been logged, clay volume is obtained from that relation, and CLP-CSGM is not a substitute for it. Gamma ray is left out of the main input because the clay-volume curve was computed from it. Without that log the Larionov relation cannot be evaluated. The task is then to reconstruct the processed curve from the other dense logs while matching a few known values of the curve. Recovery of independently acquired clay-volume measurements is outside this scope. Putting gamma ray back, in Appendix B.7, lowers the absolute error and does not change the ordering of CLP-CSGM Ridge, the direct AE, and measurement-only CSGM at step 16. That run checks the channel. It is not the log suite used for the main comparison.

The F03-4 porosity curve is density-derived, and its sparse values are taken from that same curve, not from core or laboratory plugs. Density porosity is linear in bulk density, and a standard sonic porosity is linear in slowness. If those logs have been recorded, the processed porosity is

already determined. Sonic travel time is left out because it follows this curve almost one to one, and the density curve behind the porosity was itself predicted from sonic and gamma ray. The test uses gamma ray as the only dense log. Its errors do not transfer to a porosity interpretation that also uses density, neutron, and sonic. The two F3 benchmarks therefore remove the log the conventional formula needs and compare methods that must reconstruct the processed curve from the remaining logs while matching a few known values. They are not a contest against Larionov, density porosity, or sonic porosity on a fully logged well, and they are not a test against independent core or laboratory measurements.

The Lapa permeability case [Alvarez et al., 2025] extends the evaluation to another lithology and target property, but its small sample size limits generalization. The laboratory-porosity supplement in Appendix E uses $\mathrm{Phi}_{\mathrm{lab}}$ rather than a curve computed from the wireline suite, and it scores only coordinates that were not supplied as measurements. The direct AE has the lower pooled mean RMSE. Laboratory coordinates still correct the decoded prediction of the conditional prior in every split, and that correction is enough where the test logs remain close to the blocks that fit $h(\boldsymbol{u})$. On the deepest block the logs move outside that support, the decoded prediction $G(\boldsymbol{z}_0)$ differs substantially from the target windows, and a few measurements improve on the decoded prediction without reaching the direct AE. The measurement term still lowers the error relative to the decoded prediction of the conditional prior in that block. The supplement therefore does not extend the F3 low-support advantage to a mean win on this laboratory target. Each test fold still has many overlapping training windows, so the direct map is in the regime where it was expected to be competitive. Block RMSE and the log summaries are reported in the appendix. Broader validation requires larger multi-well datasets, additional log suites, and facies-conditioned settings.

The reported metrics are not stratified by lithofacies, lithological texture, or formal stratigraphic interval because the available benchmark tables do not provide independent geological labels at each depth sample. Consequently, RMSE, MAE, and relative ℓ_2 error aggregate heterogeneous windows. The depth profiles provide qualitative audits but no geological performance breakdowns. Facies-resolved reliability would require independently interpreted or core-calibrated interval labels, particularly because the adequacy of the decoder range and the accuracy of the latent prior may vary across geological regimes.

The recovery interpretation is also conditional. Proposition 3 relates reconstruction error to

decoder mismatch, latent-prior mismatch, measurement noise, optimization quality, and local separability, but does not prove that M_Ω satisfies S-REC for the trained decoder. The diagnostics in Appendix B.1 evaluate sampled decoder outputs and should not be interpreted as an operator-level guarantee.

CLP-CSGM incurs iterative inference cost, whereas direct baselines produce predictions through a forward pass. The runtime results in Appendix B.13 show that this additional cost is most justified when explicit measurement-consistent refinement produces a meaningful accuracy gain. Faster latent solvers, warm starts, adaptive stopping, or amortized refinement are natural directions for larger-scale deployment.

8 Conclusion

We introduced CLP-CSGM. A decoder trained on windows of the target curve supplies the windows the reconstruction may use. The conditional map predicts a latent code $\boldsymbol{z}_0$ from the dense auxiliary logs. This code provides an informed initialization and serves as the reference for the quadratic penalty during latent refinement. Latent refinement balances agreement with the sparse observations against deviation from $\boldsymbol{z}_0$. A residual at the observed depths can remain. The MAP/Tikhonov reading and the local error bound separate measurement noise, decoder error, a code $\boldsymbol{z}_0$ far from a good code, and whether nearby windows can be told apart from the measurements.

Under the paired F3 protocols, CLP-CSGM Ridge has lower mean RMSE than AE $[\boldsymbol{u}, \boldsymbol{b}]$ in seven of nine cross-well cells. On the published grid that autoencoder has the lowest unweighted mean RMSE among the direct baselines MLP, PCA, and AE $[\boldsymbol{u}, \boldsymbol{b}]$. The seven-of-nine result is a comparison with this autoencoder, not a ranking against interpolation or the direct linear maps. This is a conditional advantage where training support is smallest (steps 16 and 32), not overall superiority: AE $[\boldsymbol{u}, \boldsymbol{b}]$ has lower mean RMSE at step 8 for $\rho = 0.05$ and $\rho = 0.10$. The unweighted mean of the nine cell-level mean RMSE values decreases from 0.03496 to 0.02862, a reduction of approximately 18.1%. In the F03-4 GR-only benchmark, CLP-CSGM Ridge has the lowest mean RMSE at $\rho \in \{0.20, 0.30\}$, while AE $[\boldsymbol{u}, \boldsymbol{b}]$ has lower mean RMSE at $\rho \in \{0.40, 0.60\}$. That benchmark assimilates the processed porosity curve with gamma ray as the only dense input; the depth profile at $\rho = 0.20$ follows the broad trend of this curve. It is a controlled stress test under

intentionally incomplete context and is not evidence of general porosity reconstruction performance. The full Ridge model has lower mean RMSE than prior-only in every cross-well cell and than measurement-only CSGM in seven of nine cells. It also has lower mean RMSE than both ablations throughout the F03-4 grid. The opposite Ridge–MLP rankings show that the better map from dense logs to $\boldsymbol{z}_0$ depends on the test.

At the reference noise standard deviation 0.01, depth interpolation of the sparse measurements has lower full-window RMSE than CLP-CSGM Ridge in all nine configurations and in all 27 seed comparisons. In that regime the sparse measurements already support a competitive simple solution, which limits the need for a learned reconstruction. The same windows were then rescored at noise standard deviations 0.05, 0.10, and 0.20, with the decoder and the dense-log map held fixed and the penalty reselected at each level. At 0.05, about 1.43 times the held-out curve standard deviation 0.035, CLP-CSGM Ridge has the lower mean RMSE at steps 16 and 32 and at step 8 for $\rho = 0.20$. At step 8 for $\rho = 0.05$ and $\rho = 0.10$, interpolation keeps the lower mean. At 0.10 and 0.20, about 2.86 and 5.71 times that standard deviation, CLP-CSGM Ridge has the lower error in all nine configurations. These noise levels were used for sensitivity analysis and were not calibrated to field measurement uncertainties. AE $[\boldsymbol{u}, \boldsymbol{b}]$, ordinary least squares, and direct ridge regression were not rerun across this sweep, and no smoothing comparator was added, so the change is not a general advantage under noise.

The study presents a latent reconstruction that combines dense-log context with sparse target observations. The benefit depends on the regime. The F3 evidence is reconstruction of processed target logs: cross-well V_c is Larionov-derived from gamma ray, F03-4 porosity is density-derived, and sparse observations are coordinate subsamples of these same curves rather than independently acquired core or laboratory measurements. A separate single-well supplement on laboratory porosity scores held-out laboratory coordinates and is not part of that processed-log grid. Relative to the decoded prediction $G(\boldsymbol{z}_0)$, CLP-CSGM Ridge has lower RMSE in every split. It also has the lower block-mean RMSE on the two shallower blocks. On the pooled mean, the direct AE is lower. The larger share of that error sits in the deepest block. The log differences observed there are compatible with a transfer difficulty. They do not establish that shift as the only cause. That supplement is one cored interval. The cross-well comparison is a controlled single-holdout benchmark on F03-4 and is not evidence of operational cross-well or multi-field transferability. Broader assessment will

require larger multi-well datasets, independently calibrated targets, and geological labels. It will also require uncertainty estimates and faster latent-refinement strategies.

Author Contributions

Romulo Brito da Silva: Conceptualization, Methodology, Software, Formal Analysis, Data Curation, Validation, Visualization, Writing—Original Draft, Writing—Review & Editing.

Diego Passos: Visualization, including figure organization and standardization, Writing—Review & Editing.

Irineu de Azevedo Lima Neto: Visualization, including figure organization and standardization, Writing—Review & Editing.

Herson Oliveira da Rocha: Visualization, including figure organization and standardization, Writing—Review & Editing.

Roseane Marchezi Missagia: Senior Review, Advisory Support, Writing—Review & Editing.

Marco Antonio Rodrigues de Ceia: Senior Review, Advisory Support, Writing—Review & Editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the authors used OpenAI's ChatGPT to assist with language refinement, improve the clarity and organization of the manuscript, reduce its length to meet the journal's word-count requirements, and assist in the creation of explanatory schematic figures illustrating concepts, processes, and experimental workflows. The tool was not used to generate research data, perform the reported numerical experiments, fabricate or alter quantitative results, or determine the scientific conclusions. After using this tool, the authors reviewed, edited, and validated

the text and figures as needed and take full responsibility for the content of the published article.

References

Reda Abdel Azim and Abdulrahman Aljehani. Neural network model for permeability prediction from reservoir well logs. *Processes*, 10(12):2587, 2022. doi: 10.3390/pr10122587.

Jonas Adler and Ozan Öktem. Deep bayesian inversion, 2018. URL `https://arxiv.org/abs/1811.05910`.

Mohammad Ali Ahmadi and Zhangxin Chen. Comparison of machine learning methods for estimating permeability and porosity of oil reservoirs via petro-physical logs. *Petroleum*, 5(3):271–284, 2019.

Abdulrahman Al-Fakih, SanLinn I. Kaka, and Ardiansyah I. Koeshidayatullah. Reservoir property prediction in the north sea using machine learning. *IEEE Access*, 11:140148–140160, 2023. doi: 10.1109/ACCESS.2023.3336623.

Abdulrahman Al-Fakih, A. Koeshidayatullah, Tapan Mukerji, Sadam Al-Azani, and SanLinn I. Kaka. Well log data generation and imputation using sequence based generative adversarial networks. *Scientific Reports*, 15:11000, 2025. doi: 10.1038/s41598-025-95709-0.

Abdulrahman Al-Fakih, Ardiansyah Koeshidayatullah, Nabil A. Saraih, Tapan Mukerji, Rayan Kanfar, Abdulmohsen Alali, and SanLinn I. Kaka. Pix2geomodel: A next-generation reservoir geomodeling with property-to-property translation. *Geoenergy Science and Engineering*, 258: 214342, 2026. doi: 10.1016/j.geoen.2025.214342.

Fabian Altekrüger, Paul Hagemann, and Gabriele Steidl. Conditional generative models are provably robust: Pointwise guarantees for bayesian inverse problems. *Transactions on Machine Learning Research*, 2024. ISSN 2835-8856.

Auddys Alvarez, Roseane Missagia, Romulo Brito da Silva, and Abel Carrasquilla. Non-linear regression for permeability prediction and hydraulic flow units' characterization from core and well logs of a pre-salt carbonate reservoir in the santos basin, brazil. *Journal of Applied Geophysics*, 241:105865, 2025. doi: 10.1016/j.jappgeo.2025.105865.

Muhammad Asim, Mara Daniels, Oscar Leong, Ali Ahmed, and Paul Hand. Invertible generative models for inverse problems: Mitigating representation error and dataset bias. In *Proceedings of*

the 37th International Conference on Machine Learning, volume 119 of *Proceedings of Machine Learning Research*, pages 399–409. PMLR, 2020.

George Asquith and Daniel Krygowski. *Basic Well Log Analysis*. American Association of Petroleum Geologists, 2 edition, 2004.

Matthew C. Bendel, Rizwan Ahmad, and Philip Schniter. A regularized conditional gan for posterior sampling in image recovery problems. In *Advances in Neural Information Processing Systems*, volume 36, 2024.

Christopher M. Bishop. *Pattern Recognition and Machine Learning*. Springer, 2006.

Ashish Bora, Ajil Jalal, Eric Price, and Alexandros G. Dimakis. Compressed sensing using generative models. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 537–546. PMLR, 2017.

Martin Burger and Felix Lucka. Maximum a posteriori estimates in linear inverse problems with log-concave priors are proper Bayes estimators. *Inverse Problems*, 30(11):114004, 2014. doi: 10.1088/0266-5611/30/11/114004.

Emmanuel J. Candes, Justin Romberg, and Terence Tao. Robust uncertainty principles: Exact signal reconstruction from highly incomplete frequency information. *IEEE Transactions on Information Theory*, 52(2):489–509, 2006a.

Emmanuel J. Candes, Justin K. Romberg, and Terence Tao. Stable signal recovery from incomplete and inaccurate measurements. *Communications on Pure and Applied Mathematics*, 59(8):1207–1223, 2006b.

Romulo B. da Silva, Diego Passos, Cássio M. Oishi, and J. Nathan Kutz. CS-SHRED: Enhancing SHRED for robust recovery of spatiotemporal dynamics, 2025. URL `https://arxiv.org/abs/2507.22303`.

Agnimitra Dasgupta, Harisankar Ramaswamy, Javier Murgoitio Esandi, Ken Foo, Arka Chakraborty, Phaedon-Stelios Koutsourelakis, Dennis M. Kochmann, and Yunzhu Yang. Conditional score-based diffusion models for solving inverse problems in mechanics, 2024. URL `https://arxiv.org/abs/2406.13154`.

Pedro Carlos de Assis, Romulo B. da Silva, Irineu Lima Neto, Alan C. Galante, Alexandre L. Campi, Alessandro R. de Oliveira, Roseane Marchezi Missagia, and Marco A. R. Ceia. Compressive sensing framework using the linear Radon transform for 3D ultrasonic data reconstruction in a pinch-out reservoir model. In *Second International Meeting for Applied Geoscience & Energy*, pages 2621–2625. Society of Exploration Geophysicists and American Association of Petroleum Geologists, 2022. doi: 10.1190/image2022-3749310.1.

David L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*, 52(4):1289–1306, 2006.

Darwin V. Ellis and Julian M. Singer. *Well Logging for Earth Scientists*. Springer, 2 edition, 2007.

Yawen He, Hongjun Zhang, Zhiyu Wu, Hongbo Zhang, Xin Zhang, Xiaojing Zhuo, Xiaoli Song, Sha Dai, and Wei Dang. Porosity prediction of tight reservoir rock using well logging data and machine learning. *Scientific Reports*, 15:13124, 2025. doi: 10.1038/s41598-025-95578-7.

Gilles Hennenfent and Felix J. Herrmann. Simply denoise: Wavefield reconstruction via jittered undersampling. *Geophysics*, 73(3):V19–V28, 2008. doi: 10.1190/1.2841038.

Felix J. Herrmann and Gilles Hennenfent. Non-parametric seismic data recovery with curvelet frames. *Geophysical Journal International*, 173(1):233–248, 2008. doi: 10.1111/j.1365-246X.2007.03698.x.

Arthur E. Hoerl and Robert W. Kennard. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970.

Christian Jacobsen, Yilin Zhuang, and Karthik Duraisamy. Cocogen: Physically consistent and conditioned score-based generative models for forward and inverse problems. *SIAM Journal on Scientific Computing*, 47(2):C399–C425, 2025. doi: 10.1137/24M1636071.

Akshay Kamath, Eric Price, and Jonathan Scarlett. On the power of compressed sensing with generative models. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 5101–5109. PMLR, 2020.

Kyung-Su Kim, Jung Hyun Lee, and Eunho Yang. Compressed sensing via measurement-conditional generative models. *IEEE Access*, 9:155335–155352, 2021. doi: 10.1109/ACCESS.2021.3128721.

Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization, 2015. URL https://arxiv.org/abs/1412.6980.

Ardiansyah Koeshidayatullah, Abdulrahman Al-Fakih, and SanLinn Ismael Kaka. Toward basin-agnostic well log imputation and anomaly detection via a pre-trained time-series foundation model. *Energy Geoscience*, 7(2):100536, 2026. doi: 10.1016/j.engeos.2026.100536.

Abdul-Muaizz Koray, Dung Bui, Emmanuel Appiah Kubi, William Ampomah, et al. Machine learning based reservoir characterization and numerical modeling from integrated well log and core data. *Geoenergy Science and Engineering*, 243:213296, 2024. doi: 10.1016/j.geoen.2024.213296.

Alice Lucas, Michael Iliadis, Rafael Molina, and Aggelos K. Katsaggelos. Using deep neural networks for inverse problems in imaging: Beyond analytical methods. *IEEE Signal Processing Magazine*, 35(1):20–36, 2018.

Uti Ikitsombika Markus, Jing Ba, Muhammad Abid, Faruwa Ajibola Richard, and Emo Obadiah. Hybrid optimization of interpretable ensemble machine learning for petrophysical property prediction from well logs. *Frontiers in Earth Science*, 13:1721227, 2026. doi: 10.3389/feart.2025.1721227.

Morteza Matinkia, Romina Hashami, Mohammad Mehrad, Mohammad Reza Hajsaeedi, and Arian Velayati. Prediction of permeability from well logs using a new hybrid machine learning algorithm. *Petroleum*, 9(1):108–123, 2023. doi: 10.1016/j.petlm.2022.03.003.

Mehdi Mirza and Simon Osindero. Conditional generative adversarial nets, 2014.

Ilius Mondal and Kumar Hemant Singh. Core-log integration and application of machine learning technique for better reservoir characterisation of eocene carbonates, indian offshore. *Energy Geoscience*, 3(1):49–62, 2022. doi: 10.1016/j.engeos.2021.10.006.

Jack B. Muir and Zhongwen Zhan. Seismic wavefield reconstruction using a pre-conditioned wavelet–curvelet compressive sensing approach. *Geophysical Journal International*, 227(1):303–315, 2021. doi: 10.1093/gji/ggab222.

Greg Ongie, Ajil Jalal, Christopher A. Metzler, Richard G. Baraniuk, Alexandros G. Dimakis, and Rebecca Willett. Deep learning techniques for inverse problems in imaging. *IEEE Journal on Selected Areas in Information Theory*, 1(1):39–56, 2020.

Wen Pan, Carlos Torres-Verdín, Ian J. Duncan, and Michael J. Pyrcz. Improving multiwell petrophysical interpretation from well logs via machine learning and statistical models. *Geophysics*, 88 (2), 2023. doi: 10.1190/GEO2022-0151.1.

Pengyu Shi, Pengda Shi, Kang Bie, Chuang Han, Xiaowei Ni, Zhiqiang Mao, and Peiqiang Zhao. Prediction of permeability in a tight sandstone reservoir using a gated network stacking model driven by data and physical models. *Frontiers in Earth Science*, 12:1364515, 2024. doi: 10.3389/feart.2024.1364515.

Kihyuk Sohn, Honglak Lee, and Xinchen Yan. Learning structured output representation using deep conditional generative models. In *Advances in Neural Information Processing Systems*, volume 28, 2015.

Weizhi Sun, Junxing Zhao, Shuaijie Peng, Wenlu Huang, and Hongyi Li. Physically consistent joint prediction of porosity and shale volume via core-calibrated deep learning in well-consolidated sandstones. *Scientific Reports*, 16:1737, 2026. doi: 10.1038/s41598-025-31326-1.

Andrey N. Tikhonov. Solution of incorrectly formulated problems and the regularization method. *Soviet Mathematics Doklady*, 4:1035–1038, 1963.

Zihao Wang, Yidong Cai, Dameng Liu, Jun Lu, Feng Qiu, Jinghong Hu, Zhentao Li, and Ranjith Pathegama Gamage. A review of machine learning applications to geophysical logging inversion of unconventional gas reservoir parameters. *Earth-Science Reviews*, 258:104969, 2024. doi: 10.1016/j.earscirev.2024.104969.

David A. Wood. Predicting porosity, permeability and water saturation applying an optimized nearest-neighbour, machine-learning and data-mining network of well-log data. *Journal of Petroleum Science and Engineering*, 184:106587, 2020.

Ji Zhang, Guiping Liu, Zhen Wei, Shengge Li, Yeheya Zayier, and Yuanfeng Cheng. Machine learning-based prediction of well logs guided by rock physics and its interpretation. *Sensors*, 25 (3):836, 2025. doi: 10.3390/s25030836.

A Supplementary Theoretical Details

This section collects the proofs omitted from the main text, including the detailed conditional local-recovery argument. That argument separates four contributions to the error: how far the true window is from anything the decoder can produce, how far $\boldsymbol{z}_0$ is from a good code, measurement noise, and whether the measurements can tell nearby windows apart.

A.1 MAP and Tikhonov interpretation

Proof of Proposition 1. The MAP estimate is obtained by maximizing the posterior density

$$p(z \mid b, u).$$

By Bayes' rule,

$$p(z \mid b, u) = \frac{p(b \mid z, u)p(z \mid u)}{p(b \mid u)}.$$

Using the conditional independence assumption $p(b \mid z, u) = p(b \mid z)$, and noting that $p(b \mid u)$ does not depend on z, we obtain

$$p(z \mid b, u) \propto p(b \mid z)p(z \mid u).$$

Therefore,

$$\hat{z}_{\mathrm{MAP}} \in \arg\min_z \left[-\log p(b \mid z) - \log p(z \mid u) \right].$$

Since

$$b \mid z \sim \mathcal{N}\big(M_\Omega(\mu_y + D_y G(z)), \sigma_\eta^2 I_m\big),$$

we have, up to constants independent of z,

$$-\log p(b \mid z) = \frac{1}{2\sigma_\eta^2} \|M_\Omega(\mu_y + D_y G(z)) - b\|_2^2 + \text{const}.$$

Similarly, since

$$z \mid u \sim \mathcal{N}(z_0(u), \sigma_z^2 I_k),$$

we have, up to constants independent of z,

$$-\log p(z \mid u) = \frac{1}{2\sigma_z^2}\|z - z_0(u)\|_2^2 + \text{const.}$$

Hence the MAP estimate minimizes

$$\frac{1}{2\sigma_\eta^2}\|M_\Omega(\mu_y + D_y G(z)) - b\|_2^2 + \frac{1}{2\sigma_z^2}\|z - z_0(u)\|_2^2.$$

Multiplying this objective by the positive constant $2\sigma_\eta^2$ leaves its minimizers unchanged. Thus,

$$\hat{z}_{\text{MAP}} \in \arg\min_z \|M_\Omega(\mu_y + D_y G(z)) - b\|_2^2 + \frac{\sigma_\eta^2}{\sigma_z^2}\|z - z_0(u)\|_2^2.$$

By the identity in Eq. (11), the data term is the weighted standardized residual minimized in Eq. (10). Therefore the CLP-CSGM objective is recovered on the original measurement scale with

$$\lambda = \frac{\sigma_\eta^2}{\sigma_z^2}.$$

The second term is a Tikhonov regularization term centered at the log-conditioned latent anchor $z_0(u)$, not at the origin. □

A.2 Existence of a minimizer

Proof of Proposition 2. Since G is continuous, M_Ω is linear, and D_Ω does not depend on z, the map

$$z \mapsto \|D_\Omega(M_\Omega G(z) - \tilde{b})\|_2^2$$

is continuous. The quadratic term

$$z \mapsto \lambda\|z - z_0\|_2^2$$

is also continuous. Hence

$$J(z) = \|D_\Omega(M_\Omega G(z) - \tilde{b})\|_2^2 + \lambda\|z - z_0\|_2^2$$

is continuous on $\mathbb{R}^k$.

Moreover, since $\lambda > 0$,

$$J(z) \geq \lambda \|z - z_0\|_2^2.$$

As $\|z\|_2 \to \infty$, we also have $\|z - z_0\|_2 \to \infty$. Therefore,

$$J(z) \to +\infty \qquad \text{as} \qquad \|z\|_2 \to \infty.$$

Thus J is coercive.

Choose the reference level

$$c = J(z_0),$$

and define the sublevel set

$$S_c = \{z \in \mathbb{R}^k : J(z) \leq c\}.$$

The set S_c is nonempty because $z_0 \in S_c$. It is closed since

$$S_c = J^{-1}((-\infty, c]),$$

J is continuous, and $(-\infty, c]$ is closed in $\mathbb{R}$.

Since J is coercive, there exists $R > 0$ such that

$$\|z\|_2 > R \quad \Longrightarrow \quad J(z) > c.$$

Therefore no point with $\|z\|_2 > R$ belongs to S_c, and hence

$$S_c \subset \overline{B}_R = \{z \in \mathbb{R}^k : \|z\|_2 \leq R\}.$$

Thus S_c is bounded. Since S_c is closed and bounded in the finite-dimensional space $\mathbb{R}^k$, it is compact.

By the Weierstrass theorem, the continuous function J attains its minimum on S_c. Hence there exists $\hat{z} \in S_c$ such that

$$J(\hat{z}) = \min_{z \in S_c} J(z).$$

It remains only to note that this minimum is also global. Indeed, if $z \in S_c$, then

$$J(\hat{z}) \leq J(z)$$

by the definition of $\hat{z}$. If $z \notin S_c$, then $J(z) > c$, while $\hat{z} \in S_c$ implies

$$J(\hat{z}) \leq c.$$

Therefore,

$$J(\hat{z}) \leq c < J(z)$$

for every $z \notin S_c$. Hence no point outside S_c can have objective value smaller than $J(\hat{z})$. Consequently,

$$J(\hat{z}) = \min_{z \in \mathbb{R}^k} J(z),$$

and J admits at least one global minimizer in $\mathbb{R}^k$. □

A.3 Conditional local recovery interpretation

Proposition 3 gives the conditional local recovery interpretation summarized in the main text. The full proof is followed by the corresponding extension for a numerical optimizer.

Proof of Proposition 3. The residual $\|M_\Omega G_{\text{orig}}(z) - b\|_2^2$ is the right-hand side of Eq. (11), so it is the data term of $\mathcal{J}$. By global optimality of $\hat{z}$,

$$\mathcal{J}(\hat{z}) \leq \mathcal{J}(z^\dagger).$$

Since $y^\star = G_{\text{orig}}(z^\dagger) + r^\dagger$, we have

$$b = M_\Omega y^\star + \eta = M_\Omega G_{\text{orig}}(z^\dagger) + M_\Omega r^\dagger + \eta = M_\Omega G_{\text{orig}}(z^\dagger) + e^\dagger.$$

Therefore,

$$\mathcal{J}(z^\dagger) = \|M_\Omega G_{\text{orig}}(z^\dagger) - b\|_2^2 + \lambda \|z^\dagger - z_0\|_2^2 = \|e^\dagger\|_2^2 + \lambda \|z^\dagger - z_0\|_2^2 = (\Delta^\dagger)^2.$$

Hence $\mathcal{J}(\hat{z}) \le (\Delta^\dagger)^2$. Since each term in $\mathcal{J}(\hat{z})$ is nonnegative,

$$\lambda \|\hat{z} - z_0\|_2^2 \le \mathcal{J}(\hat{z}) \le (\Delta^\dagger)^2,$$

which yields

$$\|\hat{z} - z_0\|_2 \le \frac{\Delta^\dagger}{\sqrt{\lambda}}.$$

Also, by the definition of $\Delta^\dagger$,

$$\Delta^\dagger \ge \sqrt{\lambda}\|z^\dagger - z_0\|_2,$$

so

$$\|z^\dagger - z_0\|_2 \le \frac{\Delta^\dagger}{\sqrt{\lambda}}.$$

Thus both $\hat{z}$ and $z^\dagger$ belong to $B(z_0, R^\dagger)$.

The local S-REC on $\mathcal{M}_{R^\dagger}$ therefore applies to $G_{\text{orig}}(\hat{z})$ and $G_{\text{orig}}(z^\dagger)$, giving

$$\gamma_R \|G_{\text{orig}}(\hat{z}) - G_{\text{orig}}(z^\dagger)\|_2 - \delta_R \le \|M_\Omega(G_{\text{orig}}(\hat{z}) - G_{\text{orig}}(z^\dagger))\|_2.$$

By the triangle inequality,

$$\|M_\Omega(G_{\text{orig}}(\hat{z}) - G_{\text{orig}}(z^\dagger))\|_2 \le \|M_\Omega G_{\text{orig}}(\hat{z}) - b\|_2 + \|b - M_\Omega G_{\text{orig}}(z^\dagger)\|_2.$$

The first term satisfies

$$\|M_\Omega G_{\text{orig}}(\hat{z}) - b\|_2^2 \le \mathcal{J}(\hat{z}) \le (\Delta^\dagger)^2,$$

so $\|M_\Omega G_{\text{orig}}(\hat{z}) - b\|_2 \le \Delta^\dagger$. The second term is $\|e^\dagger\|_2$. Hence

$$\|M_\Omega(G_{\text{orig}}(\hat{z}) - G_{\text{orig}}(z^\dagger))\|_2 \le \Delta^\dagger + \|e^\dagger\|_2.$$

Combining the last two displays gives

$$\|G_{\text{orig}}(\hat{z}) - G_{\text{orig}}(z^\dagger)\|_2 \le \frac{\Delta^\dagger + \|e^\dagger\|_2 + \delta_R}{\gamma_R}.$$

Finally, since $y^\star = G_{\mathrm{orig}}(z^\dagger) + r^\dagger$, a triangle inequality yields

$$\|G_{\mathrm{orig}}(\hat{z}) - y^\star\|_2 \le \|G_{\mathrm{orig}}(\hat{z}) - G_{\mathrm{orig}}(z^\dagger)\|_2 + \|r^\dagger\|_2 \le \|r^\dagger\|_2 + \frac{\Delta^\dagger + \|e^\dagger\|_2 + \delta_R}{\gamma_R}.$$

□

The same reasoning extends to a numerical solution $\hat{z}_{\mathrm{num}}$ produced by an iterative optimizer such as Adam. If

$$\mathcal{J}(\hat{z}_{\mathrm{num}}) \le \mathcal{J}(z^\dagger) + \tau,$$

then one replaces $(\Delta^\dagger)^2$ by $(\Delta^\dagger)^2 + \tau$ throughout the argument, so the bounds hold up to an additional optimization tolerance.

Coordinate restriction and incoherent secants. The local S-REC assumption above is not automatic for a coordinate-restriction operator. If $A_\Omega = \sqrt{N/m}\, M_\Omega$ denotes the rescaled restriction operator and Ω is sampled uniformly, then for any fixed $v \in \mathbb{R}^N$,

$$\mathbb{E}\|A_\Omega v\|_2^2 = \|v\|_2^2.$$

Thus, random coordinate restriction preserves the energy of a fixed vector in expectation. However, concentration around this expectation requires that the energy of v is not concentrated on a small number of coordinates. A standard sufficient condition is an incoherence bound of the form

$$\|v\|_\infty^2 \le \frac{\mu}{N}\|v\|_2^2.$$

Under such a condition, random coordinate samples concentrate $\|A_\Omega v\|_2^2$ around $\|v\|_2^2$. For a finite set of candidate signals, a union bound extends this concentration to all pairwise secants, provided the number of sampled coordinates is large enough relative to the incoherence level and the logarithm of the number of pairs. For the continuous decoder range, the analogous requirement would need to hold uniformly over the normalized secants

$$\frac{G_{\mathrm{orig}}(z_1) - G_{\mathrm{orig}}(z_2)}{\|G_{\mathrm{orig}}(z_1) - G_{\mathrm{orig}}(z_2)\|_2}, \qquad z_1, z_2 \in B(z_0, R^\dagger).$$

This strong assumption is not verified analytically for the trained decoder. We therefore do not claim that M_Ω satisfies S-REC over the learned decoder range. The local S-REC serves only as a sufficient theoretical lens; Appendix B.1 empirically assesses the practical separability of decoder-generated candidates under coordinate subsampling.

B Supplementary Diagnostics

B.1 Measurement Separability

Table 12 and Fig. 11 provide an empirical check of the separation property that is relevant for generative compressed sensing. For pairs of decoder-generated training windows, each returned to the original target units, we compute

$$r_{ij} = \frac{\sqrt{N/m}\,\|M_\Omega(y_i - y_j)\|_2}{\|y_i - y_j\|_2}.$$

Values close to one indicate that coordinate subsampling preserves pairwise distances after energy normalization. This diagnostic does not prove S-REC over the learned decoder range. It instead examines the corresponding practical question of whether sparse observations distinguish plausible generated windows. The ratios describe separation of G_{orig} windows. They are not the constants γ_R and δ_R. The median ratio remains close to one in all tested cells. Lower percentiles are generally stronger in denser settings, despite some non-monotonic variation at intermediate ρ values. The weakest case is the sparsest cross-well setting, where only four entries are observed per window, consistent with the larger ill-posedness expected at very low measurement ratios.

B.2 Full F03-4 benchmark table

Table 13 reports the complete single-well F03-4 GR-only porosity benchmark used in Section 6.2. In addition to mean RMSE, it includes seed-level RMSE variability (standard deviation and Student-t 95% half-width with df = 2), MAE, and relative ℓ_2 error for all seven methods at each tested measurement ratio.

Table 12: Empirical S-REC-style separation diagnostics on decoder-generated windows. Values are percentiles of the energy-normalized ratio $\sqrt{N/m}\|M_\Omega(y_i - y_j)\|_2/\|y_i - y_j\|_2$.

Dataset	ρ	m	P1	P5	P10	Median
Cross-well V_c	0.05	4	0.320	0.483	0.567	0.963
Cross-well V_c	0.10	6	0.291	0.527	0.646	1.004
Cross-well V_c	0.20	13	0.770	0.882	0.930	1.039
F03-4 GR-only	0.20	13	0.641	0.753	0.815	1.000
F03-4 GR-only	0.40	26	0.600	0.767	0.833	0.999
F03-4 GR-only	0.60	38	0.856	0.902	0.927	0.998

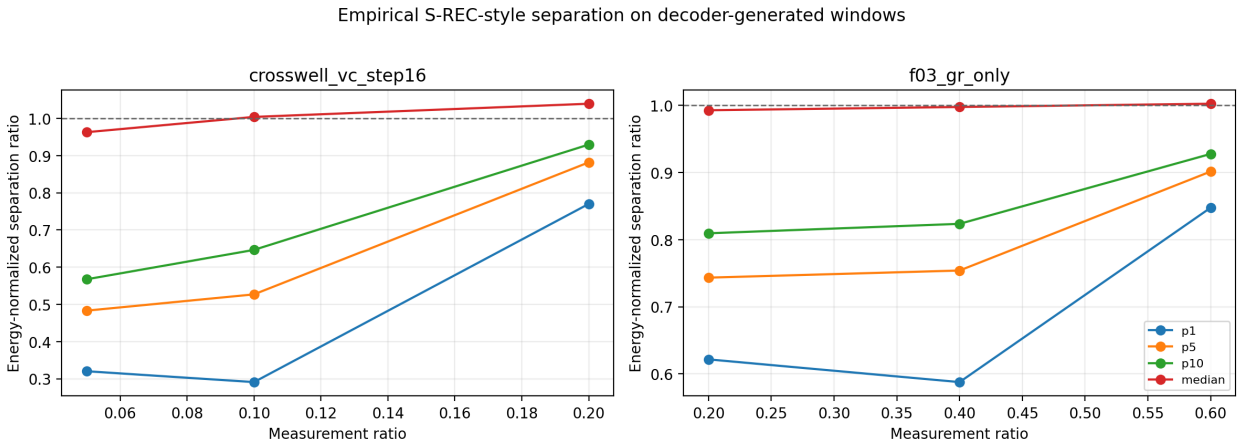


Figure 11: Empirical S-REC-style separation diagnostics. The curves show lower percentiles and medians of the energy-normalized pairwise separation ratio on decoder-generated windows returned to the original target units. Higher lower-percentile values indicate stronger separation of plausible generated candidates under coordinate subsampling.

B.3 F03-4 Metric Robustness Views

The F03-4 GR-only discussion in Section 6.2 uses RMSE because the protocol selects and reports models on that scale. Figures 12 and 13 evaluate the same interpretation under alternative error summaries. MAE downweights isolated large deviations relative to RMSE. Relative ℓ_2 error indicates whether the ranking mainly reflects a scale effect. Both plots retain the sparse-regime pattern: CLP-CSGM Ridge is lowest or close to lowest at lower measurement ratios, and AE $[u, b]$ becomes more competitive as more entries of b are available. The degradation of direct MLP $[u, b]$ also appears in more than one metric, consistent with sparse measurements acting as high-variance features when

Table 13: Full single-well F03-4 GR-only porosity benchmark by measurement ratio, including direct baselines, structural reference variants, and multi-metric summaries across seeds. The column rmse_t95_half is the Student-t 95% half-width for the seed mean (df = 2), not a normal $z = 1.96$ interval.

rho	method	rmse	rmse_std	rmse_t95_half	mae	relative_l2
0.20000	CLP-CSGM Ridge	0.01596	0.00010	0.00024	0.01409	0.06000
0.20000	AE [u,b]	0.01850	0.00067	0.00166	0.01619	0.06951
0.20000	Prior-only Ridge	0.01958	0.00009	0.00022	0.01774	0.07360
0.20000	PCA [u,b]	0.02149	0.00093	0.00230	0.01796	0.08073
0.20000	ML only	0.02359	0.00051	0.00126	0.02086	0.08866
0.20000	Measurement-only CSGM	0.02519	0.00145	0.00360	0.01993	0.09459
0.20000	MLP [u,b]	0.03022	0.00059	0.00146	0.02588	0.11356
0.30000	CLP-CSGM Ridge	0.01475	0.00067	0.00166	0.01294	0.05544
0.30000	AE [u,b]	0.01554	0.00011	0.00027	0.01351	0.05837
0.30000	Prior-only Ridge	0.01958	0.00009	0.00022	0.01774	0.07360
0.30000	PCA [u,b]	0.02187	0.00130	0.00323	0.01799	0.08214
0.30000	Measurement-only CSGM	0.02285	0.00117	0.00290	0.01786	0.08578
0.30000	ML only	0.02359	0.00051	0.00126	0.02086	0.08866
0.30000	MLP [u,b]	0.02997	0.00270	0.00670	0.02555	0.11253
0.40000	AE [u,b]	0.01304	0.00013	0.00033	0.01116	0.04897
0.40000	CLP-CSGM Ridge	0.01377	0.00052	0.00129	0.01197	0.05176
0.40000	Prior-only Ridge	0.01958	0.00009	0.00022	0.01774	0.07360
0.40000	Measurement-only CSGM	0.02033	0.00182	0.00452	0.01560	0.07635
0.40000	PCA [u,b]	0.02261	0.00172	0.00427	0.01839	0.08490
0.40000	ML only	0.02359	0.00051	0.00126	0.02086	0.08866
0.40000	MLP [u,b]	0.03712	0.00166	0.00412	0.03244	0.13932
0.50000	CLP-CSGM Ridge	0.01281	0.00048	0.00119	0.01110	0.04813
0.50000	AE [u,b]	0.01286	0.00010	0.00025	0.01107	0.04832
0.50000	Measurement-only CSGM	0.01622	0.00087	0.00217	0.01281	0.06089
0.50000	Prior-only Ridge	0.01958	0.00009	0.00022	0.01774	0.07360
0.50000	ML only	0.02359	0.00051	0.00126	0.02086	0.08866
0.50000	PCA [u,b]	0.02409	0.00301	0.00748	0.01952	0.09050
0.50000	MLP [u,b]	0.04199	0.00227	0.00564	0.03757	0.15764
0.60000	AE [u,b]	0.01178	0.00023	0.00058	0.01003	0.04425
0.60000	CLP-CSGM Ridge	0.01220	0.00041	0.00102	0.01051	0.04583
0.60000	Measurement-only CSGM	0.01494	0.00081	0.00200	0.01176	0.05611
0.60000	Prior-only Ridge	0.01958	0.00009	0.00022	0.01774	0.07360
0.60000	ML only	0.02359	0.00051	0.00126	0.02086	0.08866
0.60000	PCA [u,b]	0.02426	0.00114	0.00284	0.01976	0.09109
0.60000	MLP [u,b]	0.04627	0.00483	0.01199	0.04076	0.17381

the measurement operator is not handled explicitly.

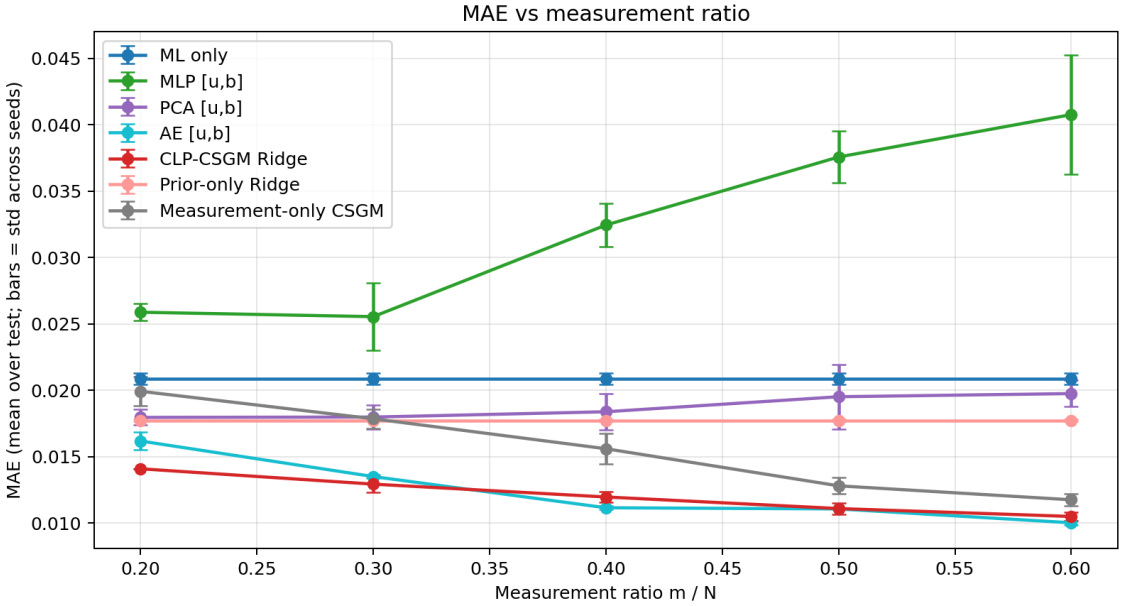


Figure 12: Single-well F03-4 GR-only MAE by measurement ratio. The ranking is consistent with the RMSE view in the main text: CLP-CSGM Ridge has the lowest MAE in the sparse regime and remains competitive as the measurement ratio increases.

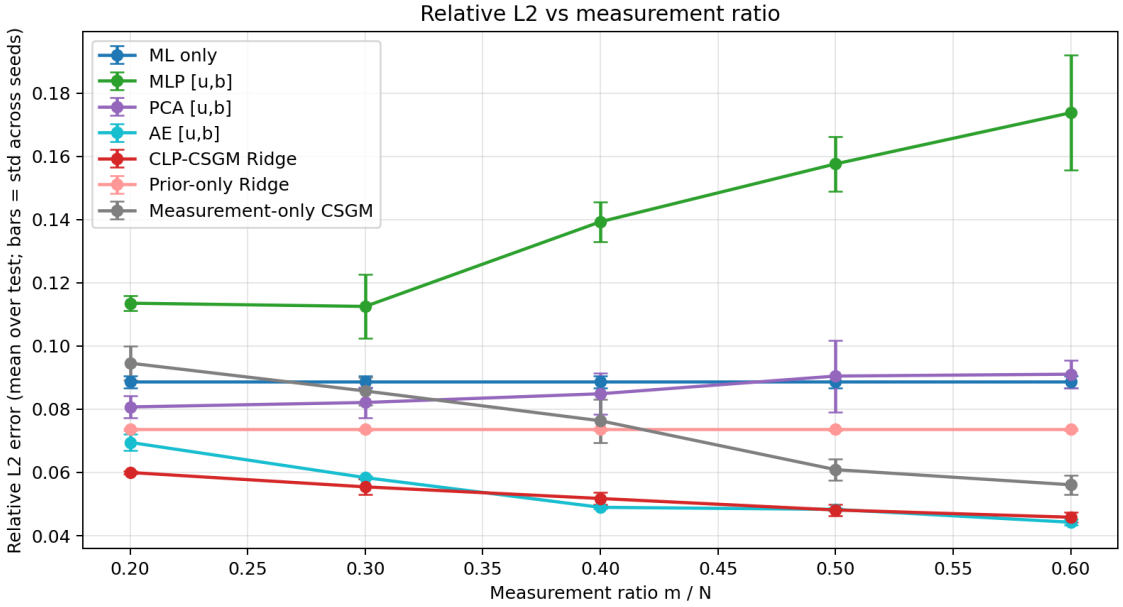


Figure 13: Relative ℓ_2 error on the single-well F03-4 GR-only benchmark. The relative-error view preserves the same conclusion as RMSE and MAE, indicating that the gain is not caused by scale effects alone.

B.4 Illustrative sparse-measurement windows

The main tables and curves in Section 6 pool scalar errors over many test windows. The examples in Figs. 14 and 15 display the target profile, the coordinates where b is observed, and the full-window predictions generated under the paired (seed, ρ) design used in the main experiments, thereby complementing the aggregate comparisons with window-level views.

The horizontal axis indexes coordinates within one length-L window ($L = 64$). The blue curve is the target window y. Black markers show the realized noisy measurements b at the sampled coordinates. The red curve is CLP-CSGM Ridge and the green curve is AE $[u, b]$, when shown. These markers are the actual sparse inputs generated under the benchmark noise model.

Each figure fixes one test-window index, one seed, and one ρ. These audit examples show that the markers alone underdetermine the reconstruction and that the methods interpolate between sparse constraints using different inductive biases. In the displayed porosity window, both CLP-CSGM Ridge and AE $[u, b]$ miss the central depression of the reference. In the displayed clay-volume window, the corresponding configuration is step 16, seed 7, $\rho = 0.10$, and window index 0. CLP-CSGM Ridge is closer to the reference, the autoencoder prediction is positively offset, and local deviations remain. These windows do not establish a uniform advantage. Aggregate tables and the main-text diagnostics provide the overall performance comparisons.

B.5 F03-4 Qualitative Reconstruction Views

This subsection complements the depth-profile reconstruction in the main text (Fig. 10). Because the depth profile follows the contiguous test interval, it provides the clearest qualitative view. The figures below present window-level and pointwise audits.

Figure 16 shows representative F03-4 test windows and reveals local shape behavior compressed by scalar metrics. Direct MLP $[u, b]$ can produce oscillatory predictions, whereas PCA, AE, and CLP-CSGM Ridge are smoother. In these displayed windows, CLP-CSGM Ridge and the autoencoder miss part of the central depression of the reference. Aggregate metrics determine the window-level rankings. Figure 17 pools pointwise predictions against ground truth over the F03-4 test windows. The figure prints a correlation of 0.4632 for CLP-CSGM Ridge and 0.3720 for AE $[u, b]$. The scatter remains broad. This comparison describes the pooled points on the plot. It

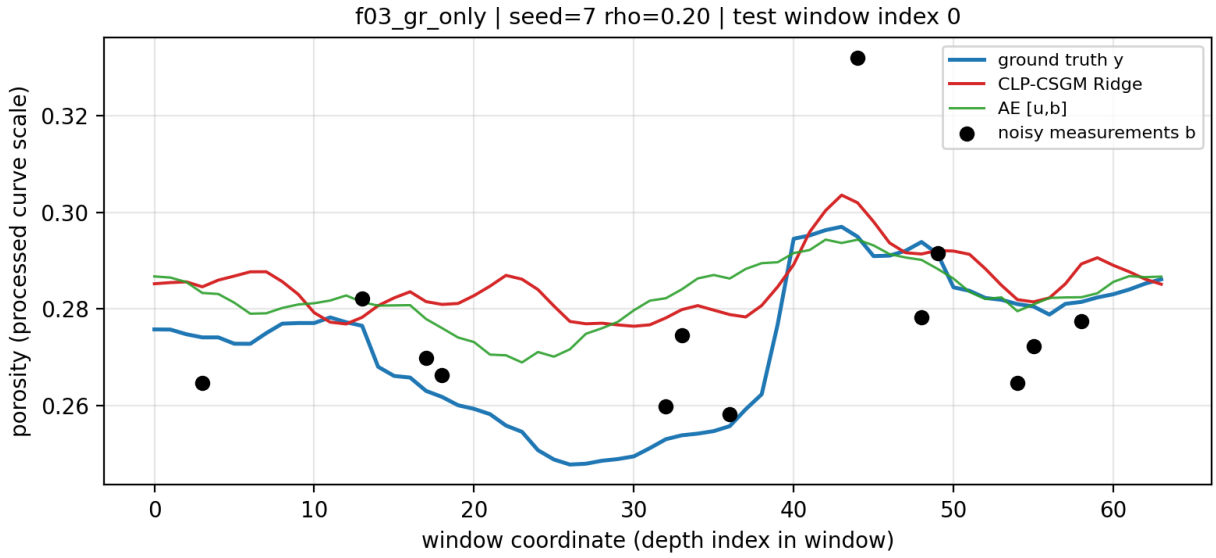


Figure 14: F03-4 GR-only porosity example (test window index 0, seed 7, $\rho = 0.20$): ground truth, sparse noisy measurements b, CLP-CSGM Ridge, and AE $[u, b]$. In this displayed window, both reconstructions miss the central depression of the reference.

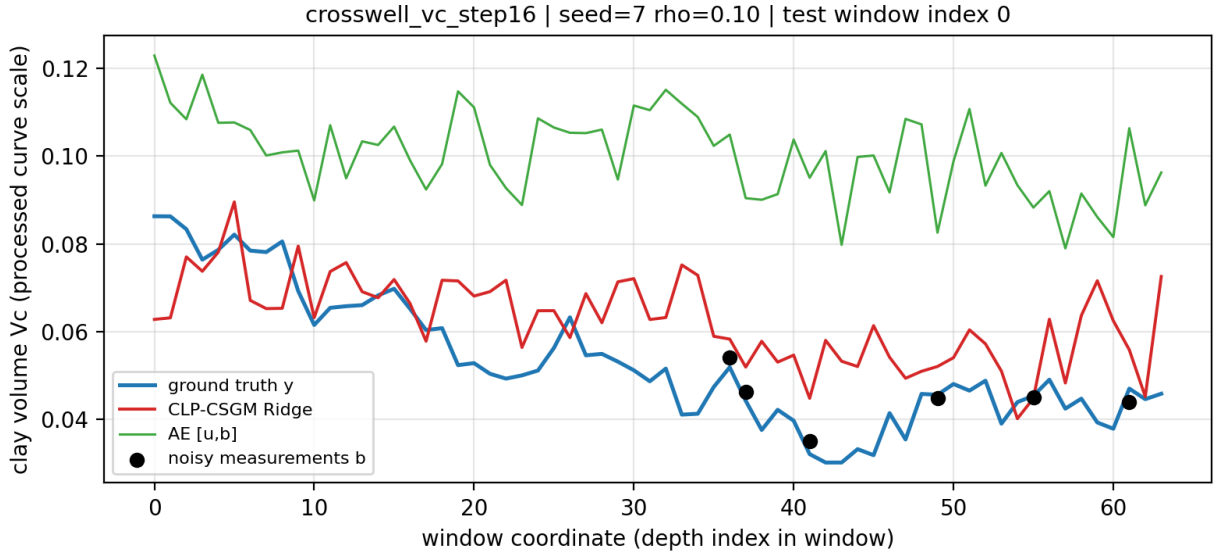


Figure 15: Cross-well V_c example (held-out F03-4 test window index 0, step 16, seed 7, $\rho = 0.10$): ground truth, sparse noisy measurements b, CLP-CSGM Ridge, and AE $[u, b]$. In this displayed window, CLP-CSGM Ridge is closer to the reference. The autoencoder prediction is positively offset, and local deviations remain.

does not establish statistical significance, and overlapping-window predictions are not independent observations.

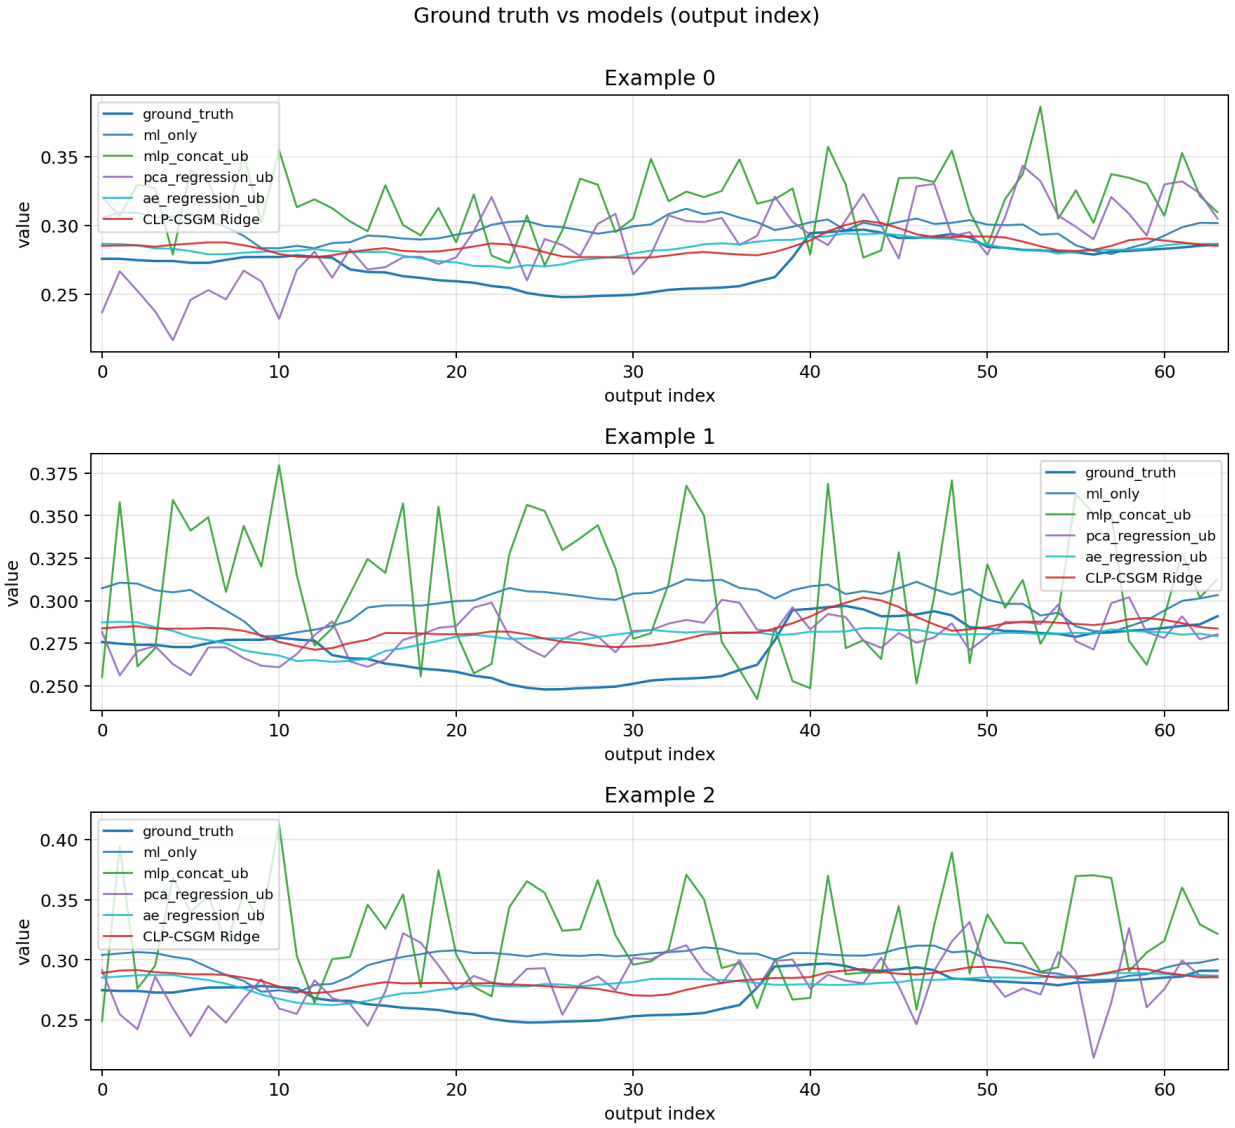


Figure 16: Representative F03-4 test windows. Direct MLP $[u, b]$ is visibly more oscillatory. CLP-CSGM Ridge and AE $[u, b]$ are smoother and, in these displayed windows, miss part of the central depression of the reference.

B.6 Validation and Generalization Diagnostics

Figures 18 and 19 audit quantities used by the experimental protocol: validation RMSE for selecting λ, selected regularization values, per-window test RMSE distributions, and the relation between autoencoder training reconstruction error and final test RMSE. These diagnostics support

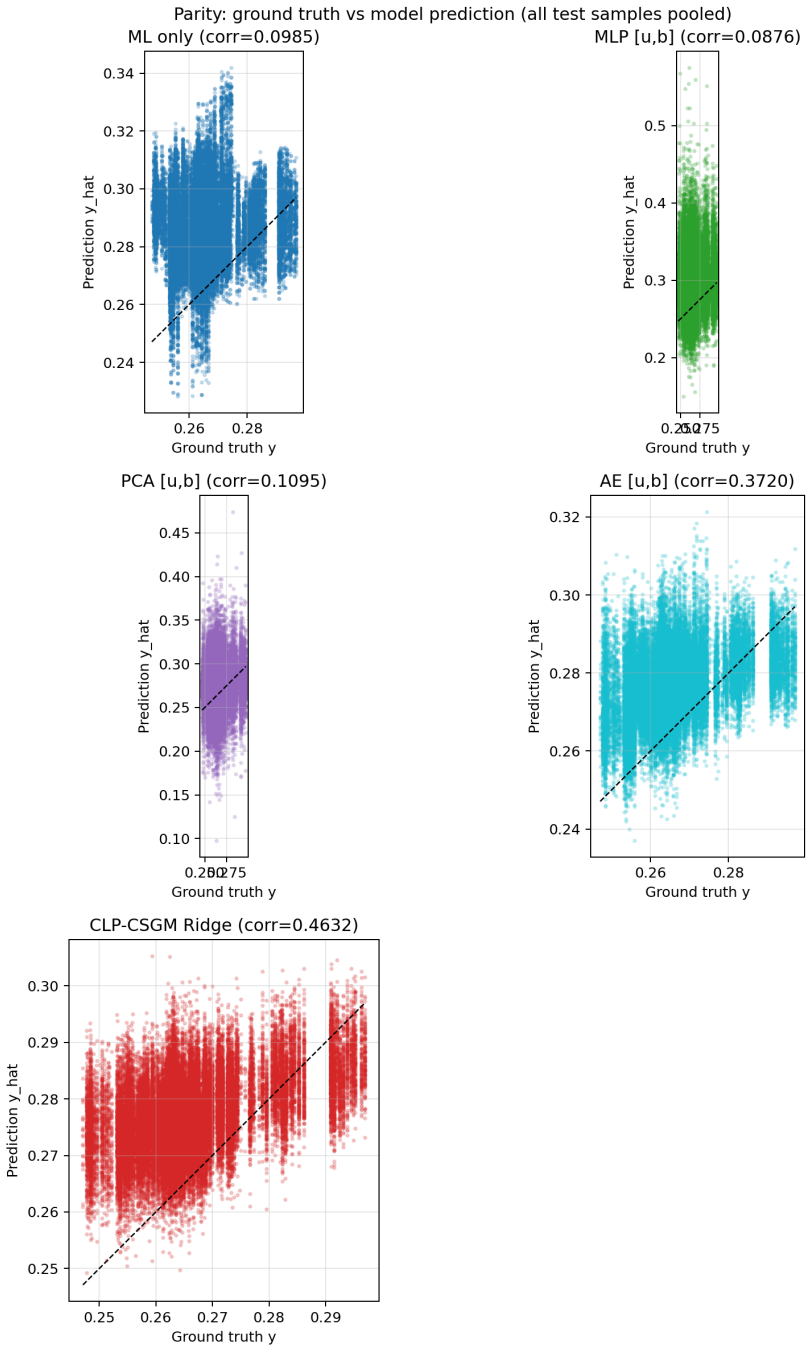


Figure 17: Pooled parity plots for F03-4 test windows. The displayed correlations are 0.4632 for CLP-CSGM Ridge and 0.3720 for AE $[u, b]$. Broad scatter remains. These values describe the pooled points on the figure and are not a significance test.

the reliability and interpretation of the main rankings.

Figure 18 summarizes the cross-well V_c case. The validation and test RMSE values remain of similar magnitude across the low-data grid, with no sign that lower validation error is obtained only by shifting error to the test well. The selected λ values are small but nonzero, consistent with sparse target-log observations providing the main target-well information while the log-conditioned latent prior acts as a stabilizing regularizer. The per-window RMSE boxplots reinforce the ablation result: CLP-CSGM Ridge has a lower median error than the latent-prior-only decoder and a substantially tighter distribution than measurement-only CSGM.

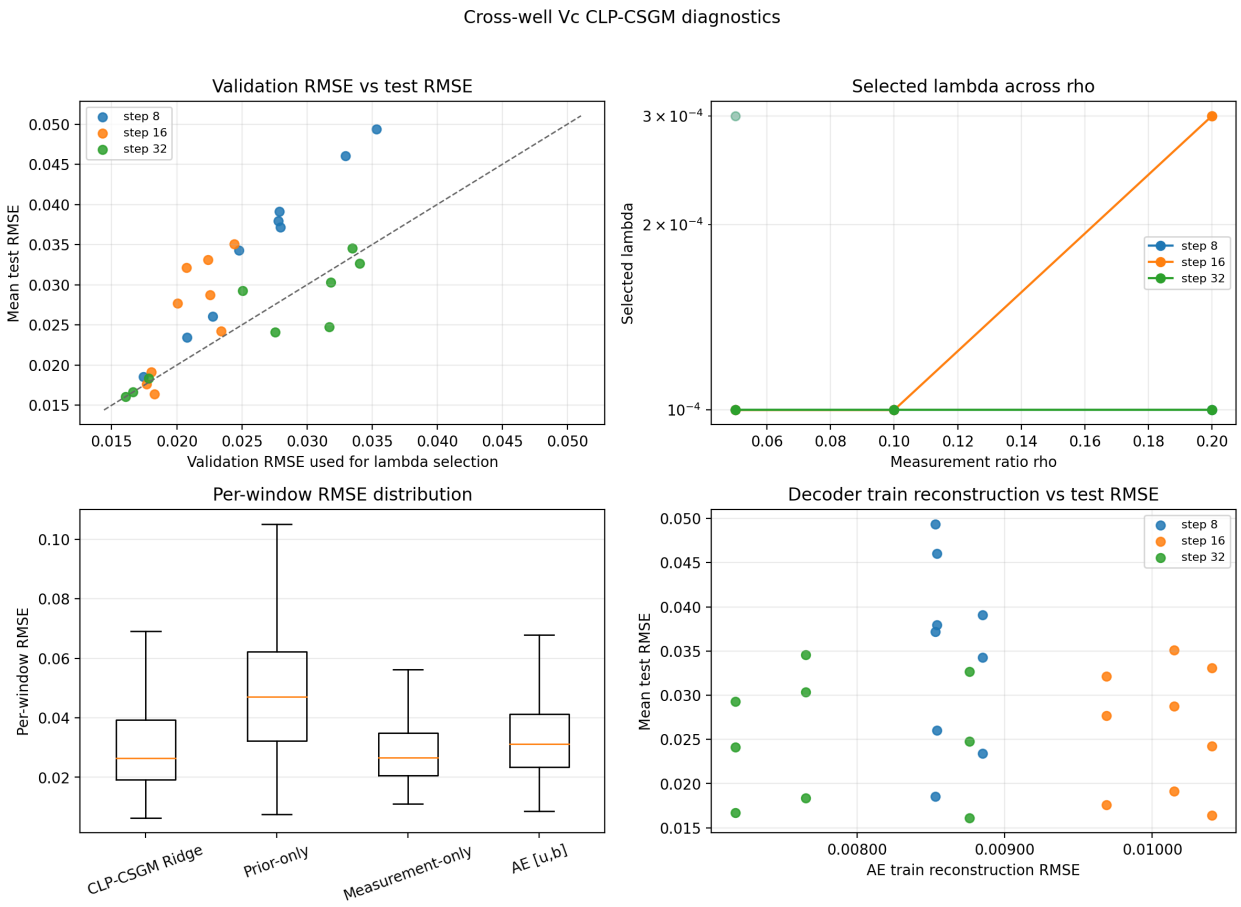


Figure 18: Cross-well V_c diagnostics for CLP-CSGM Ridge summarizing the experiments in Section 6.1. The panels compare validation RMSE with mean test RMSE for CLP-CSGM Ridge, show the selected λ values across measurement ratios, summarize per-window RMSE distributions, and relate autoencoder training reconstruction error to final test RMSE.

Figure 19 gives the analogous diagnostic view for the F03-4 porosity benchmark with GR-only dense input. The validation–test relation is less diagonal than in the cross-well case, consistent

with the harder stress test: the dense input is restricted to GR, and the test block is a contiguous depth interval, not a random split. The trend across measurement ratios remains coherent. Larger ρ generally reduces the full-model test error, and the selected λ remains stable across ratios. The per-window RMSE distributions show that CLP-CSGM Ridge improves the central error relative to both ablated variants while remaining close to AE $[u, b]$, matching the main F03 regime-dependent result.

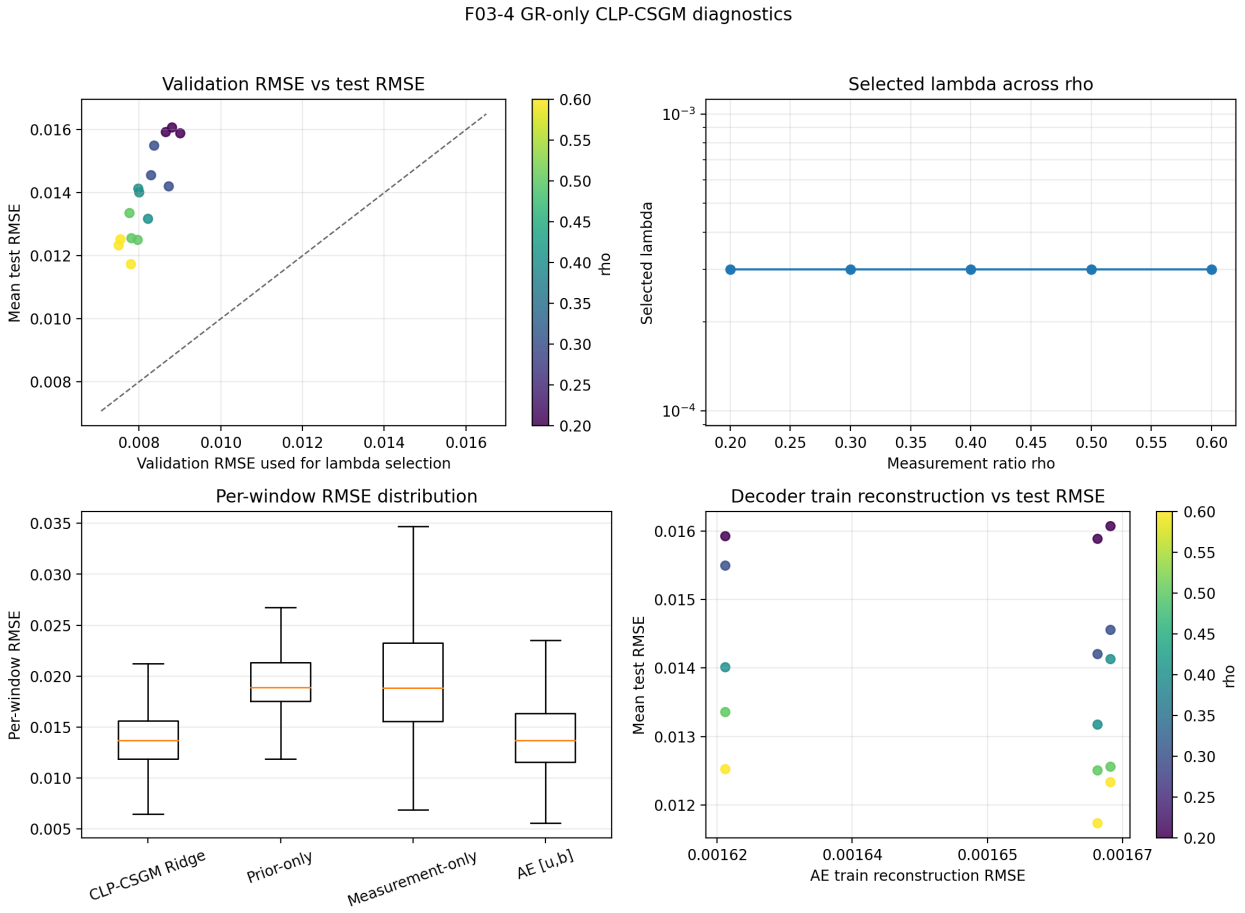


Figure 19: F03-4 GR-only diagnostics for CLP-CSGM Ridge summarizing the experiments in Section 6.2. The panels report validation RMSE versus mean test RMSE, selected λ, per-window RMSE distributions, and the relation between autoencoder training reconstruction error and final test RMSE.

B.7 Cross-well Channel-Ablation and Embargo-Style Boundary Check

This subsection reports a paired robustness check for the cross-well V_c step-16 setting. It tests two possible sensitivities: reintroducing GR into u, relative to the main no-GR protocol in Section 5, and

removing overlap-prone windows near the train/validation boundary. The latter is an embargo-style boundary check: a small band of windows adjacent to the boundary is removed before validation to create a gap between training and validation windows. We compare four scenarios under the same seeds and measurement-ratio grid, $\rho \in \{0.05, 0.10, 0.20\}$:

- No boundary check + GR in u (stronger dense context reference).
- Embargo-style removal (4 boundary windows) + GR in u.
- No boundary check + GR excluded from u (matches the main cross-well u).
- Embargo-style removal (4 boundary windows) + GR excluded from u.

Operationally, the embargo-style removal drops four overlap-prone boundary windows before the validation tail in each training well. In this setup, n_{train} decreases from 52 to 41, while n_{val} and n_{test} stay fixed.

Table 14: Cross-well step-16 robustness check across GR ablation and embargo-style boundary-check scenarios. Ridge RMSE is reported as mean ± std over three seeds. Gains are relative RMSE reductions of CLP-CSGM Ridge against AE $[u, b]$ and measurement-only CSGM.

scenario	rho	ridge rmse	ae rmse	meas-only rmse	ridge gain vs ae (%)	ridge gain vs meas-only (%)
Emb4+GR	0.05	0.01487 ± 0.00035	0.02056	0.03647	27.7	59.2
Emb4+GR	0.10	0.01406 ± 0.00032	0.01953	0.02849	28.0	50.7
Emb4+GR	0.20	0.01282 ± 0.00026	0.01932	0.02167	33.6	40.8
Emb4-noGR	0.05	0.03117 ± 0.00020	0.03781	0.03647	17.6	14.5
Emb4-noGR	0.10	0.02549 ± 0.00268	0.03439	0.02849	25.9	10.6
Emb4-noGR	0.20	0.01625 ± 0.00139	0.03771	0.02167	56.9	25.0
NoEmb+GR	0.05	0.01403 ± 0.00064	0.01893	0.03622	25.9	61.3
NoEmb+GR	0.10	0.01380 ± 0.00107	0.01802	0.02954	23.4	53.3
NoEmb+GR	0.20	0.01289 ± 0.00052	0.01822	0.02127	29.3	39.4
NoEmb-noGR	0.05	0.03344 ± 0.00151	0.03712	0.03622	9.9	7.7
NoEmb-noGR	0.10	0.02689 ± 0.00235	0.03420	0.02954	21.4	9.0
NoEmb-noGR	0.20	0.01771 ± 0.00138	0.03214	0.02127	44.9	16.7

Table 14 shows three stable patterns. Within each boundary protocol, including GR in u lowers absolute RMSE, as expected because V_c is Larionov-derived from GR. Under the no-GR protocol matching the main benchmark, CLP-CSGM Ridge has the lowest RMSE relative to AE $[u, b]$ and measurement-only CSGM throughout the tested grid. The embargo-style boundary check changes the absolute values without reversing this ordering. These checks support the main interpretation: the cross-well V_c benchmark tests sparse target-log reconstruction under distribution shift without a GR shortcut in u, and the reported gains are not solely an artifact of train/validation boundary overlap.

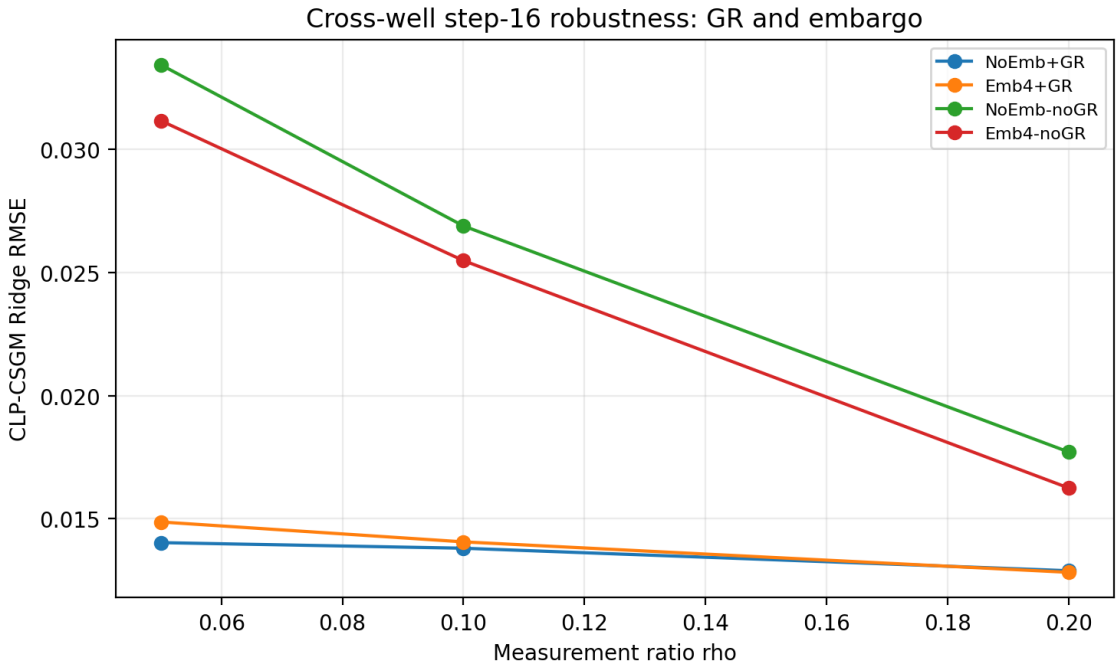


Figure 20: Cross-well step-16 RMSE sensitivity for CLP-CSGM Ridge across GR inclusion in u and embargo-style boundary-check scenarios. Omitting GR from u relative to GR-inclusive rows increases error; the boundary check shifts absolute values without overturning the ranking trends.

B.8 Step-32 development-split resampling

The main cross-well grid uses a depth-tail validation split. At step 32 the development pool is small ($n_{\text{train}} = 27$, $n_{\text{val}} = 3$ after the split). To test whether the CLP-CSGM Ridge versus AE $[u, b]$ gap in that cell depends on the particular validation windows, we redrew the same validation budget ten times inside each development well (`random_per_well`), keeping F03-4 as the held-out test well and freezing the sparse coordinates Ω to the paper seed 7. Prior-only and measurement-only ablations were not rerun.

Table 15: Cross-well step-32 resampling of the development train/validation split ($n = 10$). RMSE is the mean ± standard deviation across splits. Positive mean ΔRMSE (AE − CLP) favors CLP-CSGM Ridge. CLP win frac is the fraction of splits with lower CLP RMSE.

ρ	n_{splits}	n_{train}	n_{val}	CLP RMSE	AE RMSE	mean ΔRMSE	CLP win frac
0.05	10	27	3	0.0342 ± 0.0024	0.0407 ± 0.0025	0.0064	1.00
0.10	10	27	3	0.0305 ± 0.0032	0.0388 ± 0.0015	0.0083	1.00
0.20	10	27	3	0.0188 ± 0.0010	0.0362 ± 0.0064	0.0174	1.00

Table 15 shows that CLP-CSGM Ridge has lower mean RMSE than AE $[u, b]$ at every tested ρ, and lower RMSE in all ten splits. This check does not add a second held-out well and does not

vary Ω; it only resamples the small development split.

B.9 Ten-seed CLP versus AE panel

The main cross-well intervals use three seeds (t critical value with df = 2). Several of those paired ΔRMSE intervals include zero, and the CLP-CSGM Ridge versus AE $[u, b]$ ranking changes between step 8 and steps 16–32. We therefore reran that pairwise comparison with ten seeds $\{7, 23, 41, 13, 29, 31, 37, 43, 47, 53\}$, keeping the submitted tail validation split, the no-GR input, and F03-4 as the held-out well. Direct MLP and PCA, structural ablations, and the F03-4 grid were not rerun.

Table 16: Cross-well paired $\Delta\text{RMSE} = \text{RMSE}_{\text{AE}} - \text{RMSE}_{\text{CLP}}$ on ten seeds. SE is the seed-level standard error; the 95% half-width uses Student-t with df = 9. Positive mean ΔRMSE favors CLP-CSGM Ridge. Wins count seeds with strictly lower RMSE.

step	ρ	n	mean ΔRMSE	SE	t 95% half-width	CLP wins	AE wins	CI includes 0
8	0.05	10	-0.00826	0.00121	0.00273	0	10	no
8	0.10	10	-0.00318	0.00092	0.00209	2	8	no
8	0.20	10	0.00360	0.00129	0.00293	7	3	no
16	0.05	10	0.00434	0.00159	0.00360	9	1	no
16	0.10	10	0.00677	0.00134	0.00304	10	0	no
16	0.20	10	0.01382	0.00145	0.00328	10	0	no
32	0.05	10	0.00865	0.00129	0.00292	10	0	no
32	0.10	10	0.01217	0.00118	0.00266	10	0	no
32	0.20	10	0.01934	0.00185	0.00418	10	0	no

Table 16 shows that none of the nine ten-seed intervals include zero. The CLP-CSGM Ridge advantage is concentrated where training support is smallest (steps 16 and 32, approximately 52 and 27 windows): lower RMSE at every ρ. At step 8, where the development pool is largest (approximately 104 windows), AE $[u, b]$ has lower RMSE at $\rho = 0.05$ and $\rho = 0.10$; CLP-CSGM Ridge has lower RMSE only at $\rho = 0.20$. This is a low-training-support pattern, not a uniform gain at every small measurement ratio. The submitted three-seed tables are unchanged. Each of these seeds also draws its own coordinate set, so the panel does not separate training randomness from observation placement. Appendix B.11 holds the model seeds fixed and redraws only that coordinate set.

B.10 Explicit-mask direct AE

The main direct baselines receive $[\boldsymbol{u}, \boldsymbol{b}]$ with positional identity implicit in the column order. CLP-CSGM instead applies M_Ω to the length-N window. To test whether the CLP-CSGM gain is only that coordinate placement, we fit one additional direct AE on the same nine cells, seeds, and noise. Its input is $[\boldsymbol{u}, \boldsymbol{b}_\Omega, \text{mask}]$: $\boldsymbol{b}_\Omega \in \mathbb{R}^N$ holds the noisy measurement at each observed coordinate and zero elsewhere, and the mask is the binary indicator of Ω. MLP, PCA, ablations, and the F03-4 grid were not rerun, and this method is not added to the main tables.

Table 17: Cross-well CLP-CSGM Ridge versus AE $[\boldsymbol{u}, \boldsymbol{b}]$ and an AE that also receives an explicit window mask ($n = 3$ seeds). Δ is the other method's RMSE minus CLP RMSE. Positive Δ favors CLP-CSGM Ridge. Wins count seeds with strictly lower CLP RMSE.

step	ρ	CLP	AE $[u, b]$	AE mask	Δ AE	Δ mask	CLP>AE	CLP>mask
8	0.05	0.0448	0.0357	0.0340	-0.0092	-0.0108	0/3	0/3
8	0.10	0.0365	0.0320	0.0314	-0.0045	-0.0051	1/3	0/3
8	0.20	0.0227	0.0260	0.0258	0.0033	0.0032	2/3	2/3
16	0.05	0.0334	0.0371	0.0352	0.0037	0.0018	3/3	2/3
16	0.10	0.0269	0.0342	0.0340	0.0073	0.0071	3/3	3/3
16	0.20	0.0177	0.0321	0.0277	0.0144	0.0100	3/3	3/3
32	0.05	0.0322	0.0390	0.0426	0.0068	0.0105	3/3	3/3
32	0.10	0.0264	0.0389	0.0386	0.0125	0.0122	3/3	3/3
32	0.20	0.0170	0.0396	0.0366	0.0225	0.0196	3/3	3/3

Table 17 shows that the explicit mask does not remove the CLP-CSGM advantage where training support is smallest. At steps 16 and 32, CLP-CSGM Ridge has lower mean RMSE than the masked AE at every ρ. At step 8, the masked AE remains better at $\rho = 0.05$ and $\rho = 0.10$, as the unmasked AE does. The mask shifts the direct AE only modestly and not always downward: at step 32, $\rho = 0.05$, the masked AE has higher mean RMSE than AE $[\boldsymbol{u}, \boldsymbol{b}]$. This check keeps Ω fixed inside each cell; it does not test sensitivity to observation placement. That placement check is Appendix B.11.

B.11 Sparse-coordinate placement

The main protocol samples one coordinate set Ω for each model seed and measurement ratio and then holds it fixed across windows. To test whether the CLP-CSGM Ridge versus AE $[\boldsymbol{u}, \boldsymbol{b}]$ ranking depends on those locations, we redrew Ω five times on the cross-well V_c grid. The model seeds stay at $\{7, 23, 41\}$, so the decoder and the ridge anchor keep the same initialization. The first draw is

the published coordinate set. The other four draws change only the generator that samples Ω and the noise on $\boldsymbol{b}$. The direct AE is fit again on each draw, because $\boldsymbol{b}$ is part of its input, and the CLP-CSGM penalty weight is selected again on the validation windows of that draw. The score is the same full-window RMSE as the main grid. Direct MLP, PCA, structural ablations, and the F03-4 grid were not rerun, and these draws are not merged into the submitted three-seed tables.

Table 18: Cross-well CLP-CSGM Ridge versus AE $[\boldsymbol{u}, \boldsymbol{b}]$ across five sparse-coordinate placements and the three paper seeds ($n = 15$ per cell). RMSE values are means over those comparisons. Positive mean ΔRMSE (AE − CLP) favors CLP-CSGM Ridge. Wins count comparisons with strictly lower RMSE.

step	ρ	n	CLP RMSE	AE RMSE	mean ΔRMSE	CLP wins	AE wins
8	0.05	15	0.04231	0.03586	-0.00645	0	15
8	0.10	15	0.03581	0.03225	-0.00356	5	10
8	0.20	15	0.02383	0.02591	0.00208	11	4
16	0.05	15	0.03072	0.03710	0.00638	15	0
16	0.10	15	0.02615	0.03412	0.00797	15	0
16	0.20	15	0.01840	0.03036	0.01196	15	0
32	0.05	15	0.02895	0.03873	0.00977	15	0
32	0.10	15	0.02470	0.03870	0.01401	15	0
32	0.20	15	0.01795	0.04023	0.02228	15	0

Table 18 shows that the cell-level winner matches the published ranking in all nine cells. At steps 16 and 32, CLP-CSGM Ridge has lower RMSE in all fifteen comparisons at every ρ. At step 8 and $\rho = 0.05$, AE $[\boldsymbol{u}, \boldsymbol{b}]$ is lower in all fifteen. The two closer step-8 cells are not unanimous: five of fifteen comparisons favor CLP-CSGM Ridge at $\rho = 0.10$, and four of fifteen favor the direct AE at $\rho = 0.20$. The low-training-support advantage therefore does not depend on the one published coordinate set. The step-8 ranking is the less stable of the two regimes, and a single extra draw can reverse an individual seed there. This check does not add seeds beyond $\{7, 23, 41\}$.

B.12 Direct linear maps, depth interpolation, and measurement noise

The main grid does not include a direct linear map from $[\boldsymbol{u}, \boldsymbol{b}]$ or an interpolation of the sparse measurements. This supplement adds both on the cross-well V_c grid, with the same windows, seeds $\{7, 23, 41\}$, measurement ratios, and published coordinate draws. Ordinary least squares and ridge regression map standardized $[\boldsymbol{u}, \boldsymbol{b}]$ to the window. The ridge penalty is chosen on validation full-window RMSE. Interpolation connects the noisy measurements along the depth index of the window and keeps the nearest measured value at the ends. It does not use the dense logs. The score is the

same full-window RMSE as the main grid. CLP-CSGM Ridge and AE $[\boldsymbol{u}, \boldsymbol{b}]$ in Table 19 are the published offset-0 runs. These comparators are not merged into the submitted tables.

Table 19: Cross-well full-window RMSE at the published noise standard deviation 0.01. Values are means over seeds $\{7, 23, 41\}$. Linear $[u, b]$ is ordinary least squares and Ridge direct is validation-selected ridge, both from $[\boldsymbol{u}, \boldsymbol{b}]$ to the window. Interpolation uses only the sparse measurements. The lowest value in every cell is interpolation.

step	ρ	Linear $[u, b]$	Ridge direct	Interpolation	CLP-CSGM	AE $[u, b]$
8	0.05	0.03122	0.02916	0.01779	0.04483	0.03567
8	0.10	0.02790	0.02384	0.01537	0.03646	0.03195
8	0.20	0.02158	0.01862	0.01100	0.02267	0.02600
16	0.05	0.02903	0.02910	0.01751	0.03344	0.03712
16	0.10	0.02634	0.02661	0.01478	0.02689	0.03420
16	0.20	0.02062	0.02089	0.01115	0.01771	0.03214
32	0.05	0.03382	0.03980	0.01755	0.03217	0.03901
32	0.10	0.03083	0.03878	0.01533	0.02640	0.03894
32	0.20	0.02421	0.02648	0.01125	0.01705	0.03959

Table 19 shows that interpolation has the lowest RMSE in all nine cells. Seed by seed, it is also lowest in all 27 comparisons. CLP-CSGM Ridge is lower than both direct linear maps at step 16, $\rho = 0.20$, and at step 32 for every ρ. At step 8, and at step 16 for $\rho = 0.05$ and $\rho = 0.10$, one or both linear maps are lower than CLP-CSGM Ridge. The measurements at this noise level are a slightly perturbed copy of the same smooth curve, so connecting them is enough.

The same windows and the same standard-normal draws are scaled to noise standard deviations 0.01, 0.05, 0.10, and 0.20. Table 20 and Fig. 6 use one refit of that series. The decoder and the dense-log map are fit once per step and seed, without the sparse measurements. For each measurement fraction and each noise level, the penalty weight is selected again on validation full-window RMSE. The reported CLP-CSGM values are that selected-penalty result, not a reconstruction with one fixed weight. The sensitivity sweep uses a different, wider grid, with candidates 10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1, 10, and 100. It omits the factor-of-three candidates from the published grid. At 0.01 and at 0.05 every selected weight also belongs to the published grid. That membership does not mean the two searches were the same. At noise 0.10, step 16 and $\rho = 0.10$, one seed selects the weight 1, above the published grid and still inside the sweep. At noise 0.20, some cells select 100, the largest candidate that was tried. That selection is not an optimum outside the searched grid. At noise 0.01 the ordering matches Table 19, but the cell means are those of this refit. At step 16 and $\rho = 0.20$, the published CLP-CSGM mean is 0.01771 and the refit mean is 0.0159. In that cell the selected

weight is 0.0001, inside the published grid. The held-out curve has standard deviation 0.035. The window predicted from the dense logs alone, with no measurement term, has mean RMSE 0.0564 at step 8, 0.0426 at step 16, and 0.0432 at step 32.

Table 20: Interpolation versus CLP-CSGM Ridge when only the measurement-noise standard deviation changes. RMSE values are means over the three paper seeds. Wins count seeds in which CLP-CSGM Ridge has strictly lower RMSE. The held-out V_c standard deviation is 0.035, on the same scale as the noise. Noise 0.05 is about 1.43 times that deviation. Noise levels 0.10 and 0.20 are about 2.86 and 5.71 times that deviation. These noise levels were used for sensitivity analysis and were not calibrated to field measurement uncertainties. The decoder and the dense-log map are fit once per step and seed, and the penalty is selected on validation at each noise level. Noise 0.01 in this table is that refit. At 0.20 some selected weights equal 100, the largest candidate on this grid, which is not an optimum outside the grid.

		$\sigma_\eta = 0.01$			$\sigma_\eta = 0.05$			$\sigma_\eta = 0.10$			$\sigma_\eta = 0.20$		
step	ρ	Interp.	CLP	wins	Interp.	CLP	wins	Interp.	CLP	wins	Interp.	CLP	wins
8	0.05	0.0178	0.0456	0	0.0456	0.0508	1	0.0864	0.0547	3	0.1703	0.0557	3
8	0.10	0.0154	0.0365	0	0.0442	0.0456	1	0.0853	0.0539	3	0.1690	0.0555	3
8	0.20	0.0110	0.0212	0	0.0423	0.0372	3	0.0838	0.0513	3	0.1673	0.0548	3
16	0.05	0.0175	0.0334	0	0.0442	0.0386	3	0.0838	0.0416	3	0.1653	0.0424	3
16	0.10	0.0148	0.0269	0	0.0427	0.0352	3	0.0825	0.0416	3	0.1636	0.0424	3
16	0.20	0.0112	0.0159	0	0.0428	0.0286	3	0.0847	0.0387	3	0.1690	0.0424	3
32	0.05	0.0175	0.0317	0	0.0430	0.0382	3	0.0820	0.0423	3	0.1625	0.0430	3
32	0.10	0.0153	0.0264	0	0.0445	0.0365	3	0.0860	0.0414	3	0.1706	0.0431	3
32	0.20	0.0112	0.0170	0	0.0425	0.0279	3	0.0840	0.0381	3	0.1674	0.0409	3

At noise standard deviation 0.05, CLP-CSGM Ridge is lower in every seed at steps 16 and 32 and at step 8 for $\rho = 0.20$. At step 8 for $\rho = 0.05$ and $\rho = 0.10$, interpolation still has the lower mean. At 0.10 and 0.20, CLP-CSGM Ridge is lower in all nine cells and in all three seeds. Figure 6 shows this ordering with one panel per step. The interpolation error grows because the interpolated curve follows the measurements. The CLP-CSGM error moves less and stays near the dense-log window. Fig. 21 shows one held-out interval at the published noise and at noise 0.05. The two panels use separate horizontal scales. Panel (b) is not one of the larger sensitivity levels. In this displayed interval, interpolation closely follows the reference at noise standard deviation 0.01. At 0.05, it exhibits larger excursions, while CLP-CSGM Ridge retains more of the broad profile structure. The example is consistent with the cell-level result for the corresponding configuration, step 32 and $\rho = 0.20$, and it does not extend to every configuration at that noise level.

On the same splits, the retrospective decoder reconstruction $G(E(\boldsymbol{y}))$ has mean test RMSE 0.0073 at step 8, 0.0099 at step 16, and 0.0110 at step 32. That reconstruction is given the full target window, so it is a check that the decoder can represent the curve. It is not a predictor and

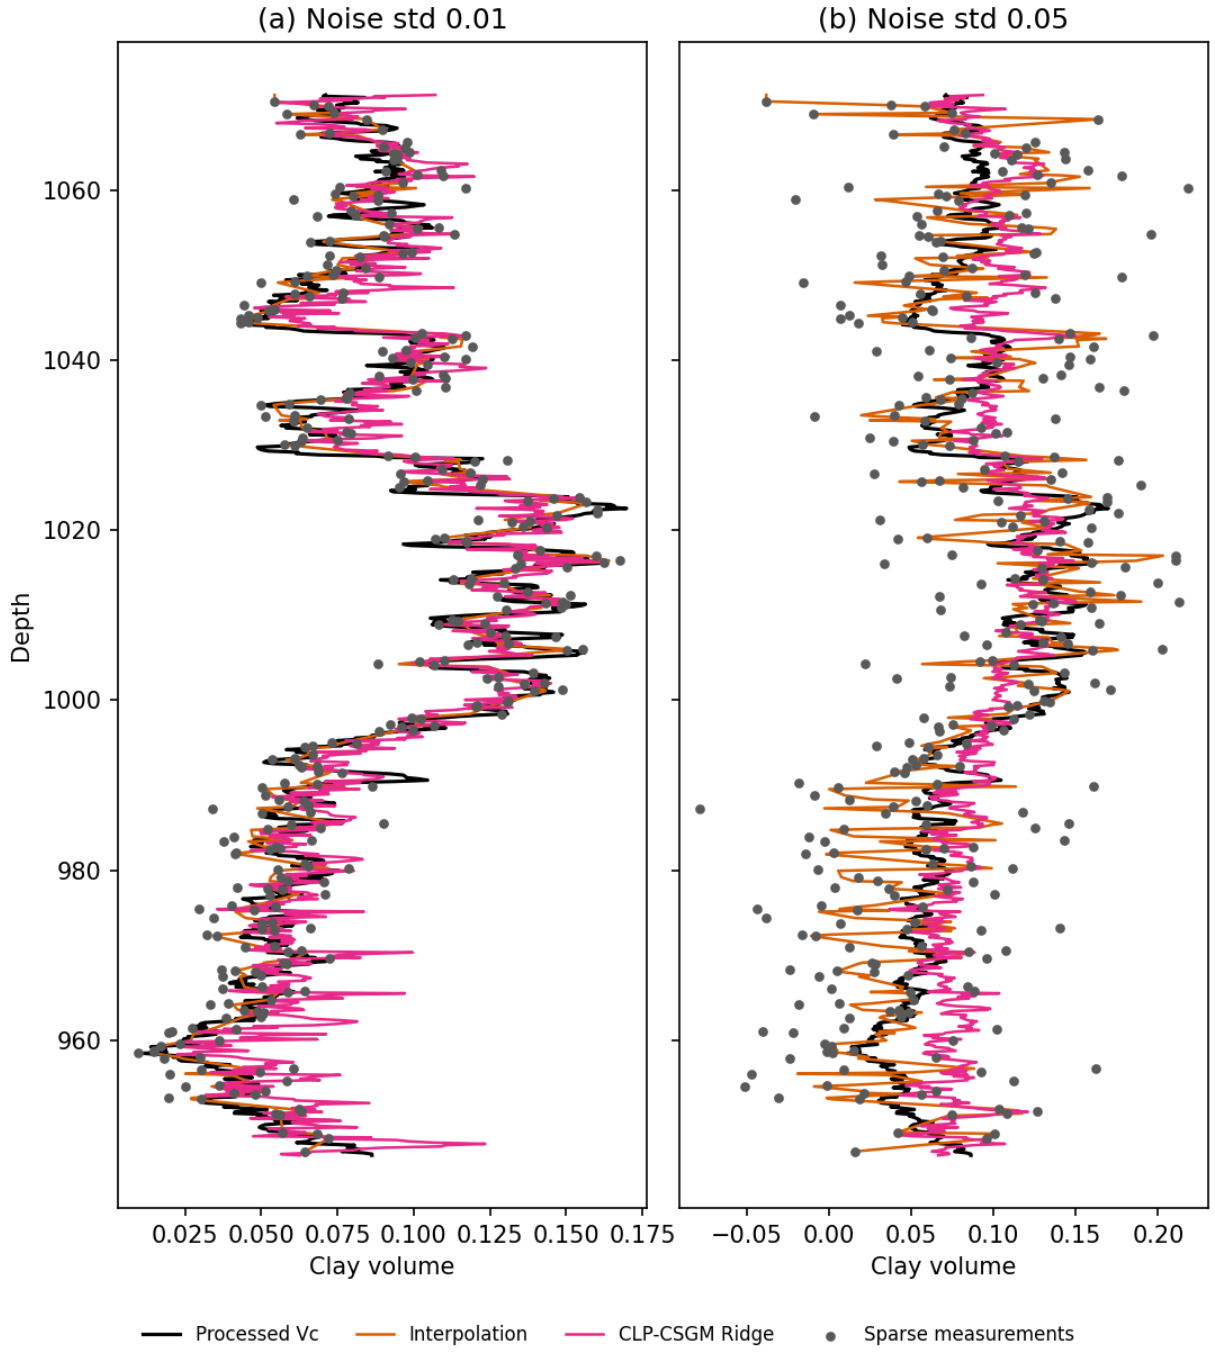


Figure 21: Held-out F03-4 depth tracks, seed 7, step 32, $\rho = 0.20$. Panel (a) uses the published noise standard deviation 0.01. Panel (b) uses 0.05, the same order as the curve's standard deviation 0.035. Each panel has its own clay-volume scale. This interval is one example. Tables 19 and 20 are the comparison.

it is not ranked with the methods above.

The step-8 pattern at noise standard deviation 0.05 is compatible with a limit of the dense-log map. It is not explained by the decoder check alone, and these comparisons do not isolate that map as the only cause. At step 8 the decoder represents the held-out curve with mean test RMSE 0.0073. The window predicted from the dense logs alone has mean RMSE 0.0564. Interpolation at the same step and noise has mean RMSE 0.0456 for $\rho = 0.05$ and 0.0442 for $\rho = 0.10$. Those ratios observe 4 and 6 coordinates in a window of length 64. CLP-CSGM Ridge then stays between the two, at 0.0508 and 0.0456. At $\rho = 0.20$ the same window observes 13 coordinates, and CLP-CSGM Ridge has mean RMSE 0.0372 against an interpolation RMSE of 0.0423. At steps 16 and 32 the dense-log window itself has mean RMSE 0.0426 and 0.0432, already near the interpolation error at this noise. The pattern agrees with the ablation in Section 6.3, where the conditional code $\boldsymbol{z}_0$ is not helpful at step 8.

B.13 Computational Trade-Offs

Table 21 and Fig. 22 quantify the practical computational trade-off. CLP-CSGM Ridge is more expensive than AE $[u, b]$ because it solves a latent inverse problem at validation and test time instead of producing a single direct forward prediction. Table 21 reports both mean wall-clock time per test window and a separate normalized runtime diagnostic relative to AE; the latter is the diagnostic shown in Fig. 22; it is not the exact quotient of the averaged seconds in the previous column. In the representative cross-well case, the additional cost is substantial relative to AE, but it comes with the lowest mean RMSE among the compared methods. In the F03-4 GR-only case, CLP-CSGM Ridge and AE $[u, b]$ have closer average RMSE, so the runtime gap is less clearly justified. The extra optimization is most readily justified when explicit sparse-measurement consistency produces a meaningful reconstruction gain, especially under stronger ill-posedness or weaker direct supervision.

C Supplementary F03-4 GR-only stress tests

The main F03-4 benchmark in Section 6 uses $\rho \in \{0.20, 0.30, \ldots, 0.60\}$. This appendix summarizes controlled supplements on the same well with GR-only u, L=64, depth-first splits before windowing, coordinate subsampling, the same three fixed random-seed realizations, and the same paired

Table 21: Representative computational-cost summary. Time is reported as average wall-clock seconds per test window across the tested measurement ratios. The normalized runtime diagnostic column is computed relative to AE $[u, b]$ within each dataset and should not be read as the exact quotient of the averaged wall-clock times.

Dataset	Method	Time/window (s)	Normalized runtime diagnostic vs. AE	Mean RMSE
Cross-well V_c	ML only	0.00787	0.32	0.02582
Cross-well V_c	MLP $[u, b]$	0.04348	2.33	0.21399
Cross-well V_c	PCA $[u, b]$	0.07129	3.15	0.14464
Cross-well V_c	AE $[u, b]$	0.03156	1.00	0.01793
Cross-well V_c	CLP-CSGM Ridge	0.23517	9.79	0.01480
F03-4 GR-only	ML only	0.06010	2.13	0.02391
F03-4 GR-only	MLP $[u, b]$	0.04771	1.73	0.03981
F03-4 GR-only	PCA $[u, b]$	0.31793	11.67	0.02149
F03-4 GR-only	AE $[u, b]$	0.02999	1.00	0.01429
F03-4 GR-only	CLP-CSGM Ridge	0.08838	3.23	0.01418

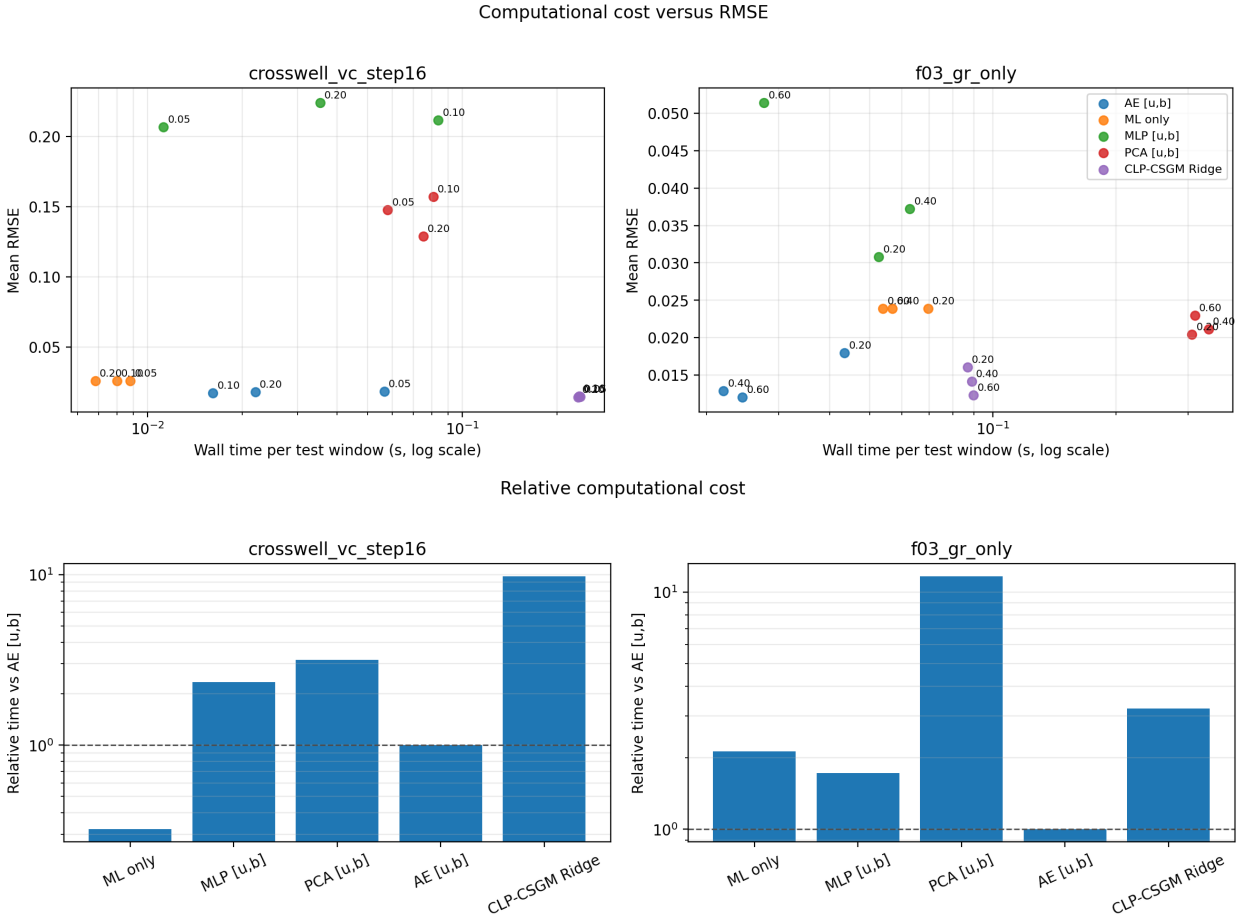


Figure 22: Representative computational-cost diagnostics. The upper panel compares mean RMSE with wall-clock time per test window, and the lower panel reports the normalized runtime diagnostic relative to AE $[u, b]$ within each dataset.

measurement protocol as Section 5. To align the grid with the sparse end of the cross-well study, we use $\rho \in \{0.05, 0.10, 0.20\}$ and compare CLP-CSGM Ridge with AE $[u, b]$.

Cases A–D differ only by the stress axes enumerated below:

- **A (baseline noise):** η=0.01 on b=$My + \eta$; standard depth fractions 0.6/0.2/0.2; step 1.
- **B (noisy b):** η=0.05; otherwise as A.
- **C (smaller train pool):** depth fractions 0.45/0.25/0.30; η=0.01; step 1.
- **D (coarse sliding step):** step 4 with standard fractions; η=0.01. This yields far fewer test windows (n_{test}=56 versus 224 in A/B), yielding high-variance paired RMSE summaries at each ρ.

Table 22: F03-4 GR-only stress tests: mean RMSE across seeds for AE $[u, b]$ versus CLP-CSGM Ridge. Relative RMSE gap is $(\text{RMSE}_{\text{AE}} - \text{RMSE}_{\text{Ridge}})/\text{RMSE}_{\text{AE}}$ in percent (positive favors Ridge). Test counts n_{test} count windows in the test depth segment after sliding-window construction.

Case	ρ	AE $[u, b]$	CLP-CSGM Ridge	Rel. RMSE gap (%)	n_{test}
A	0.05	0.02011	0.01682	+16.36	224
	0.10	0.01900	0.01596	+15.98	224
	0.20	0.01489	0.01314	+11.74	224
B	0.05	0.02444	0.01919	+21.51	224
	0.10	0.02397	0.01897	+20.87	224
	0.20	0.02125	0.01846	+13.13	224
C	0.05	0.01810	0.01596	+11.84	369
	0.10	0.01712	0.01509	+11.85	369
	0.20	0.01433	0.01243	+13.26	369
D	0.05	0.01676	0.01672	+0.24	56
	0.10	0.01493	0.01547	-3.62	56
	0.20	0.01115	0.01298	-16.33	56

In A–C, CLP-CSGM Ridge has lower mean RMSE than AE $[u, b]$ at every tested ρ, including under larger η and smaller training fractions. In D, mean RMSE crosses over at $\rho \in \{0.10, 0.20\}$ and is similar to or slightly lower at ρ=0.05. The small effective test count is consistent with this behavior because aggressive overlap reduction amplifies seed variability. The A–D results position the coarse-step F03 stress as an exploratory supplement to the primary ρ sweep in Table 6.

D Supplementary Lapa Pre-salt Permeability Case Study

This appendix reports a complementary case study built from the well-log and permeability data context described by Alvarez et al. [2025] for the Lapa pre-salt carbonate reservoir in the Santos Basin. Geologically, the Lapa oilfield targets the Aptian Barra Velha Formation, a heterogeneous pre-salt carbonate succession in which laminated microbialites, spherulite, stromatolite, rudstone, and grainstone facies coexist. In such rocks, permeability is controlled by pore structure and diagenesis as much as by porosity. Permeability is therefore the sparse reconstruction target in this supplementary check. Full core–log integration, hydraulic flow units, and nonlinear permeability prediction details are reported in Alvarez et al. [2025]. The benchmark adaptation uses the same CLP-CSGM protocol as the main paper: sliding windows, contiguous train/validation/test split, coordinate subsampling for sparse target observations, paired seeds and measurement ratios, and matched direct $[u, b]$ baselines.

This limited-sample case provides supplementary evidence. The configuration uses $L = 16$, step $= 4$, seven dense channels in u, target k_{lab} in $\log_{10}$ scale, $\rho \in \{0.20, 0.30, 0.40, 0.50, 0.60\}$, and the same three fixed random-seed realizations used only for reproducible paired averaging.

Table 23 shows that both CLP-CSGM prior variants remain competitive with direct $[u, b]$ baselines across the tested ρ grid. CLP-CSGM MLP and CLP-CSGM Ridge have lower mean RMSE than AE $[u, b]$, PCA $[u, b]$, ML only, and direct MLP $[u, b]$ at all reported ratios. The comparison with measurement-only CSGM is ratio-dependent: measurement-only has lower mean RMSE at some lower or intermediate ratios, while the prior-guided variants have lower mean RMSE at others. This variation is consistent with the expected regime dependence when sparse measurements are informative but the sample size is limited.

E Laboratory porosity on a cored interval

This supplement uses the 87 laboratory-porosity samples of well 861 over 5205.91–5233.72 m. The target is Phi_{lab} as recorded at those depths. It is not a porosity curve computed from density or sonic. Dense context is gamma ray, deep resistivity, shallow resistivity, and lithotype. Density, neutron porosity, sonic porosity, and Phi_{ND} are excluded because they track Phi_{lab}.

Depth is split into three contiguous blocks before windowing, so training and test windows do

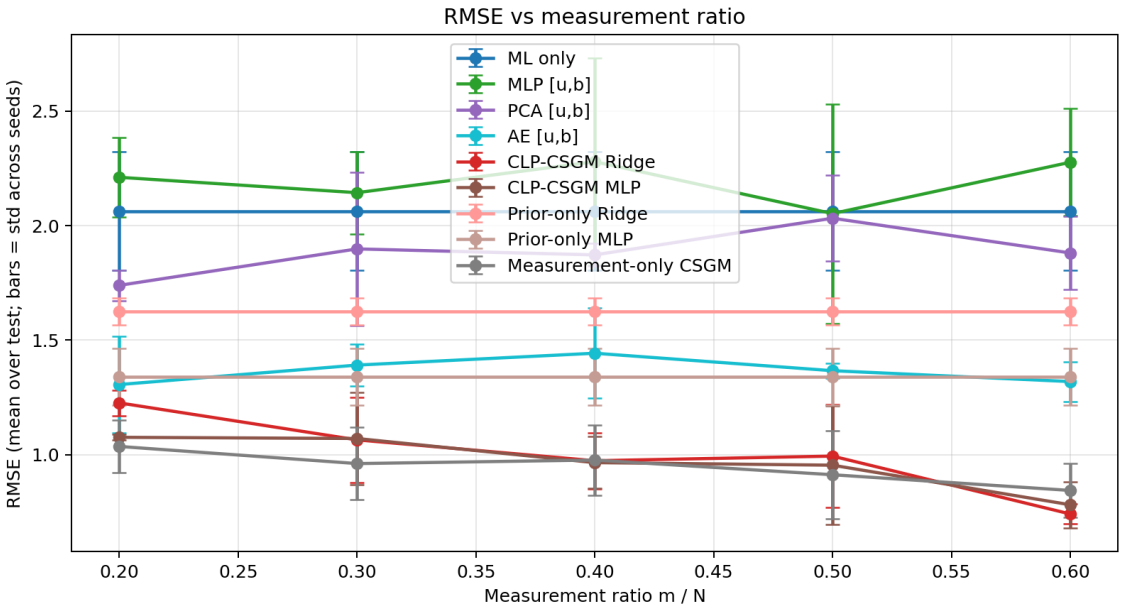


Figure 23: Supplementary Lapa case study RMSE by measurement ratio. Prior-guided CLP-CSGM variants remain competitive across the sparse-measurement grid.

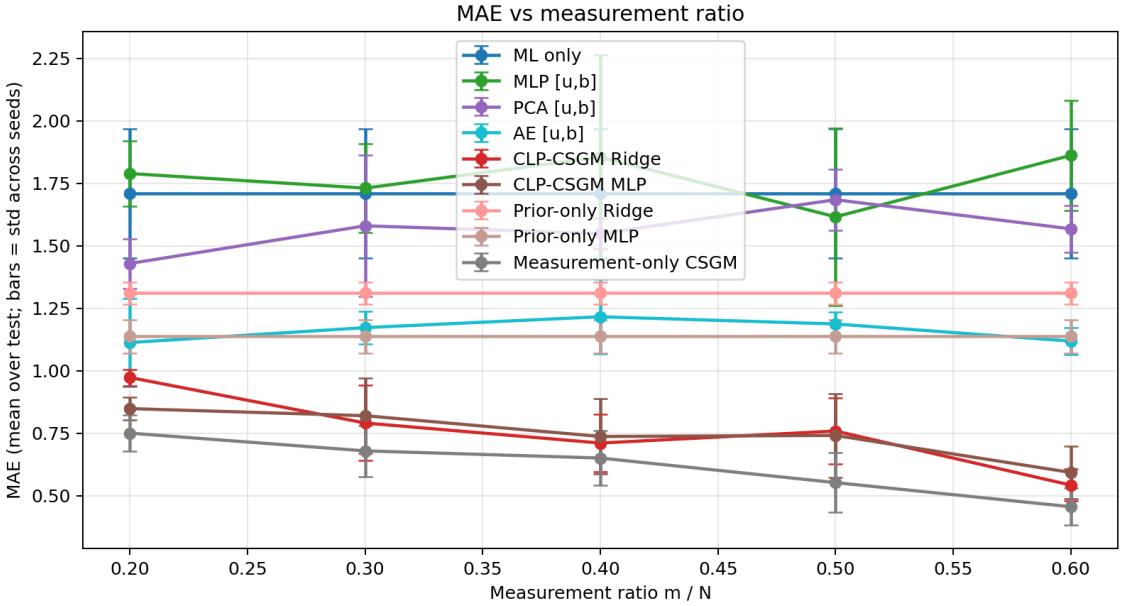


Figure 24: Supplementary Lapa case study MAE by measurement ratio. The ranking pattern is consistent with the RMSE view.

Table 23: Supplementary Lapa case study by measurement ratio. Values are means over three seeds.

rho	method	rmse	mae	relative_l2
0.20000	CLP-CSGM MLP	1.07615	0.84823	1.00530
0.20000	CLP-CSGM Ridge	1.22639	0.97273	1.08660
0.20000	Measurement-only CSGM	1.03628	0.75013	0.93124
0.20000	AE [u,b]	1.30639	1.11253	1.22956
0.20000	PCA [u,b]	1.73942	1.42881	1.60571
0.20000	ML only	2.06280	1.70838	2.00587
0.20000	MLP [u,b]	2.21106	1.78860	2.13801
0.30000	CLP-CSGM MLP	1.07099	0.81987	0.96733
0.30000	CLP-CSGM Ridge	1.06503	0.79065	0.91318
0.30000	Measurement-only CSGM	0.96152	0.67917	0.86494
0.30000	AE [u,b]	1.39133	1.17212	1.30771
0.30000	PCA [u,b]	1.89848	1.57954	1.72206
0.30000	ML only	2.06280	1.70838	2.00587
0.30000	MLP [u,b]	2.14391	1.73034	2.00236
0.40000	CLP-CSGM MLP	0.96555	0.73670	0.88465
0.40000	CLP-CSGM Ridge	0.97383	0.71050	0.86519
0.40000	Measurement-only CSGM	0.97680	0.65052	0.86968
0.40000	AE [u,b]	1.44311	1.21575	1.35981
0.40000	PCA [u,b]	1.87246	1.54857	1.79426
0.40000	ML only	2.06280	1.70838	2.00587
0.40000	MLP [u,b]	2.27913	1.85405	2.18847
0.50000	CLP-CSGM MLP	0.95447	0.74089	0.84011
0.50000	CLP-CSGM Ridge	0.99413	0.75851	0.85865
0.50000	Measurement-only CSGM	0.91277	0.55221	0.80389
0.50000	AE [u,b]	1.36684	1.18683	1.24828
0.50000	PCA [u,b]	2.03242	1.68385	1.85003
0.50000	ML only	2.06280	1.70838	2.00587
0.50000	MLP [u,b]	2.05267	1.61547	2.03388
0.60000	CLP-CSGM MLP	0.78164	0.59283	0.70811
0.60000	CLP-CSGM Ridge	0.74144	0.54218	0.67252
0.60000	Measurement-only CSGM	0.84417	0.45574	0.74392
0.60000	AE [u,b]	1.31942	1.11830	1.22740
0.60000	PCA [u,b]	1.88084	1.56736	1.77341
0.60000	ML only	2.06280	1.70838	2.00587
0.60000	MLP [u,b]	2.27679	1.86123	2.16631

not share rows. Windows have length 8 and step 1. The last 16 rows of the deepest training block form the validation segment. Seeds are $\{7, 23, 41\}$. For each fold, seed, and ratio, one coordinate set is shared by every window. The laboratory value on those coordinates is the measurement, with no added noise. RMSE is computed only on the remaining coordinates. $\rho = 0.25$ uses $m = 2$ measurements and $\rho = 0.50$ uses $m = 4$. The direct AE receives $[\boldsymbol{u}, \boldsymbol{b}]$ with the same coordinates. The decoded prediction of the conditional prior is $G(h(\boldsymbol{u}))$, with no latent search. It is reported below as the log-only anchor. Neither method is added to the main F3 tables.

Table 24 shows that the direct AE has the lower mean held-out RMSE at both ratios. CLP-

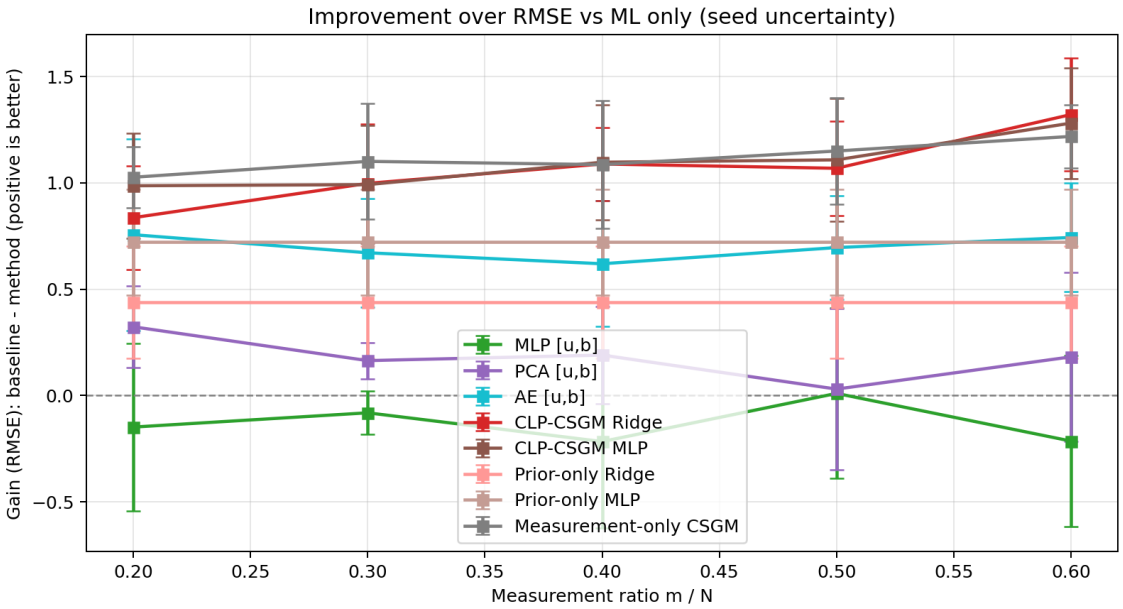


Figure 25: Supplementary Lapa case study RMSE gain over ML-only. Positive values indicate improvement over dense-context-only prediction.

Table 24: Well 861 laboratory-porosity supplement. Entries are the mean held-out RMSE across three depth blocks and three seeds ($n = 9$). Δ is AE minus CLP-CSGM Ridge; a negative value means the direct AE has the lower mean. CLP>AE counts splits in which CLP-CSGM Ridge has the lower RMSE.

ρ	CLP	AE $[u, b]$	prior-only	Δ AE	CLP>AE
0.25	0.0602	0.0537	0.0729	-0.0065	5/9
0.50	0.0580	0.0523	0.0753	-0.0057	6/9

CSGM Ridge has the lower RMSE in five of nine splits at $\rho = 0.25$ and in six of nine at $\rho = 0.50$, and it has lower RMSE than the log-only anchor in all nine splits at both ratios. The two summaries answer different questions. The mean asks which method is lower after every split is weighted equally. The split count asks how often CLP-CSGM Ridge is the lower of the two. They disagree because one depth block carries errors large enough to set the mean.

Table 25: Well 861 held-out laboratory-porosity RMSE by depth block. Each entry is the mean across the three seeds. Prior-only is $G(h(\boldsymbol{u}))$ with no measurement term.

Block	ρ	CLP	AE $[u, b]$	prior-only
5205.91–5214.48 m	0.25	0.0338	0.0371	0.0486
5205.91–5214.48 m	0.50	0.0228	0.0350	0.0472
5214.67–5223.63 m	0.25	0.0536	0.0548	0.0559
5214.67–5223.63 m	0.50	0.0497	0.0543	0.0554
5224.39–5233.72 m	0.25	0.0931	0.0691	0.1142
5224.39–5233.72 m	0.50	0.1016	0.0676	0.1234

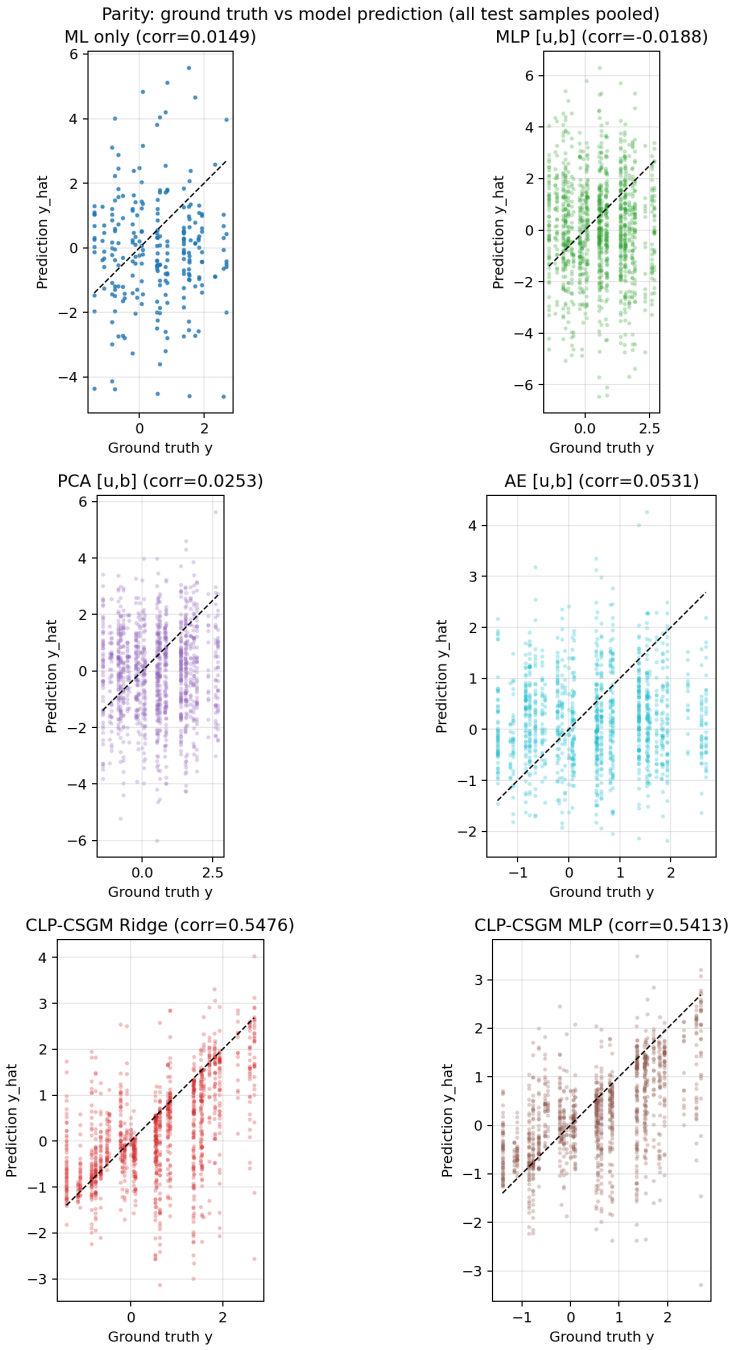


Figure 26: Supplementary Lapa pooled parity plot. The visual agreement is consistent with the quantitative summary, while still reflecting variability expected under small test counts.

Table 25 locates the tradeoff. On 5205.91–5214.48 m, CLP-CSGM Ridge has the lower block-mean RMSE at both ratios, and raising ρ from 0.25 to 0.50 lowers that RMSE from 0.0338 to 0.0228 while the direct AE stays near 0.035. On 5214.67–5223.63 m the two methods are close, with CLP-CSGM Ridge still slightly lower. On 5224.39–5233.72 m the direct AE is lower by a wide margin (0.0691 and 0.0676 against 0.0931 and 0.1016).

The gain where CLP-CSGM Ridge is ahead has two parts. First, the measurement term is doing the work the method is built for: in every block and at both ratios, CLP-CSGM Ridge has lower RMSE than the log-only anchor. The laboratory coordinates pull the latent code off $h(\boldsymbol{u})$. Second, that pull is most useful when $\boldsymbol{z}_0$ is already near a useful code. In the shallowest block, mean $\mathrm{Phi}_{\mathrm{lab}}$ is 0.092 with standard deviation 0.019, mean gamma ray is about 15 API, and the median deep resistivity is about 1900. Adding two further laboratory coordinates ($\rho = 0.50$ rather than 0.25) then cuts the CLP-CSGM Ridge error by about one third. The direct AE already receives those coordinates as input features, so the same increase in ρ changes it little. The explicit residual, not a larger training pool, is what improves CLP-CSGM Ridge there. Each test fold still has 28 to 35 overlapping training windows, so this is not the scarce-window regime of the F3 steps 16 and 32.

The loss on the mean comes from the deepest block, and it is a context shift relative to the two blocks that train the anchor when that block is held out. Mean $\mathrm{Phi}_{\mathrm{lab}}$ rises to 0.145, mean gamma ray rises to about 30 API, and the median deep resistivity falls to about 600. The lithotype mix also changes: codes 2 and 4 account for 25 of the 29 rows, whereas the middle block is mostly code 1. The ridge anchor is a linear map from $\boldsymbol{u}$ fitted on the other blocks, so the decoded prediction $G(\boldsymbol{z}_0)$ differs substantially from the target windows. With only $m = 2$ or $m = 4$ measurements in a length-8 window, the residual still improves on that anchor (0.0931 versus 0.1142, and 0.1016 versus 0.1234) but does not catch the direct AE. The AE never has to pass through the shifted anchor: the same $\boldsymbol{b}$ enters as features and the map is applied in one step. Windows overlap inside each block, so the block contrast, not the count of windows, is the depth evidence.

The tradeoff is therefore where the anchor sits. CLP-CSGM Ridge gains when dense logs seen in training still describe the test interval and the laboratory coordinates can finish the correction. It loses the pooled mean when the test logs move outside that support and only a few measurements are available to repair $\boldsymbol{z}_0$. The result does not show that the measurement term fails, and it does not extend the F3 low-support advantage to a mean win on this laboratory target. It is one cored

interval.